

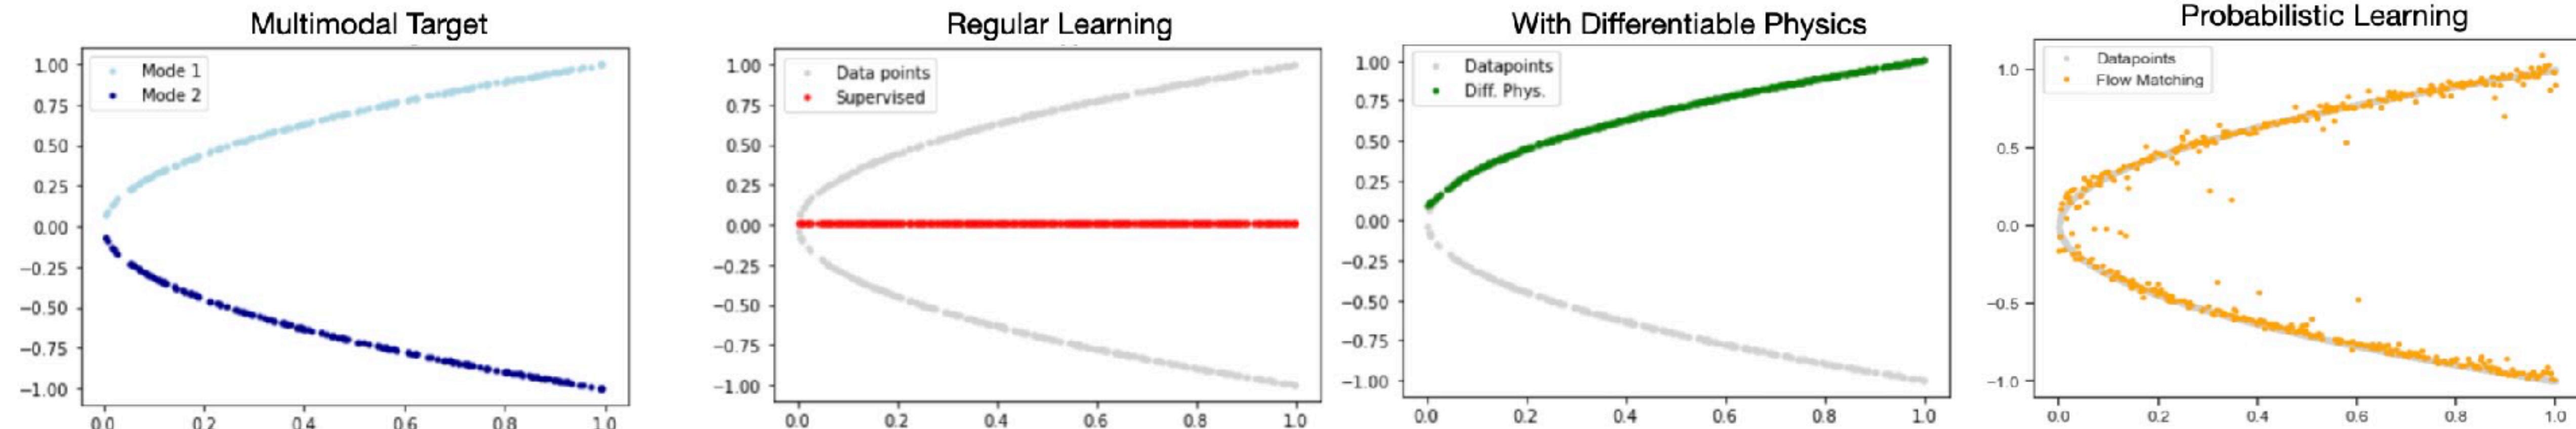
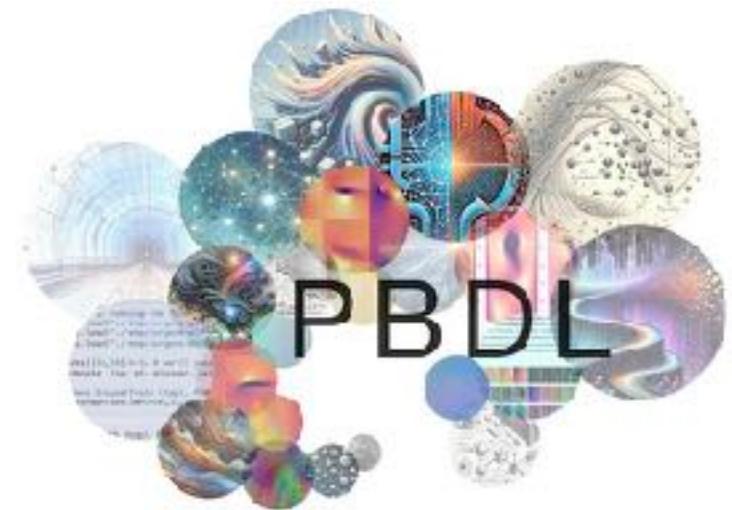


Supervised Learning

ADVANCED DEEP LEARNING FOR PHYSICS

Teaser Example from PBDL

- DL extremely powerful...
- ... but sometimes surprisingly wrong. Solve map: $y^2 \rightarrow x$



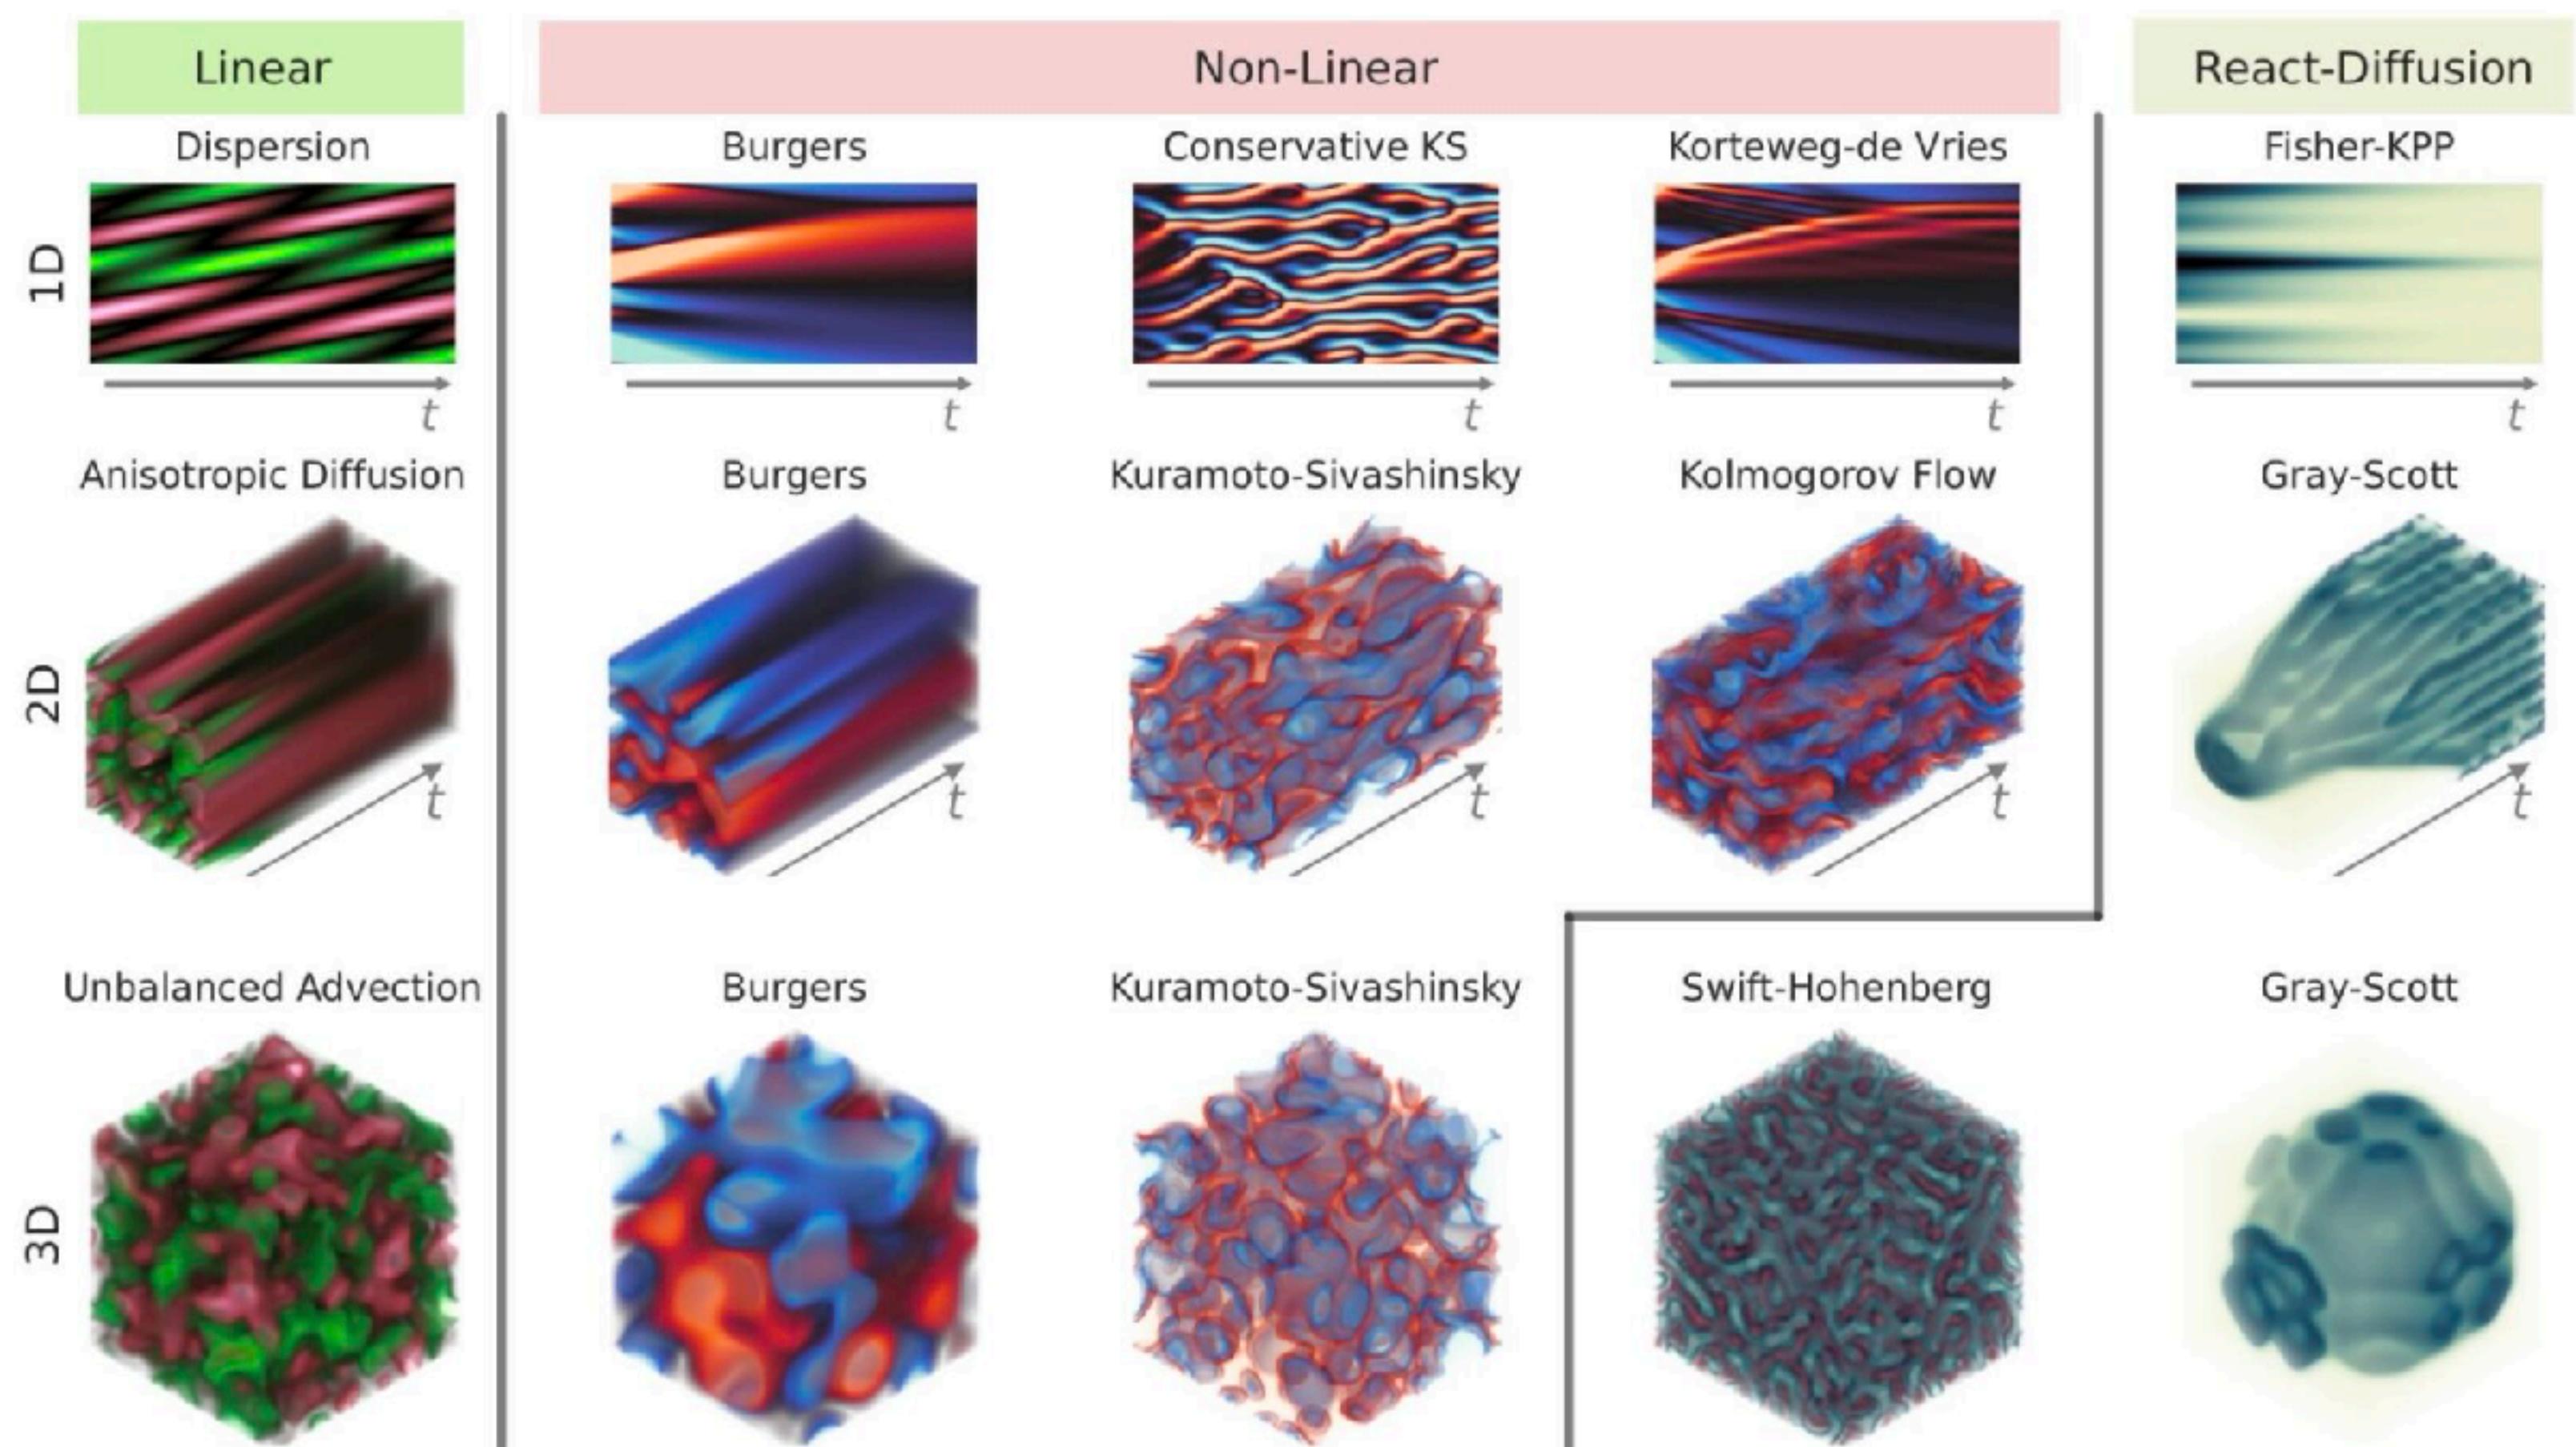
<https://colab.research.google.com/github/tum-pbs/pbdl-book/blob/main/intro-teaser.ipynb>

Model Equations

Model Equations

Basic PDEs

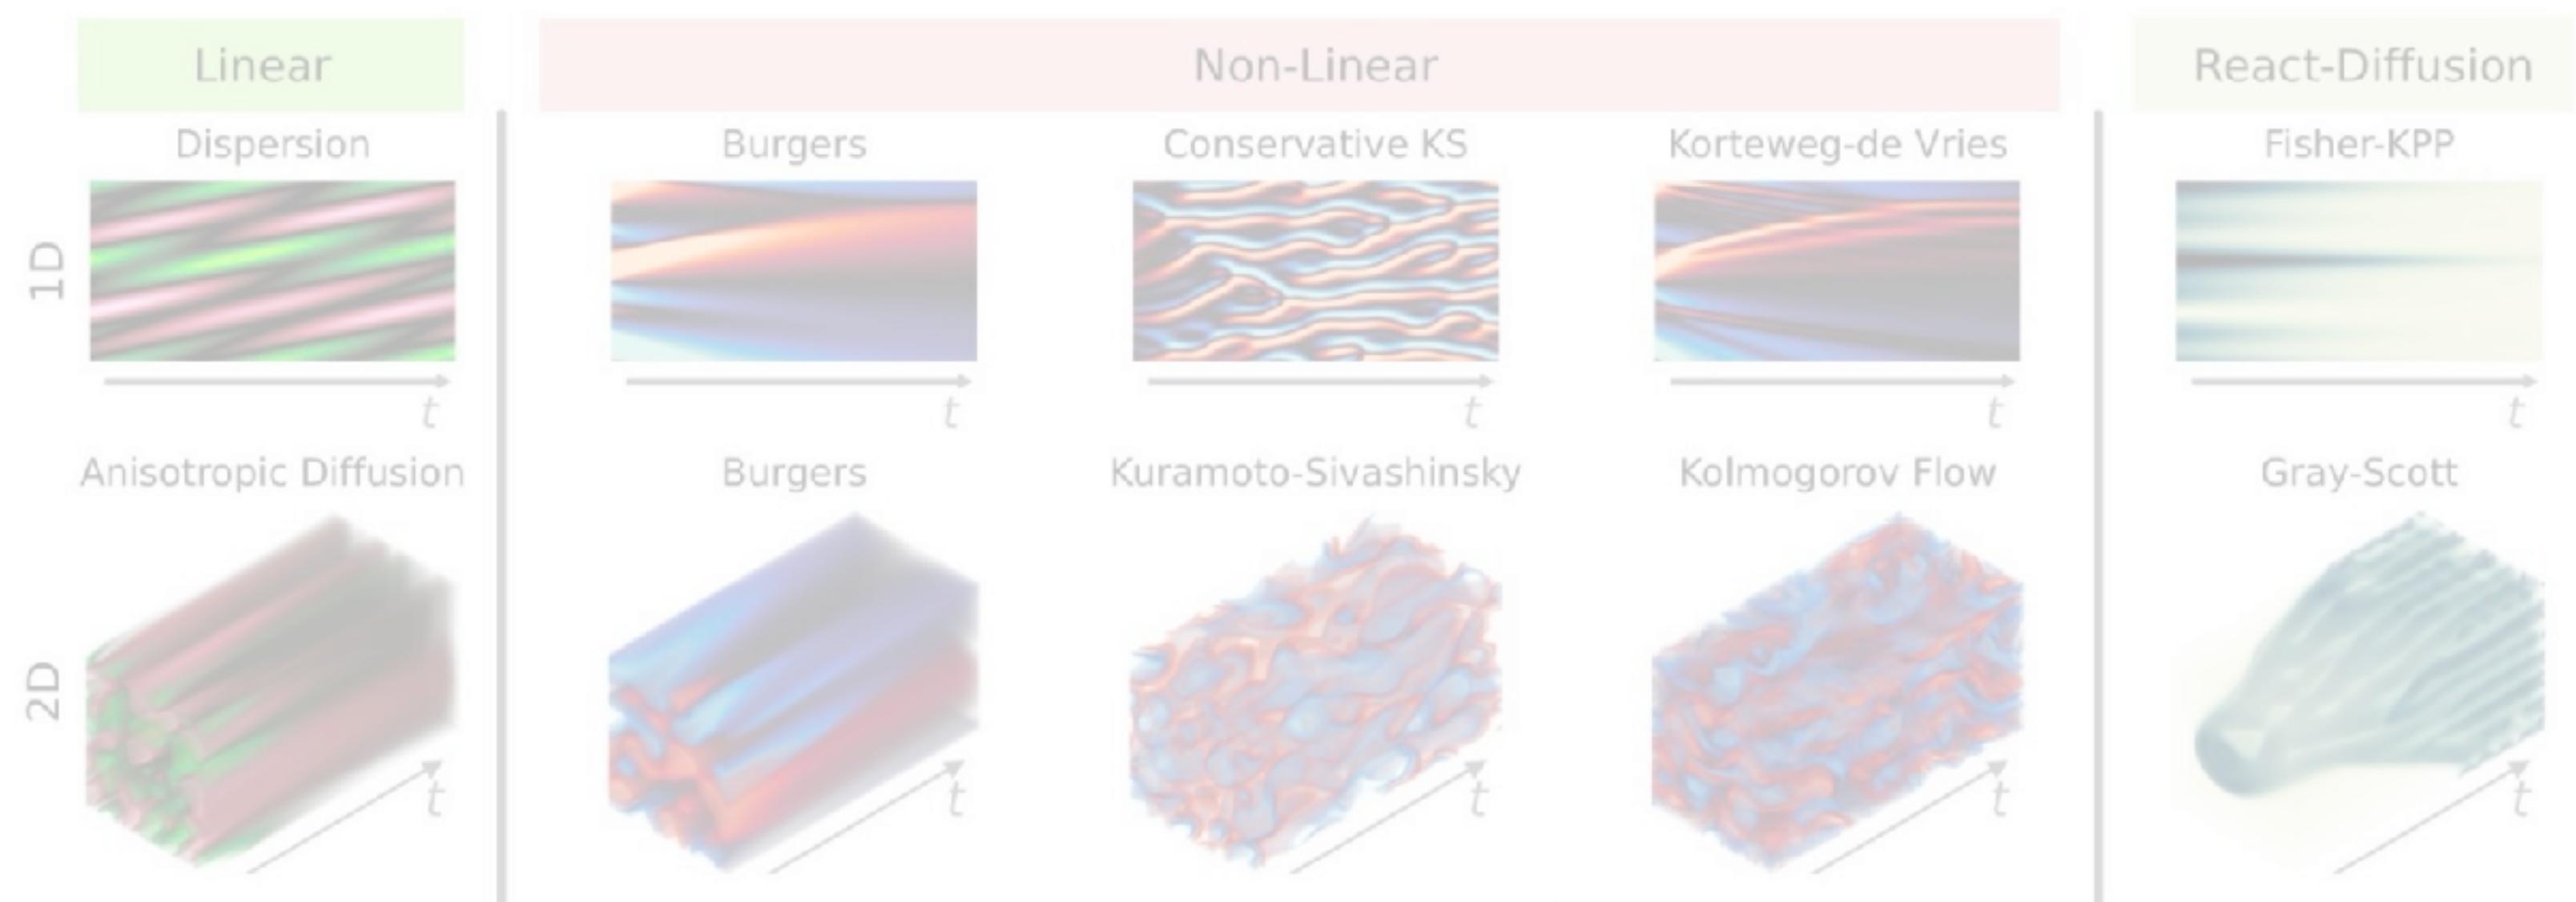
- Diffusion
- Burgers
- Navier-Stokes



Model Equations

Basic PDEs

- Diffusion
- Burgers
- Navier-Stokes



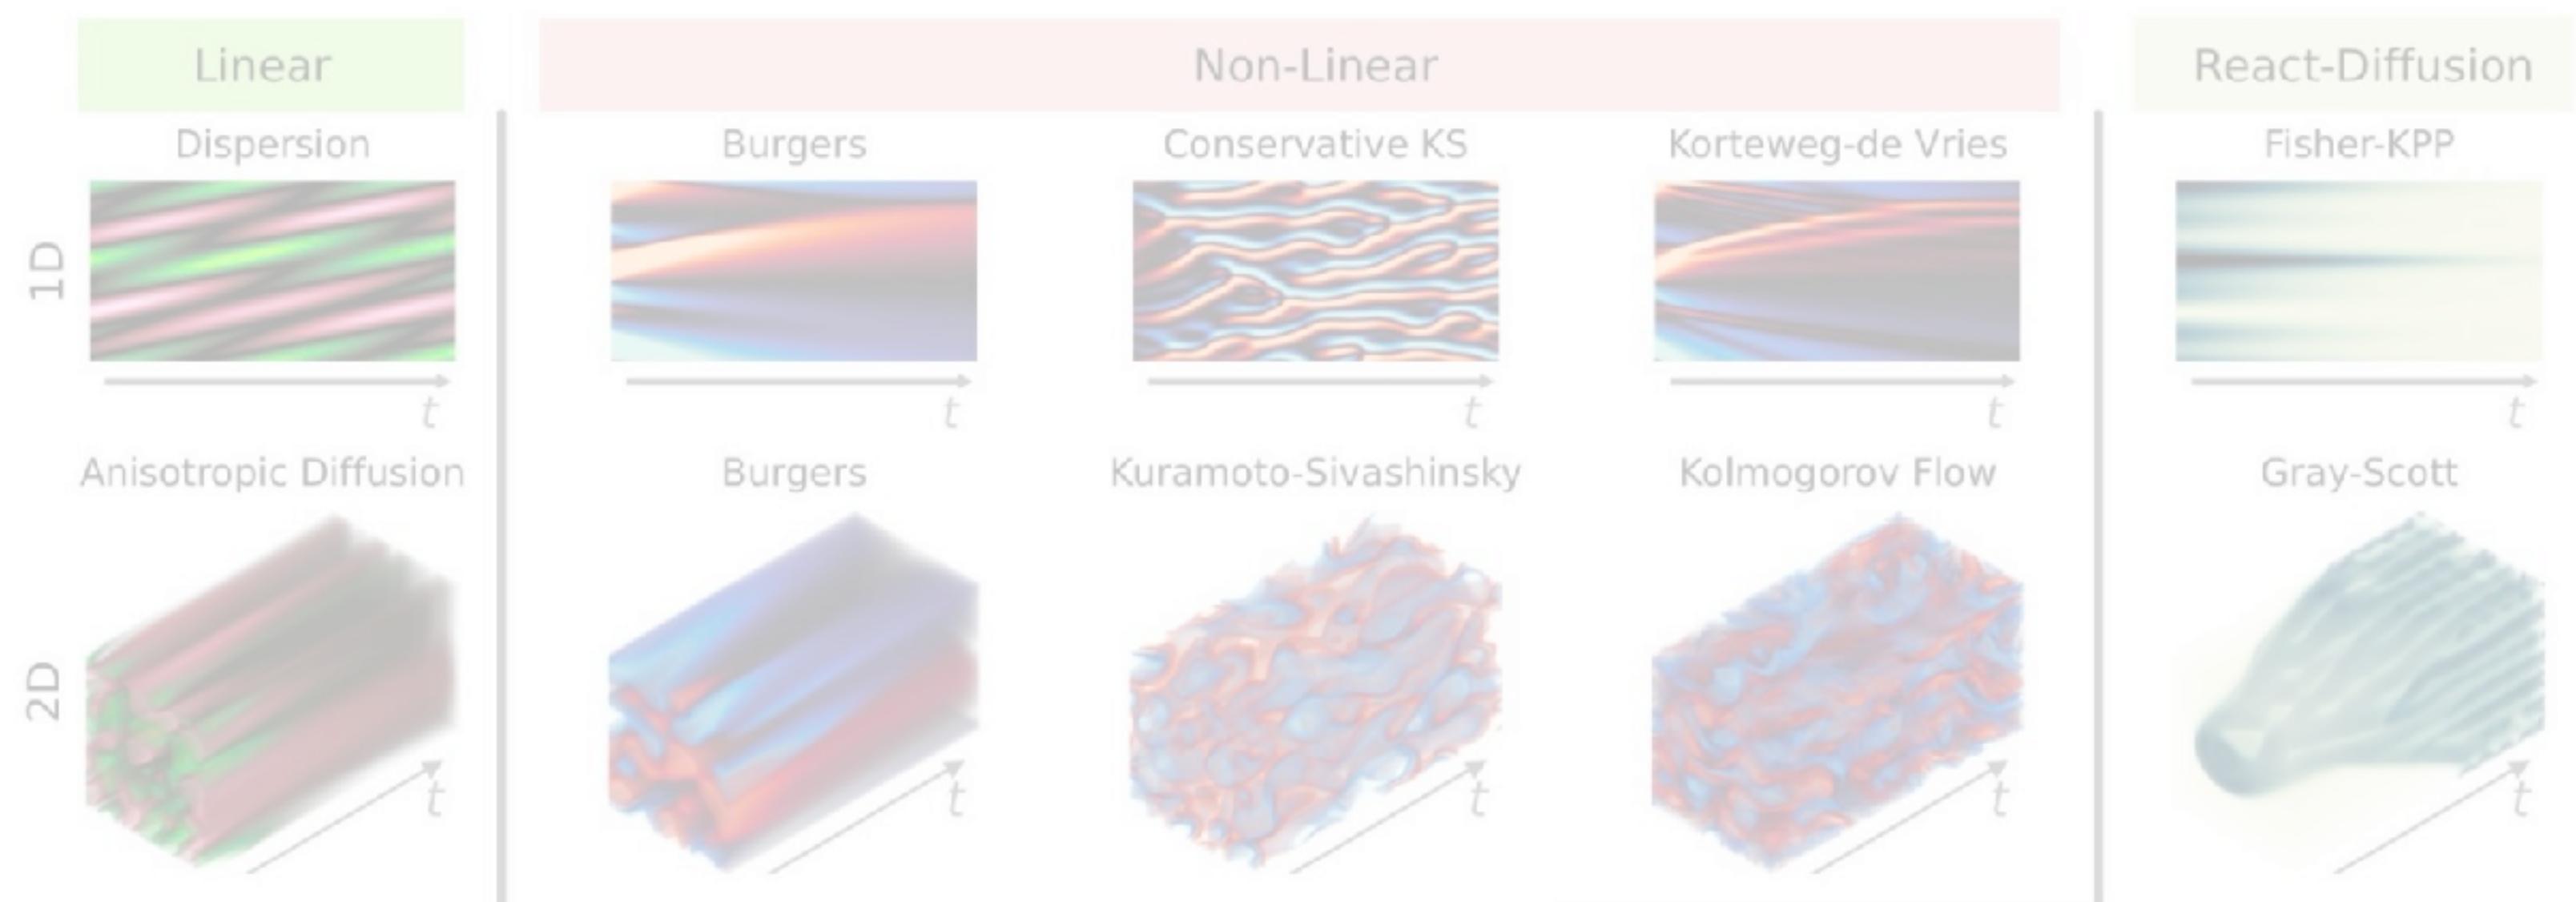
$$\frac{\partial u}{\partial t} - \alpha \nabla^2 u = 0$$

Diffusion constant α

Model Equations

Basic PDEs

- Diffusion
- Burgers (in 2D)
- Navier-Stokes



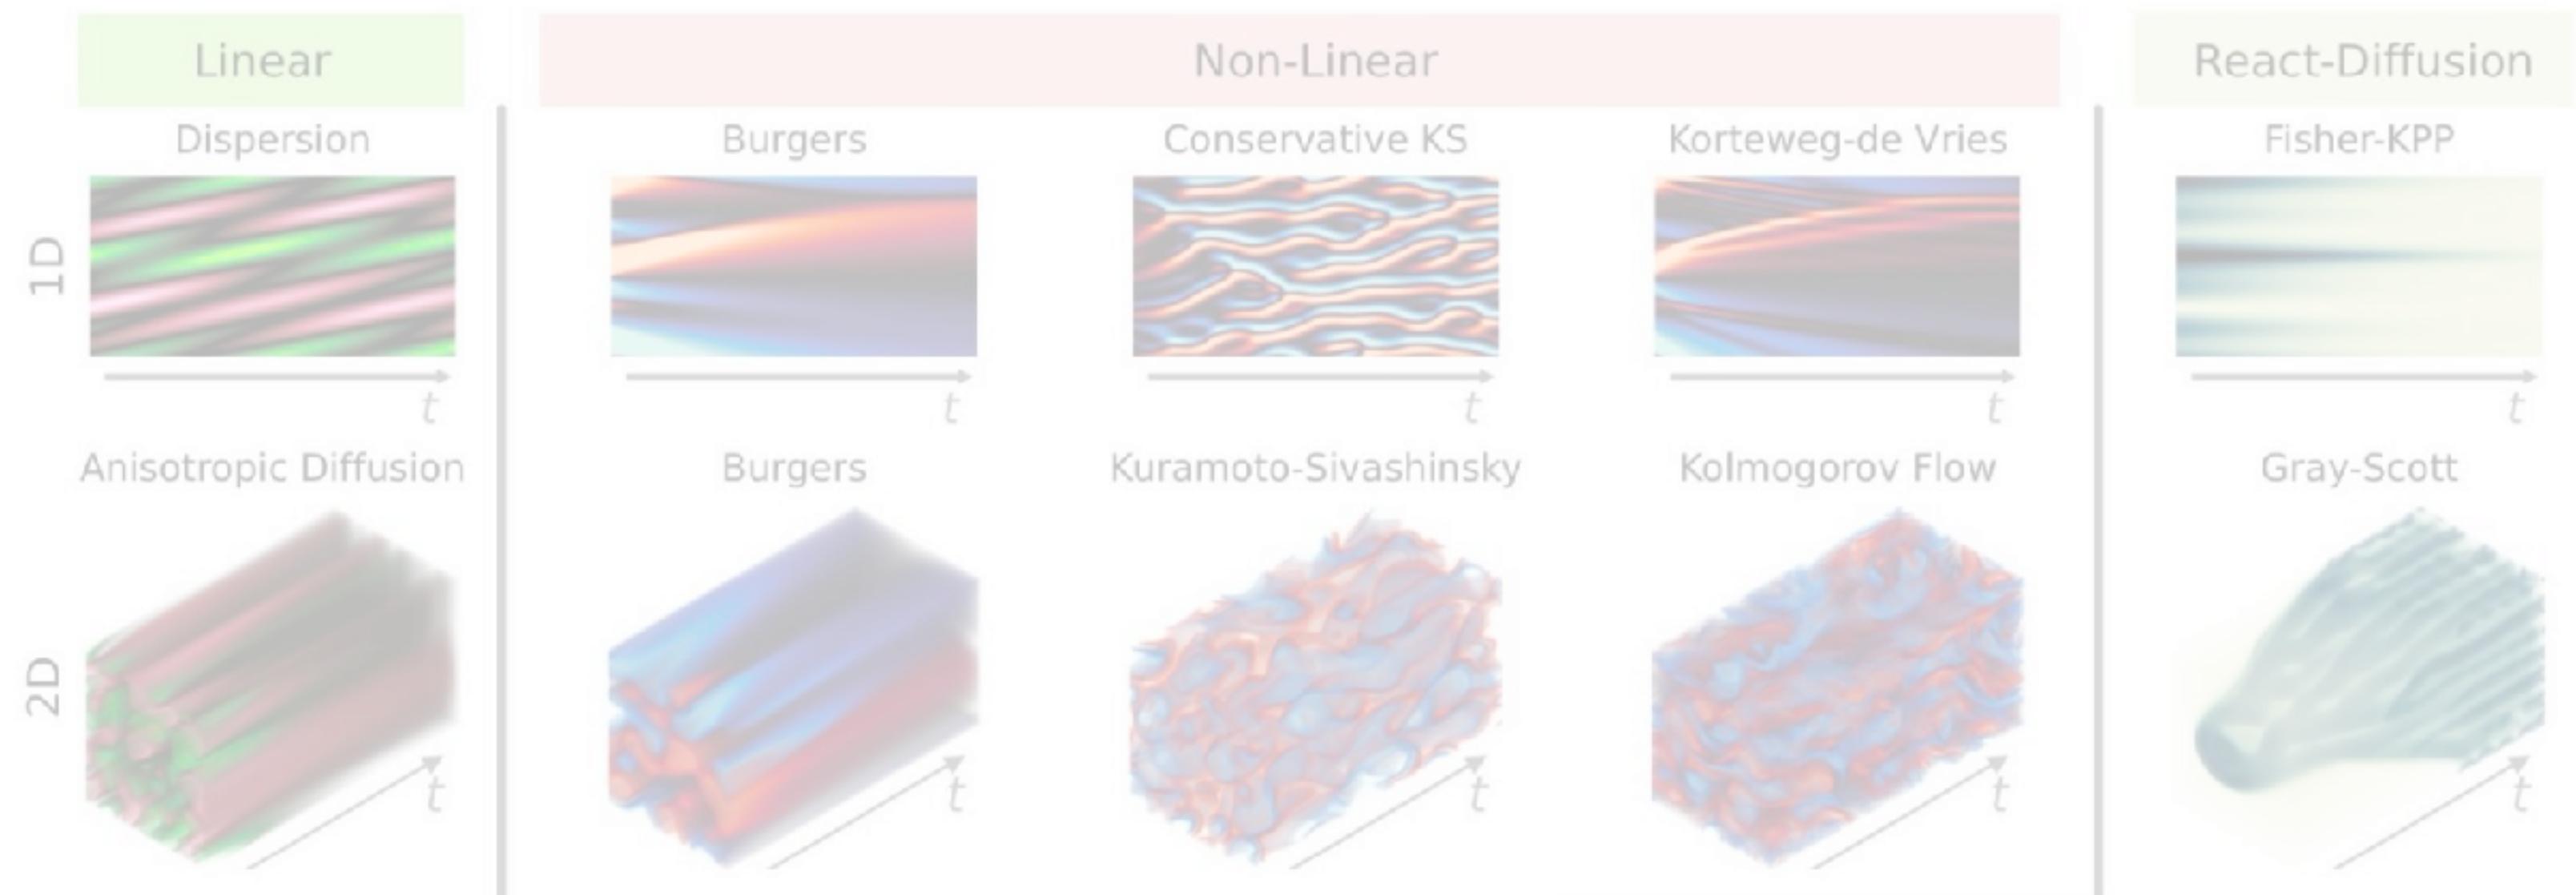
$$\frac{\partial u_x}{\partial t} + \mathbf{u} \cdot \nabla u_x = \nu \nabla \cdot \nabla u_x$$
$$\frac{\partial u_y}{\partial t} + \mathbf{u} \cdot \nabla u_y = \nu \nabla \cdot \nabla u_y$$

Kinematic Viscosity ν

Model Equations

Basic PDEs

- Diffusion
- Burgers
- Navier-Stokes (2D)



$$\frac{\partial u_x}{\partial t} + \mathbf{u} \cdot \nabla u_x = -\frac{1}{\rho} \nabla p + \nu \nabla \cdot \nabla u_x + g_x$$

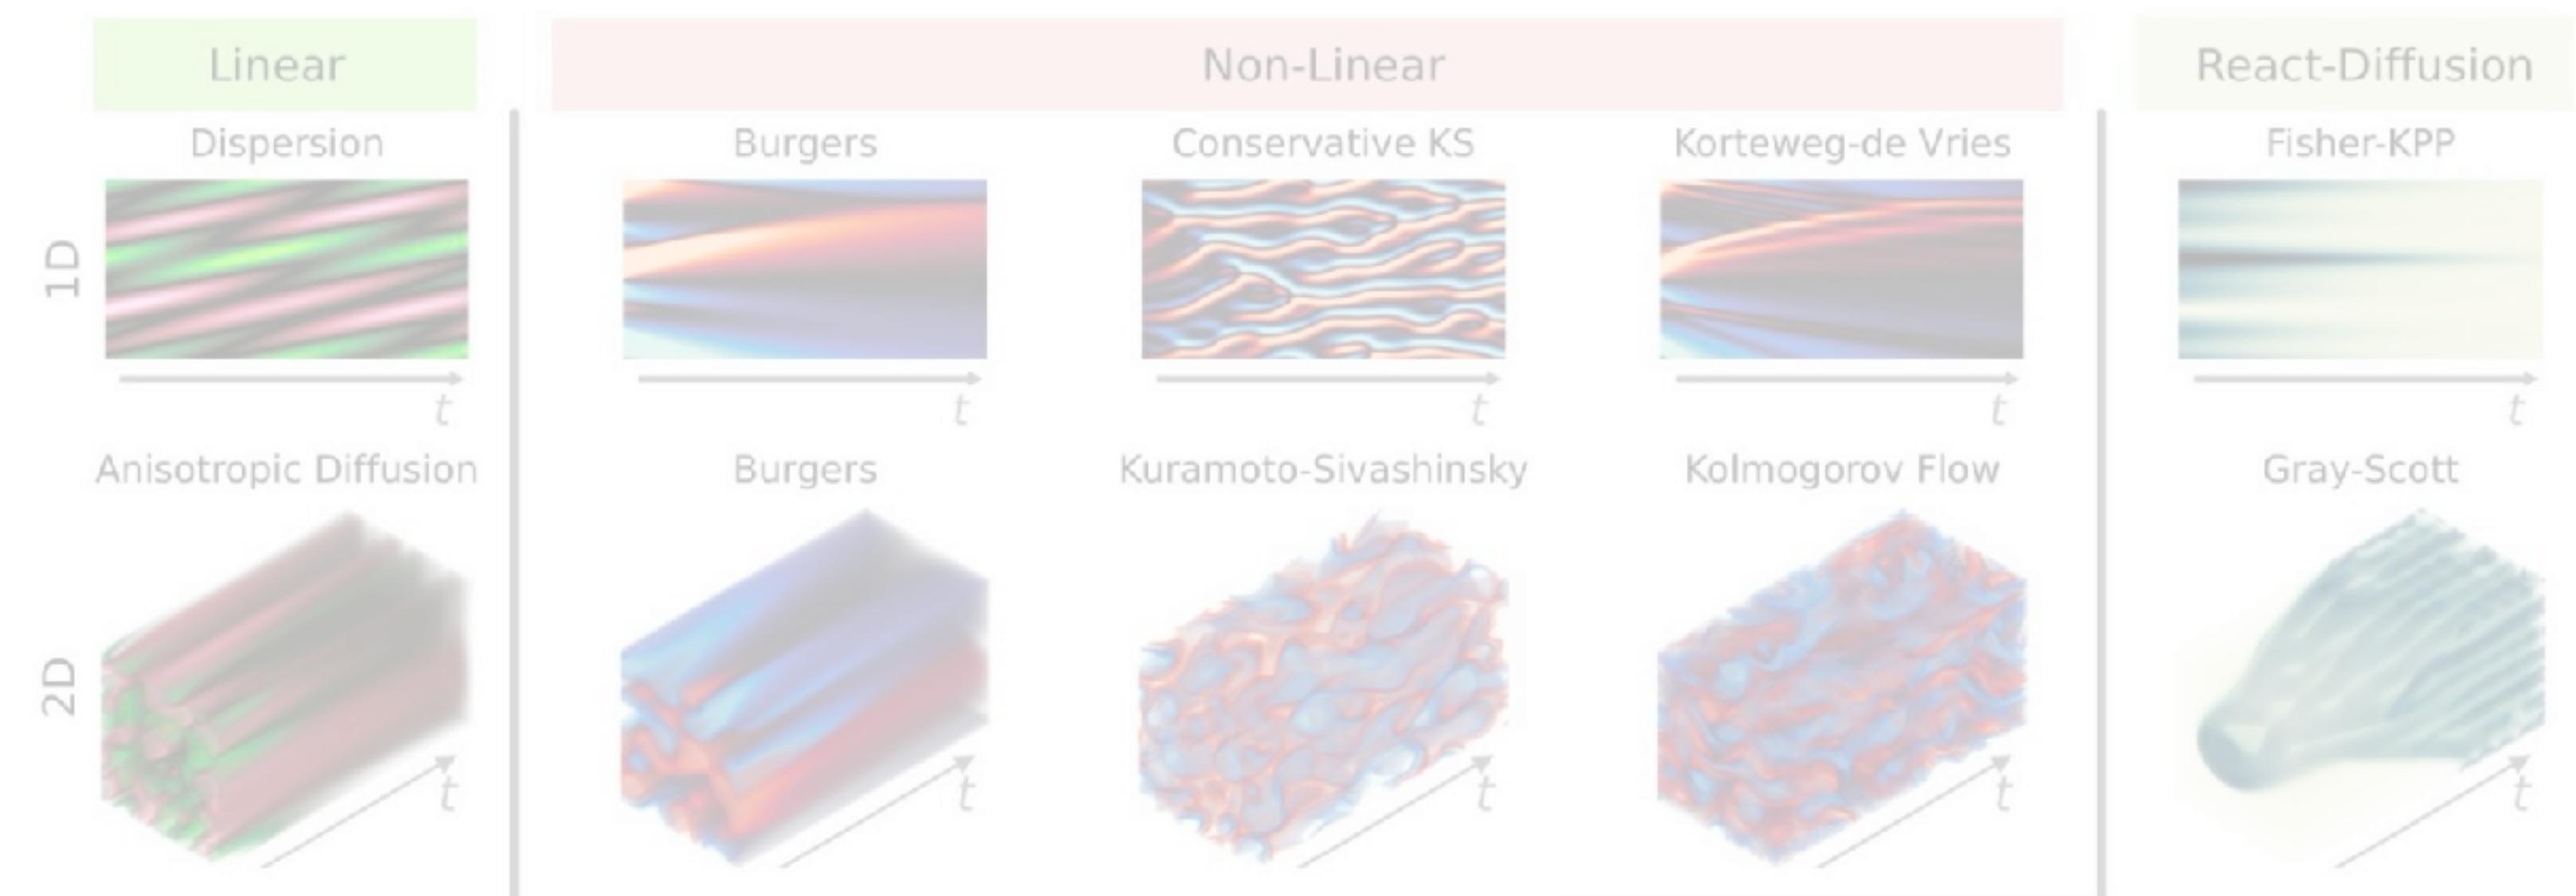
$$\frac{\partial u_y}{\partial t} + \mathbf{u} \cdot \nabla u_y = -\frac{1}{\rho} \nabla p + \nu \nabla \cdot \nabla u_y + g_y$$

$$\text{s.t. } \nabla \cdot \mathbf{u} = 0$$

Model Equations

Basic PDEs

- Diffusion
- Burgers
- Navier-Stokes (2D)



*Distinguish **forward** and **inverse** problems*

Forward: initial & boundary conditions, solve from time t_0 to end time

Inverse: from data/observations solve for state (e.g., $\mathbf{u}(t_0)$) or parameter (e.g., viscosity ν)

Supervised Learning - The Basics

Deep Learning Basics

- Approximate unknown function $f^*(x) = y^*$
- Star super-script * denotes ground truth (often intractable)
- Find approximation $f(x)$ over training data set with (x_i, y_i^*) pairs
- Minimizing error $e(x, y)$
- In the simplest case L^2 : $\arg \min_{\theta} \|f(x; \theta) - y^*\|_2^2$
- Solve non-linear minimization problem with gradient based optimizer (Adam)

Types of Machine Learning

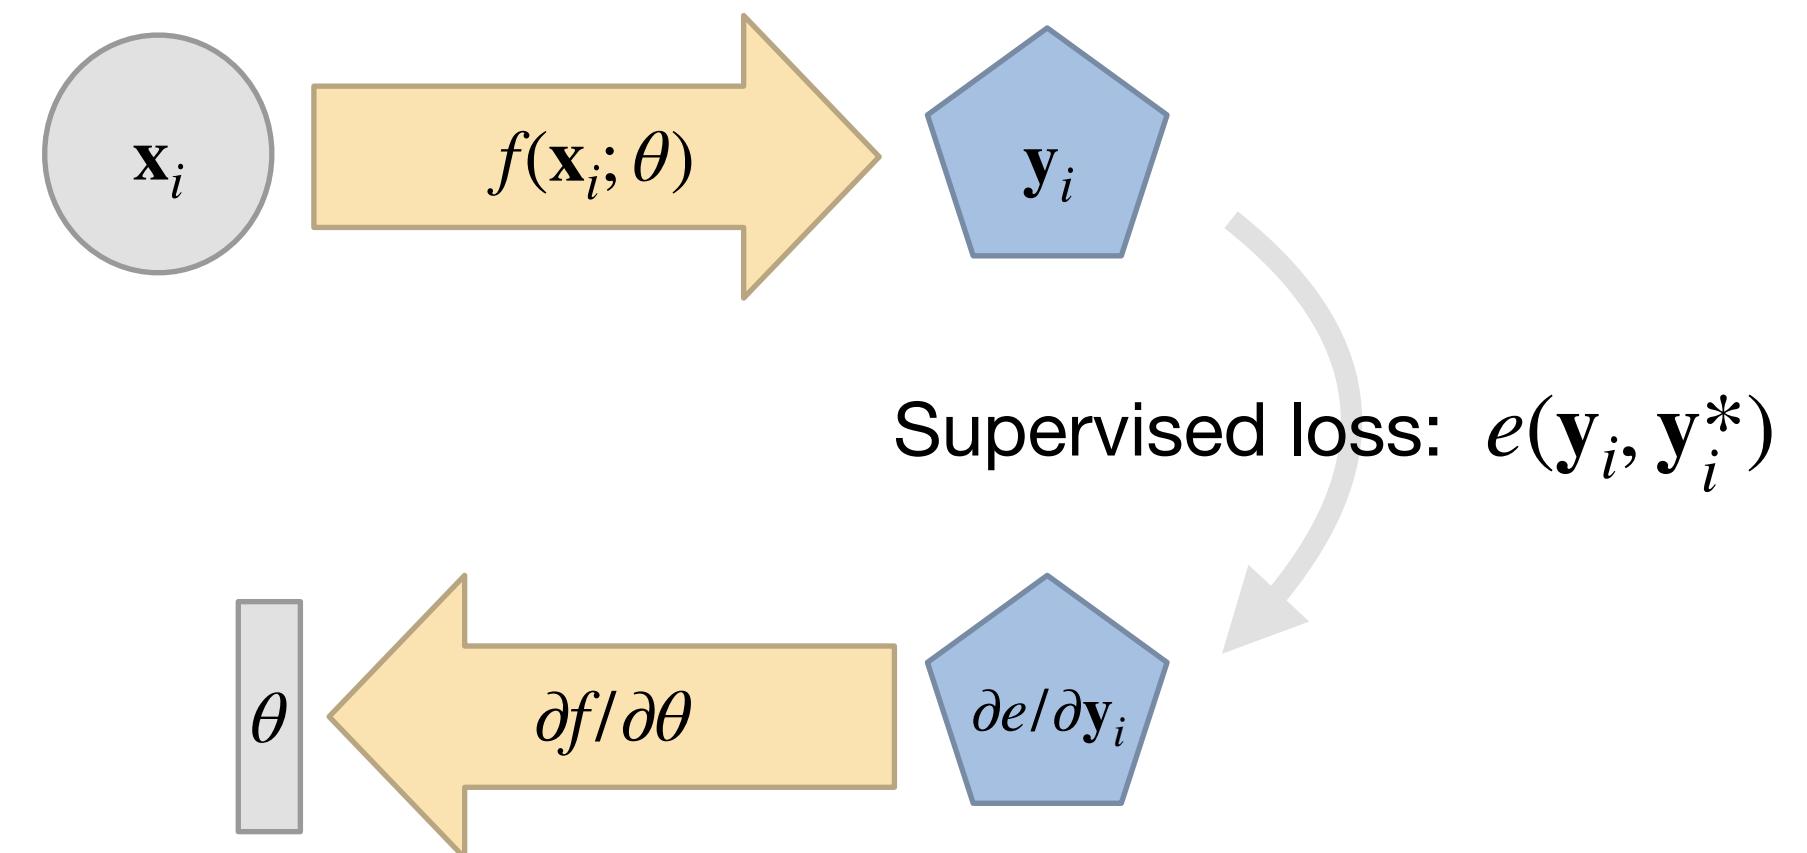
Traditional Viewpoints

- Traditional ML distinction: *classification* VS *regression*
 - In the following: **regression**, $f(x) = y$, with x, y continuous functions
- Later on ***physics regression*** $\mathcal{P}(f(x)) = y$:
physical model \mathcal{P} combined with regression problem;
typically involves highly non-linear functions that cause uneven scaling

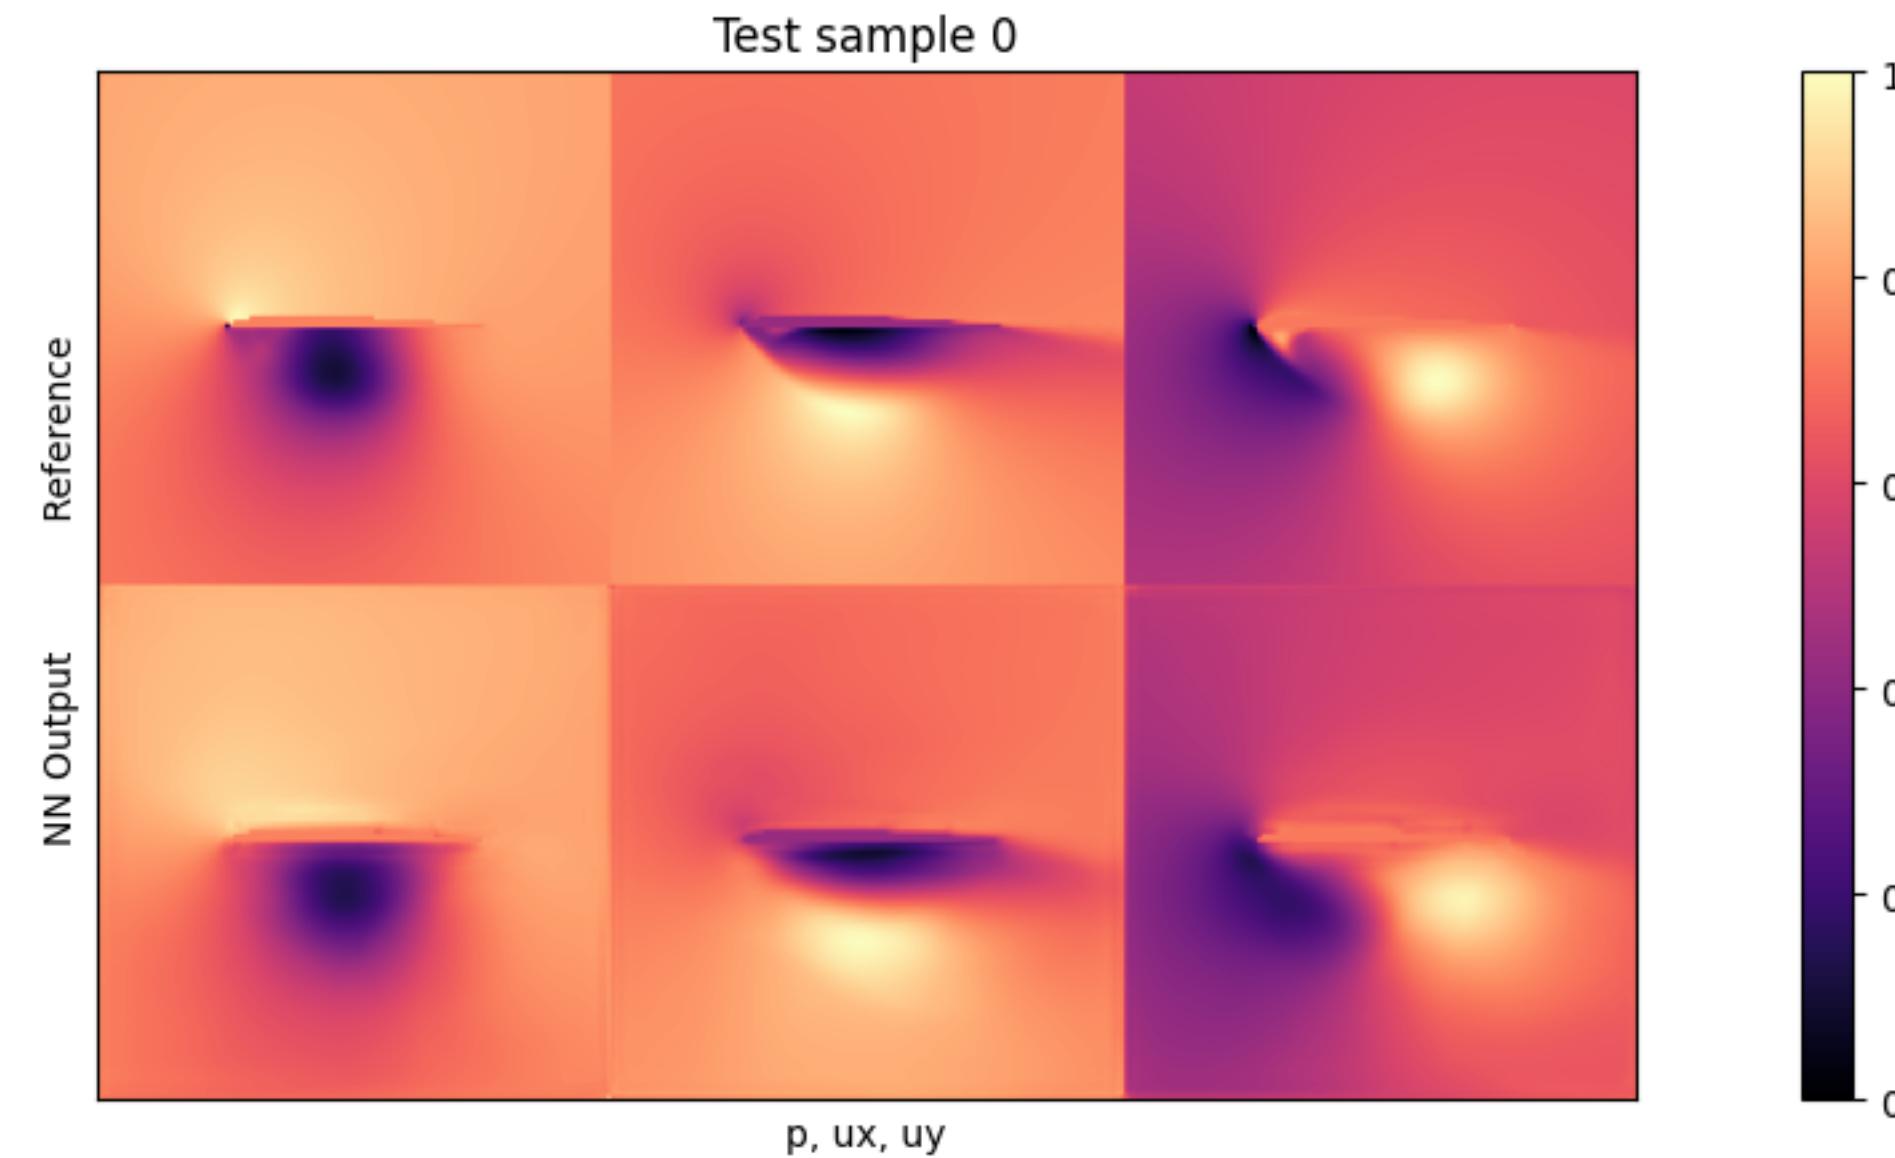


Re-cap Supervised Training

- Definition *Supervised Training* := purely data-driven, pre-computed x, y , with simple loss (e.g. L^2)
- Fully data-driven
 - Physical model not taken into account
 - Sub-optimal accuracy and generalization
- Exactly as before: $\arg \min_{\theta} \|f(x; \theta) - y^*\|_2^2$
- 😍 Beautiful from an ML perspective: no “*inductive biases*” needed
- 😱 Horrible from a computational perspective: no existing knowledge used



Supervised Training



<https://colab.research.google.com/github/tum-pbs/pbdl-book/blob/main/supervised-airfoils.ipynb>

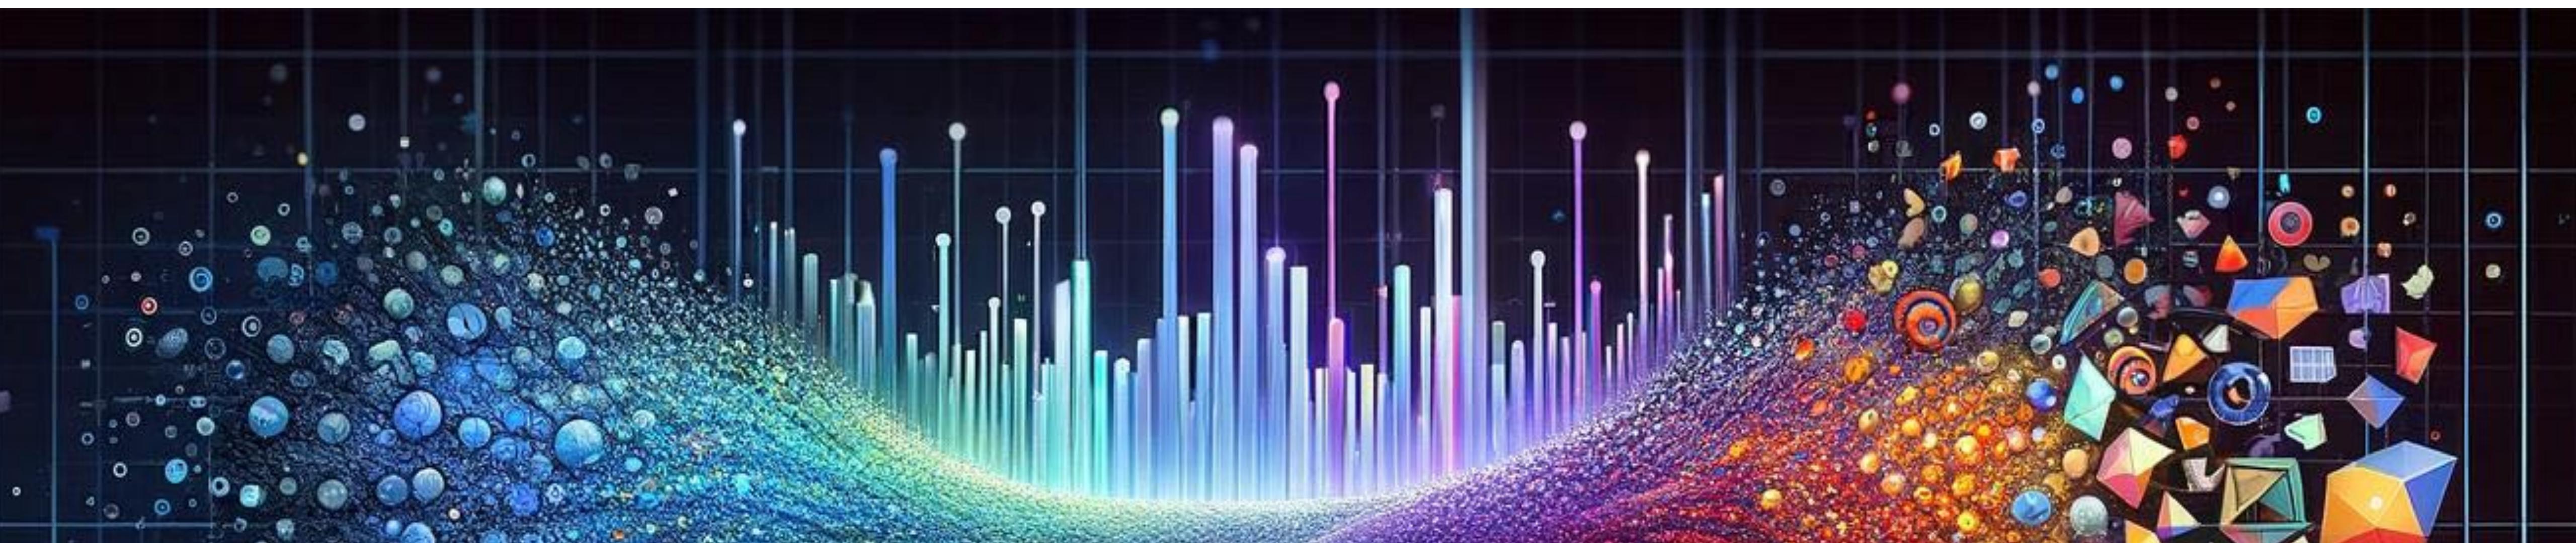
Supervised Training for Time Integration

- Precompute time series data: given states over time $[u^0, u^1, \dots, u^N]$
- Consider batches representing a **single time step forward** $x := u^t; y^* = u^{t+1}$
- Then, just like before: $\arg \min_{\theta} \|f(x; \theta) - y^*\|_2^2$
- Given u^0 approximate any state u^i by **i recurrent / autoregressive evaluations** of $f(\cdot)$

Training Surrogate Models

Recurrent Evaluation

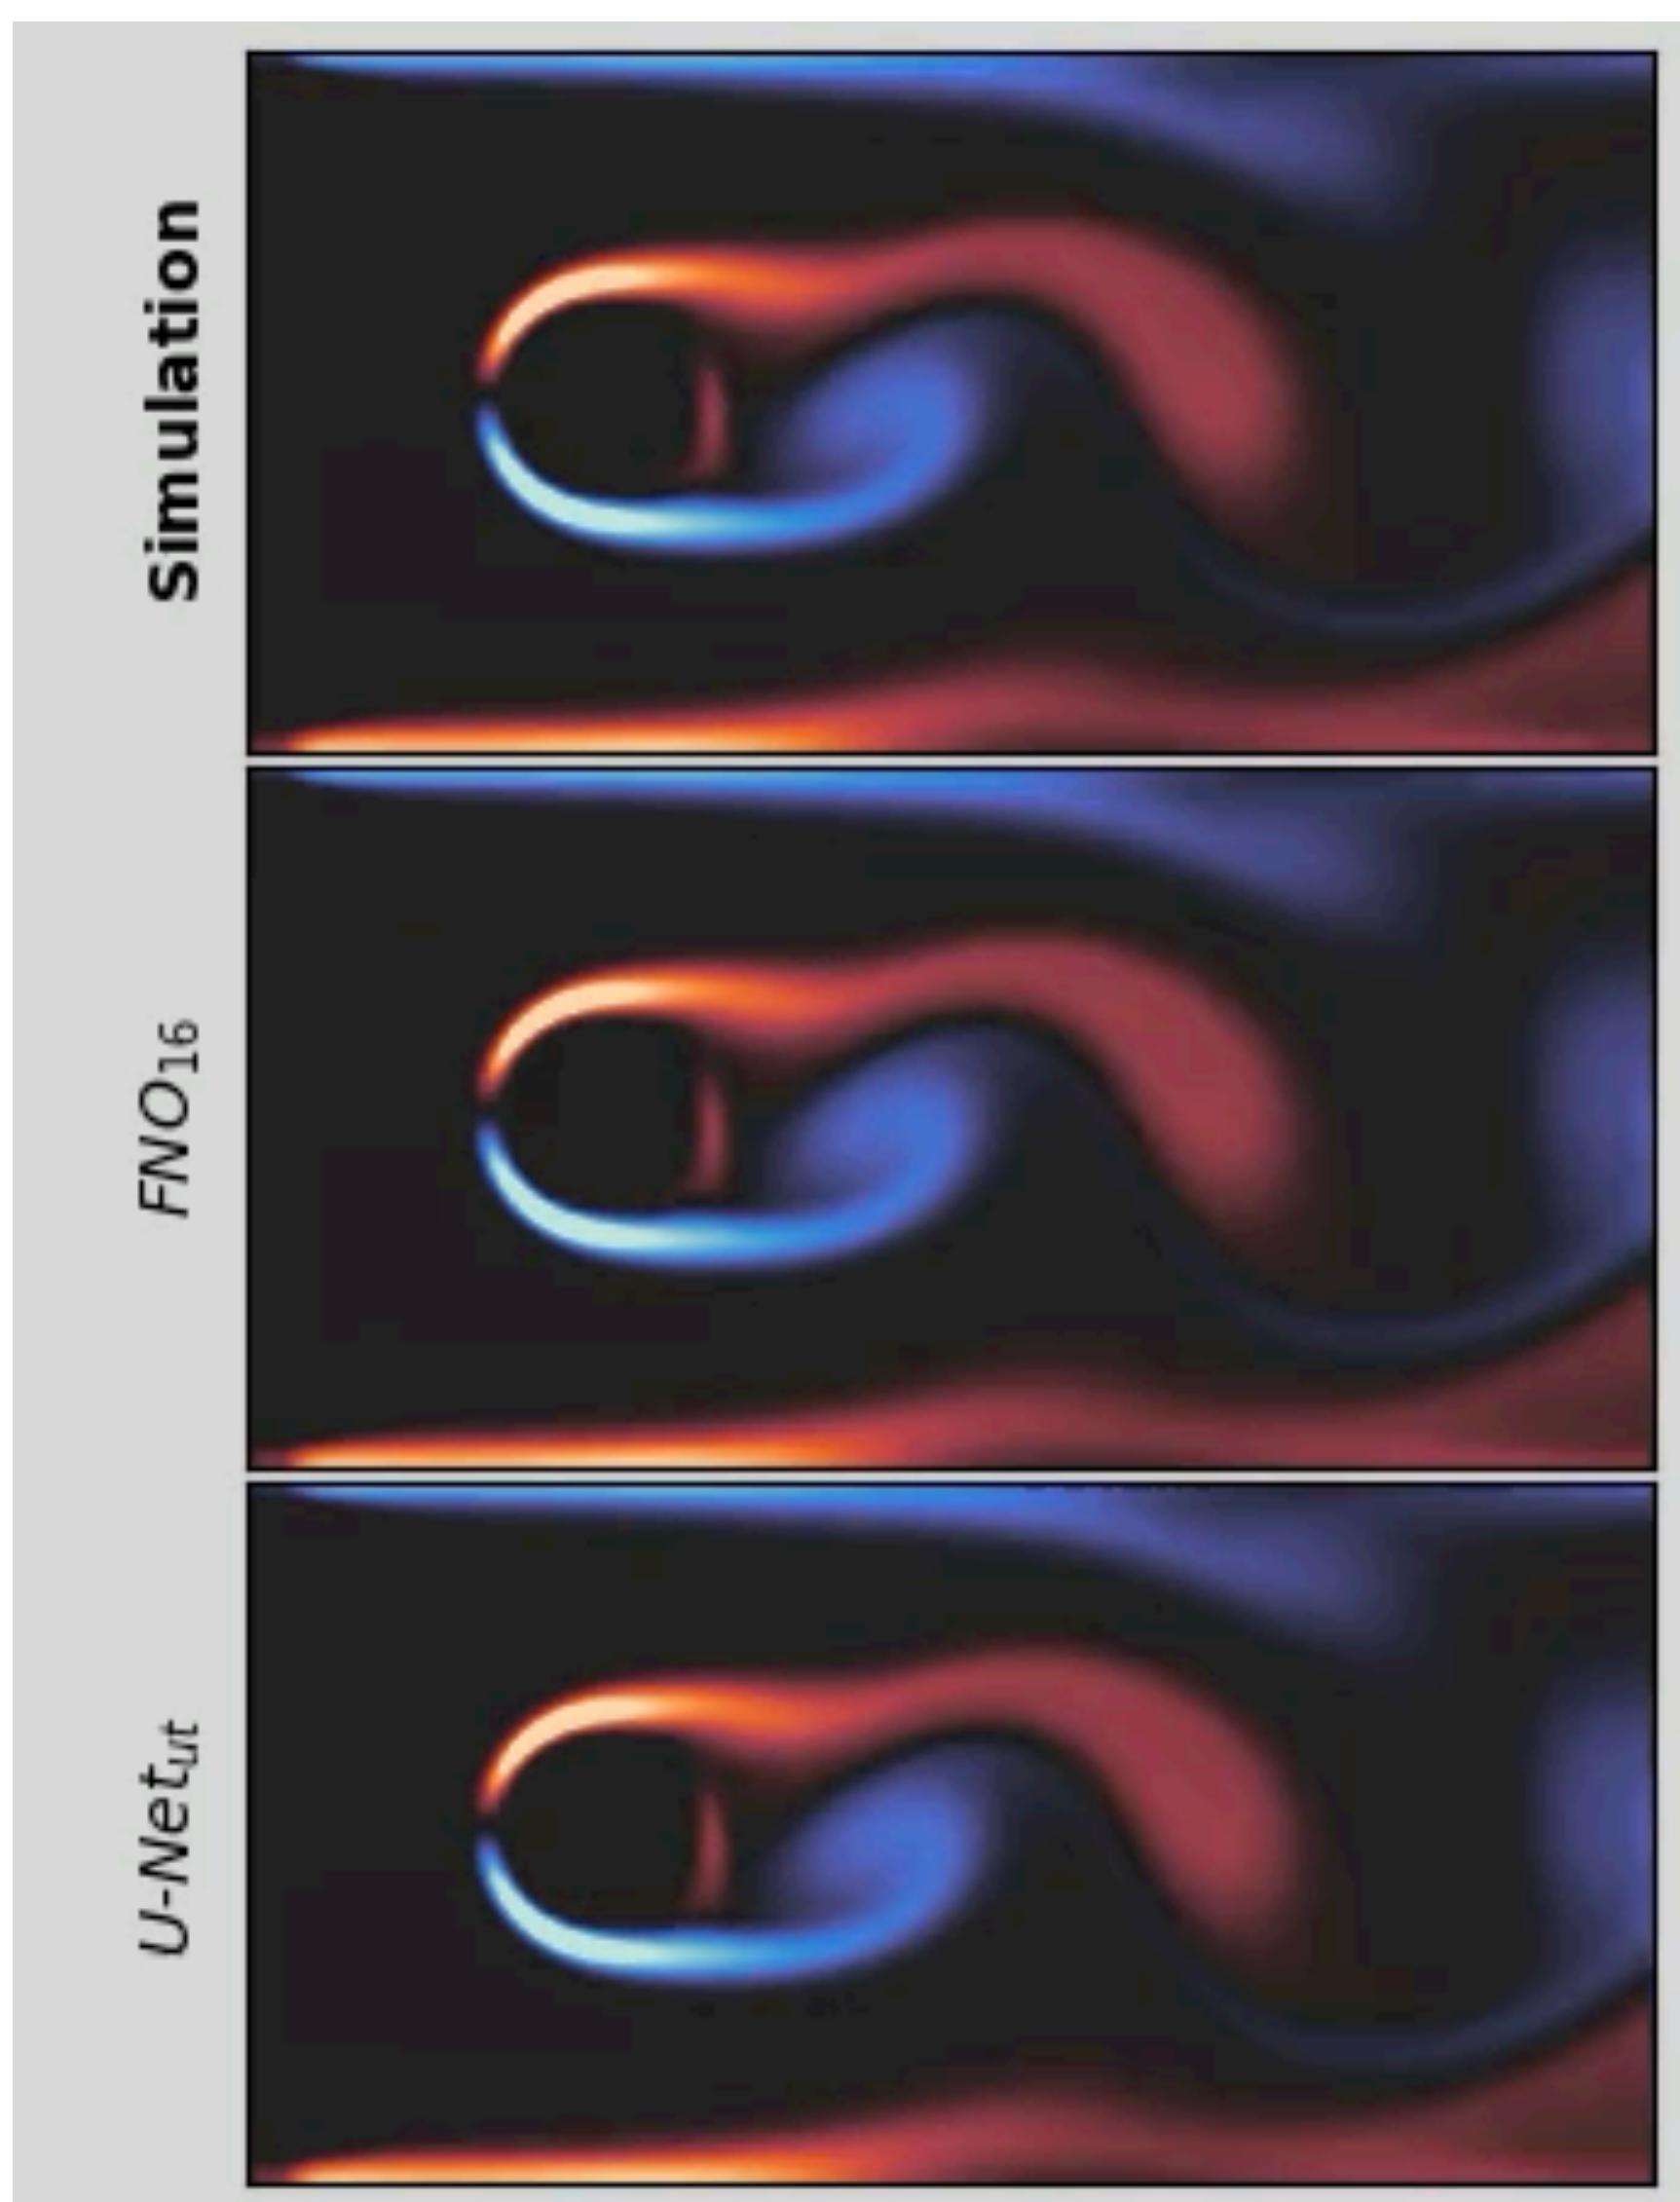
- Per step **approximation errors** will grow; dynamical systems perspective: reference states u^* evolves on *attractor* of PDE \mathcal{P} , with $u_{t+1}^* = \mathcal{P}(u_t^*)$; $u_{t>t_0}^* \in A_{\mathcal{P}}$
- Attractor of NN f doesn't match the one from \mathcal{P} : $u_{t+1} = f(u_t)$; $A_f \neq A_{\mathcal{P}}$
- Classic “**“data shift”** problem from ML , causes instabilities!



Training Surrogate Models

Growing Errors - Example

- Simple Navier Stokes “wake flow”
- Here: all models are quite good
- “Drift” from G.T. is very slow
- (Not shown: eventual complete blow up)



Training Surrogate Models

Outlook

- Obvious fix: include **time evolution** in training to improve attractor, ideally include solver
- Train with **unrolling** , more details later on...



Best Practices

- *Always* start here
- *Always* start with overfitting 1 data point
- *Always* check number of NN parameters
- *Always* adjust hyper parameters at this stage
- ... then slowly introduce more data and beautiful physics models

Supervised Training

Best Practices

- ✓ fast, reliable (builds on established DL methods)
- ✓ Great starting point
- ✗ Sub-optimal performance, accuracy and generalization.
- ✗ Fundamental problems in multi-modal settings
- ✗ Requires precomputed data (data shift problem)



Supervised Learning

ADVANCED DEEP LEARNING FOR PHYSICS

Lectures / Exercises

- Ex1: Phiflow done
- Ex2: First “real” one coming up
- Feedback session again on Monday
- Lecture slides updated just now
- Links to be fixed

Best Practices

- *Always* start here
- *Always* start with overfitting 1 data point
- *Always* check number of NN parameters
- *Always* adjust hyper parameters at this stage
- ... then slowly introduce more data and beautiful physics models



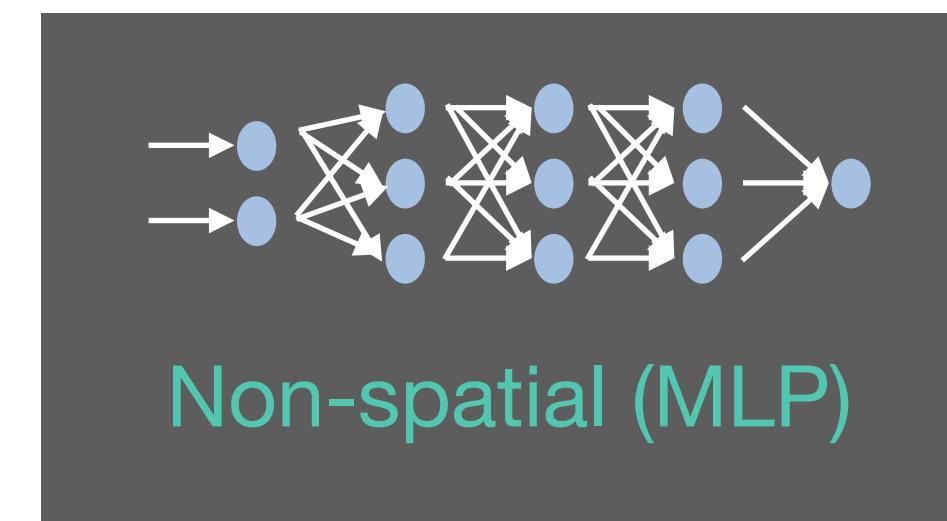
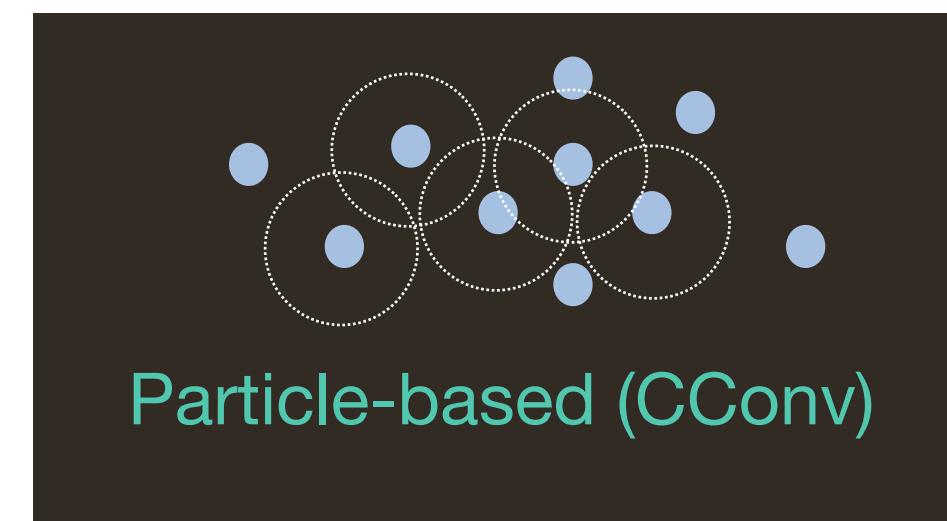
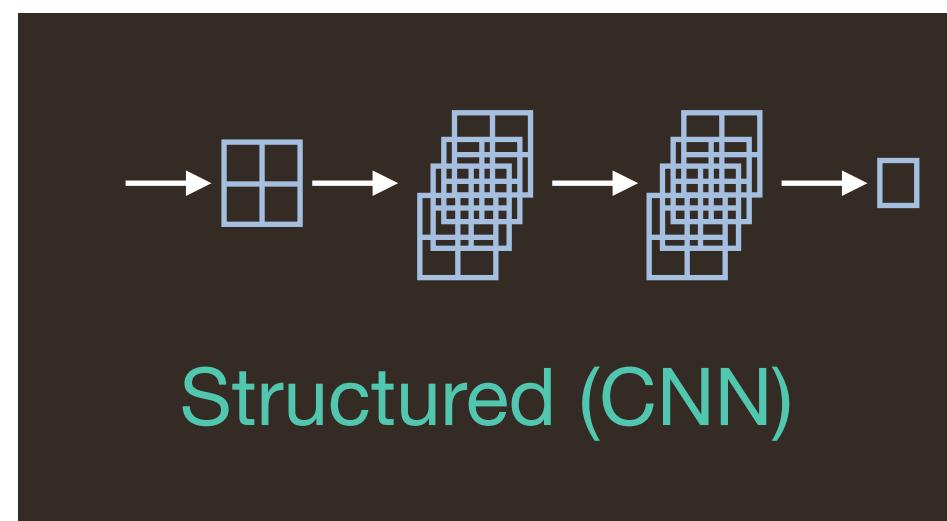
Supervised Learning

ADVANCED DEEP LEARNING FOR PHYSICS

Neural Network Architectures

Overview

Categorization



- Regular spacing on a grid (*structured*)
- Irregular arrangement (*unstructured*)
- Irregular positions without connectivity (*particles*)
- [No spatial arrangement at all]

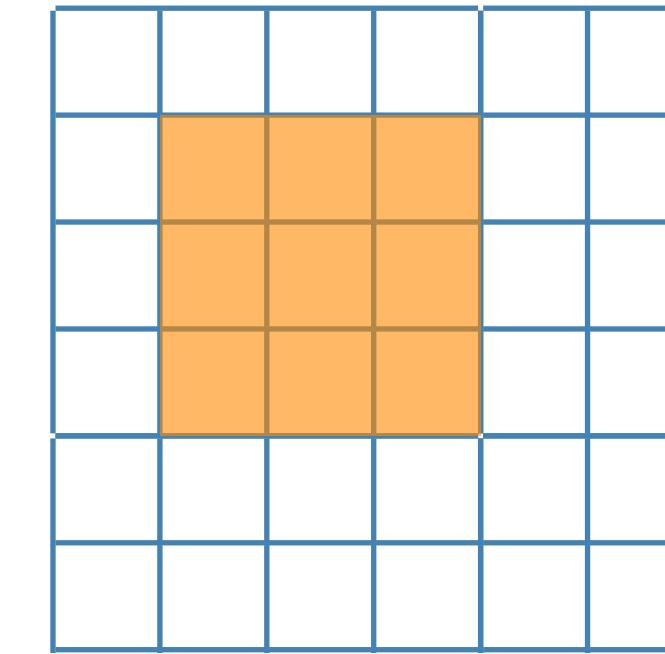
Receptive Fields

- Distinguish Local vs Global interactions
- Similar to hyperbolic PDEs (e.g. waves) and parabolic/elliptic PDEs (e.g. heat)
- No surprise: capabilities of NN should match requirements of PDE... 
- MLPs: trivially global, but scale badly with $\mathcal{O}(N^2)$

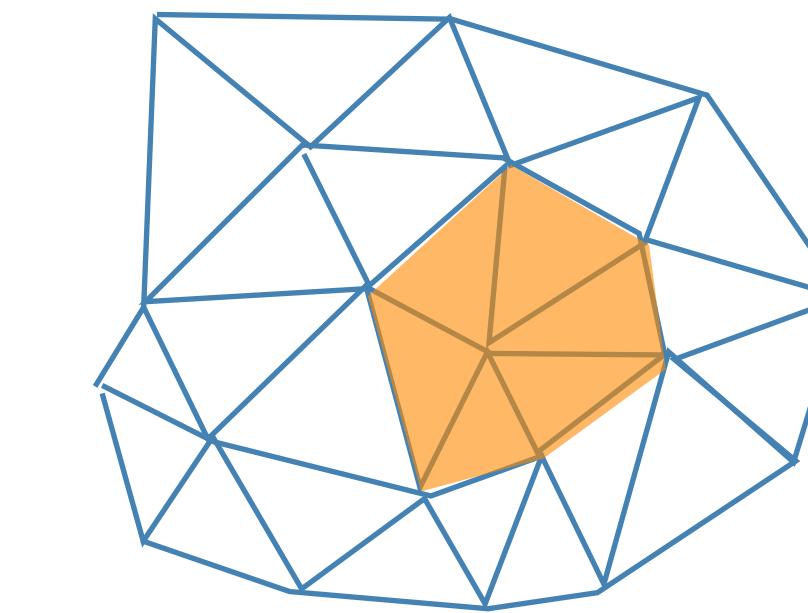
Convolutions & Message Passing

Inherently Local

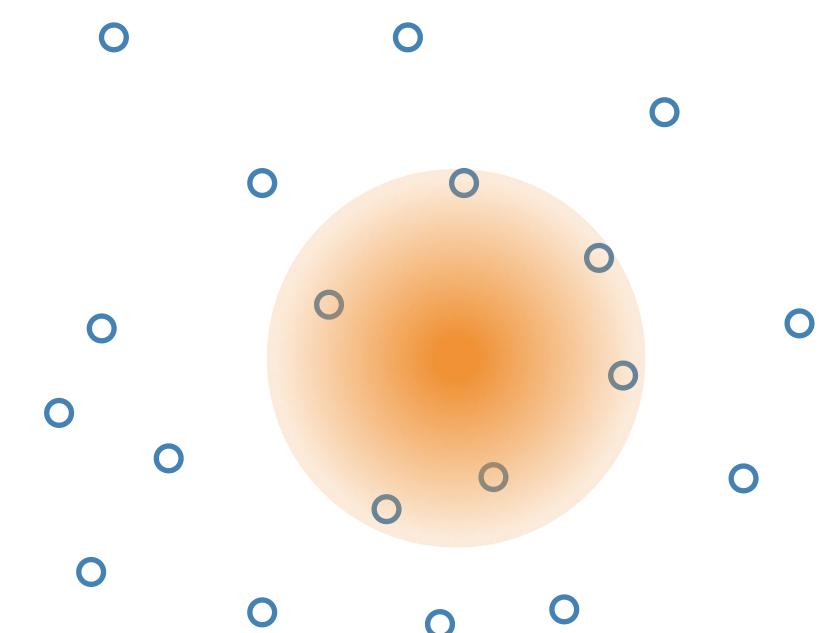
- Grids, graphs, particles
- Stencil operations
- Inductive bias on grids (simplifies implementation) , same concepts on graphs
- Can be non-trivial



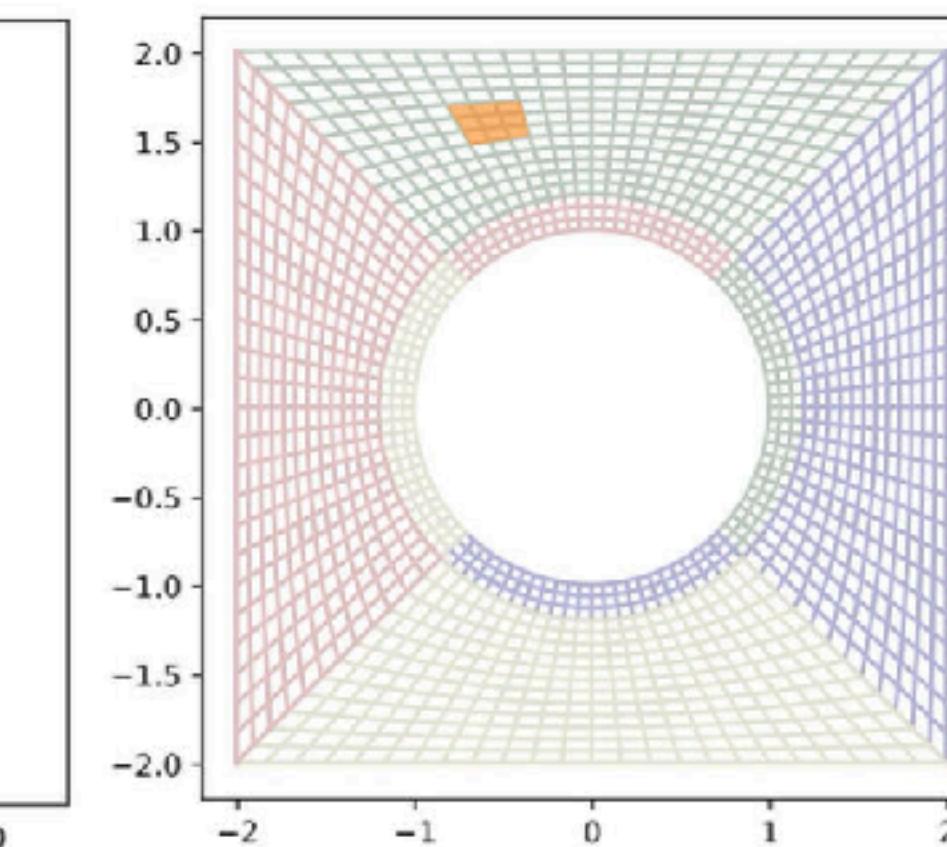
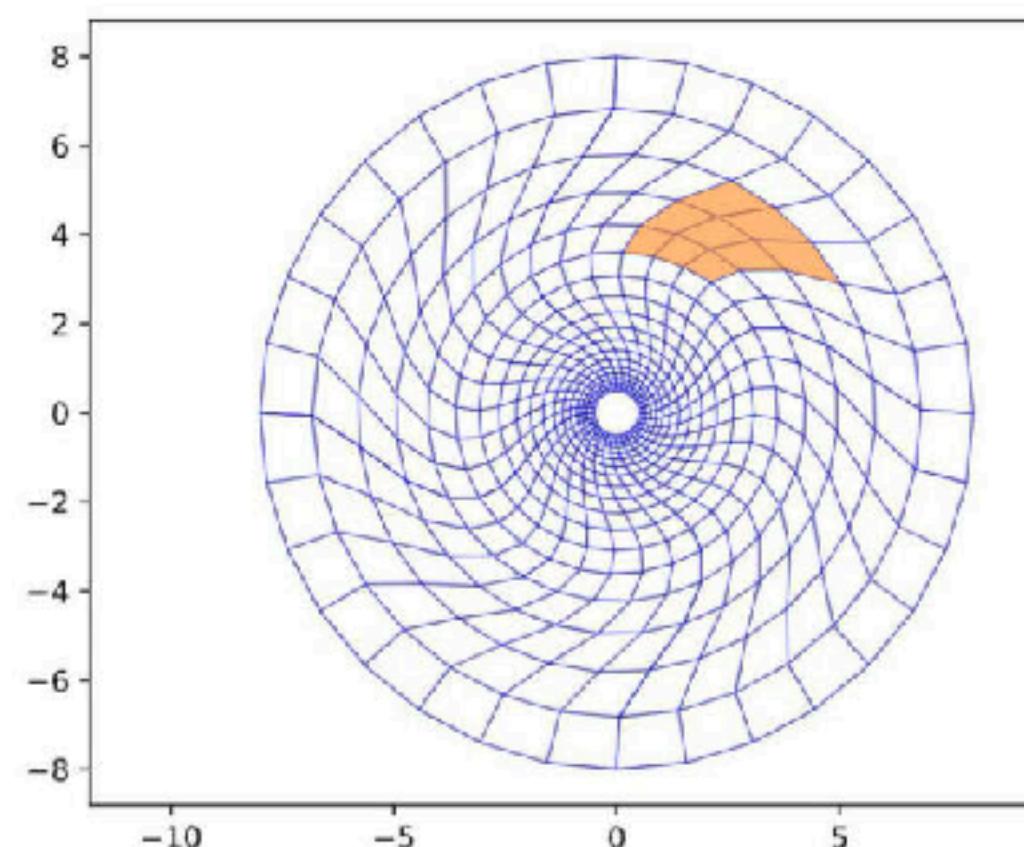
Regular convolution



Message passing /
Graph convolution

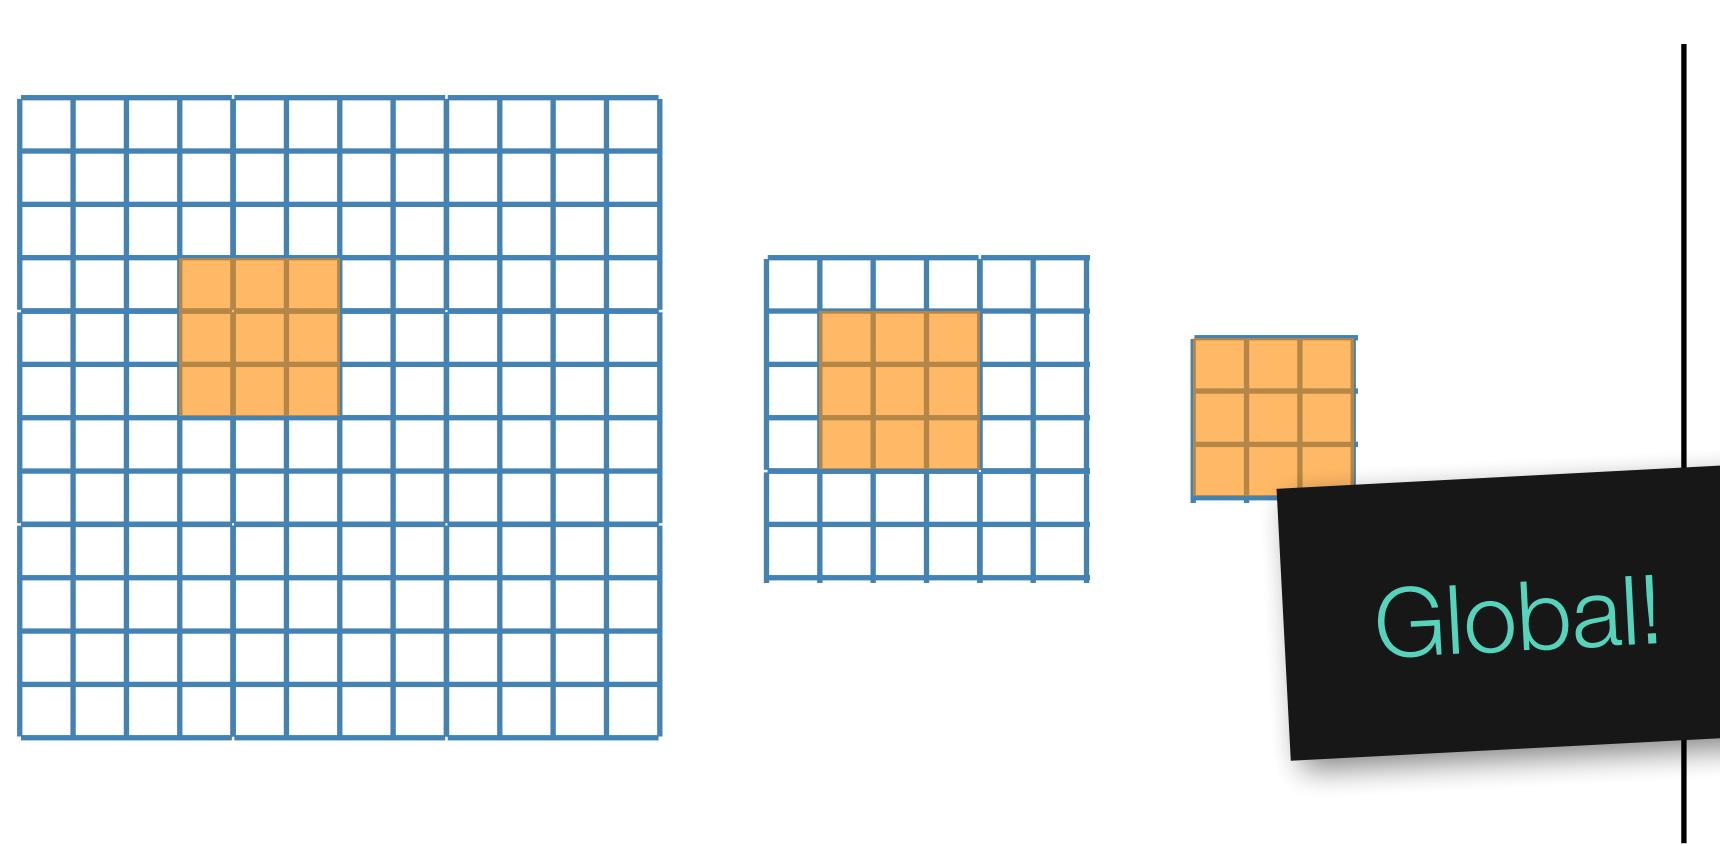


Learned Kernel /
Continuous convolution

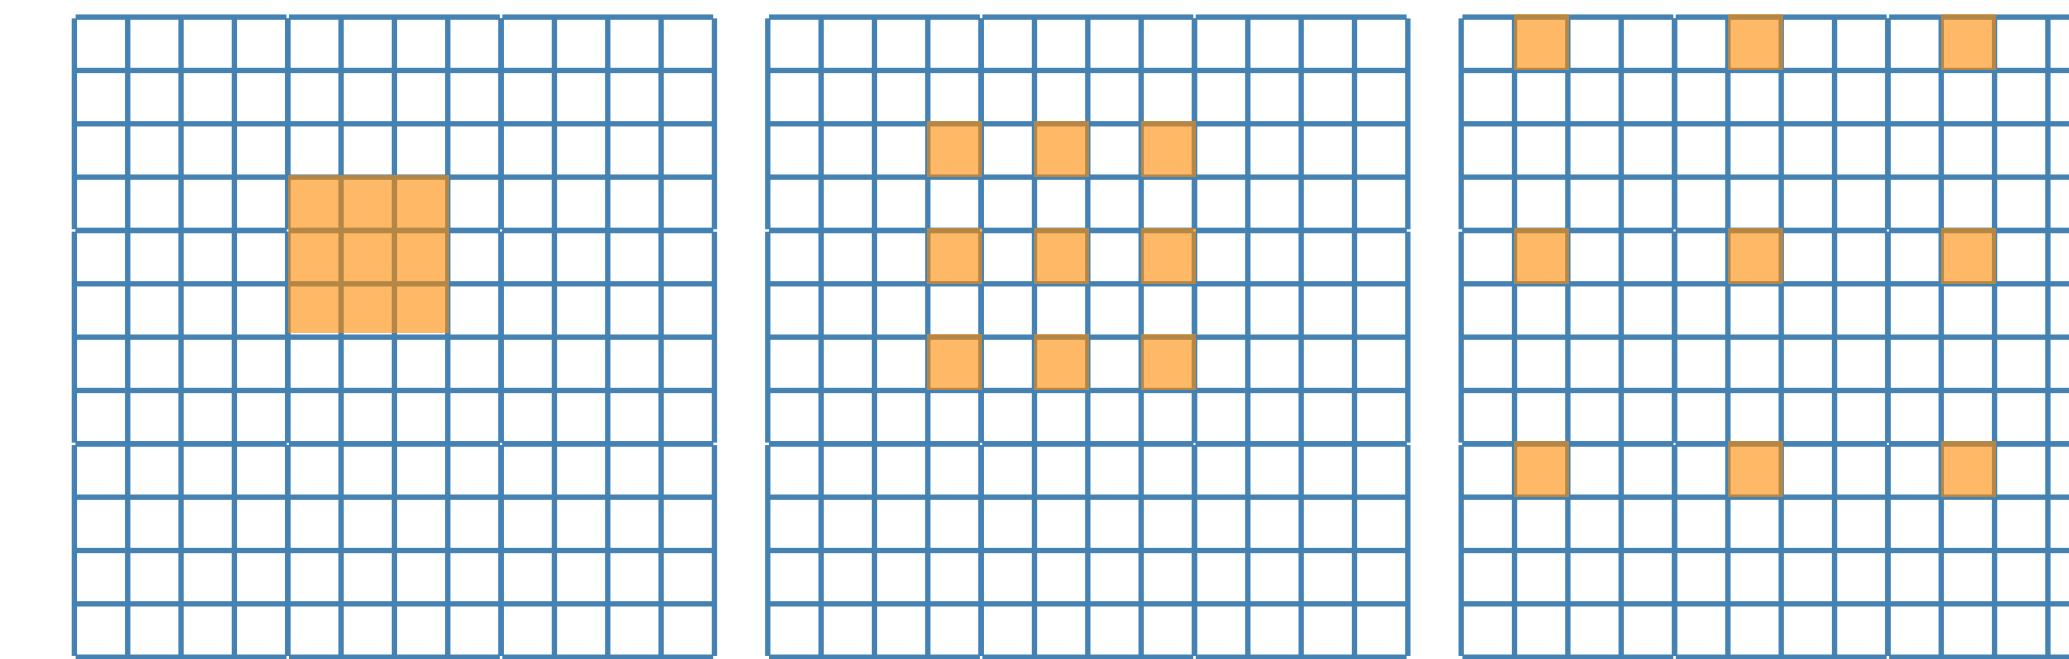


Larger Receptive Fields

- Hierarchy via **Pooling** (similar to *restriction / prolongation* of multigrid)
- Graphs require **clustering** step (ideally as preprocessing)
- Grid-based variant *dilation*: larger receptive field, but less aggregation



Pooling with 3x3 Convolution

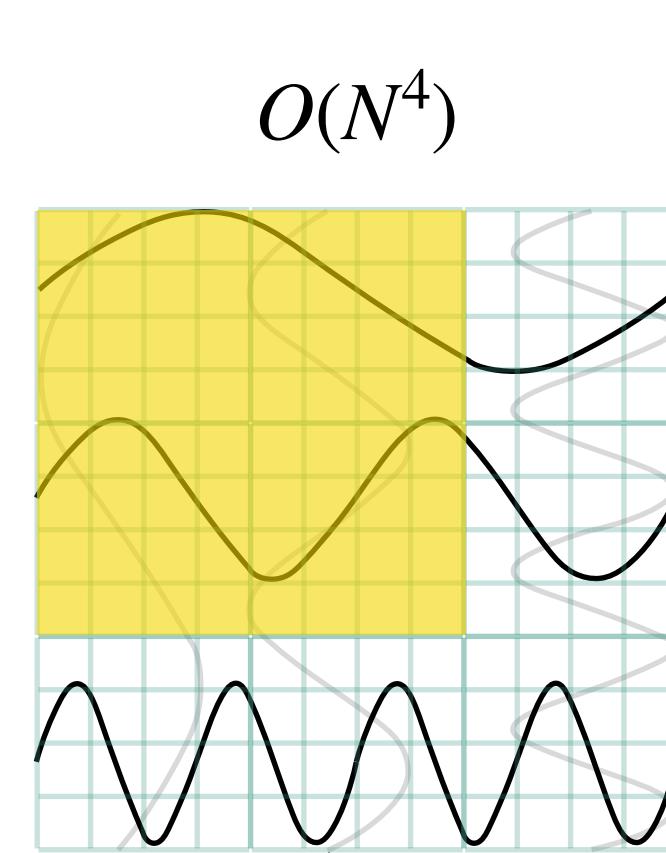


Dilated Convolutions

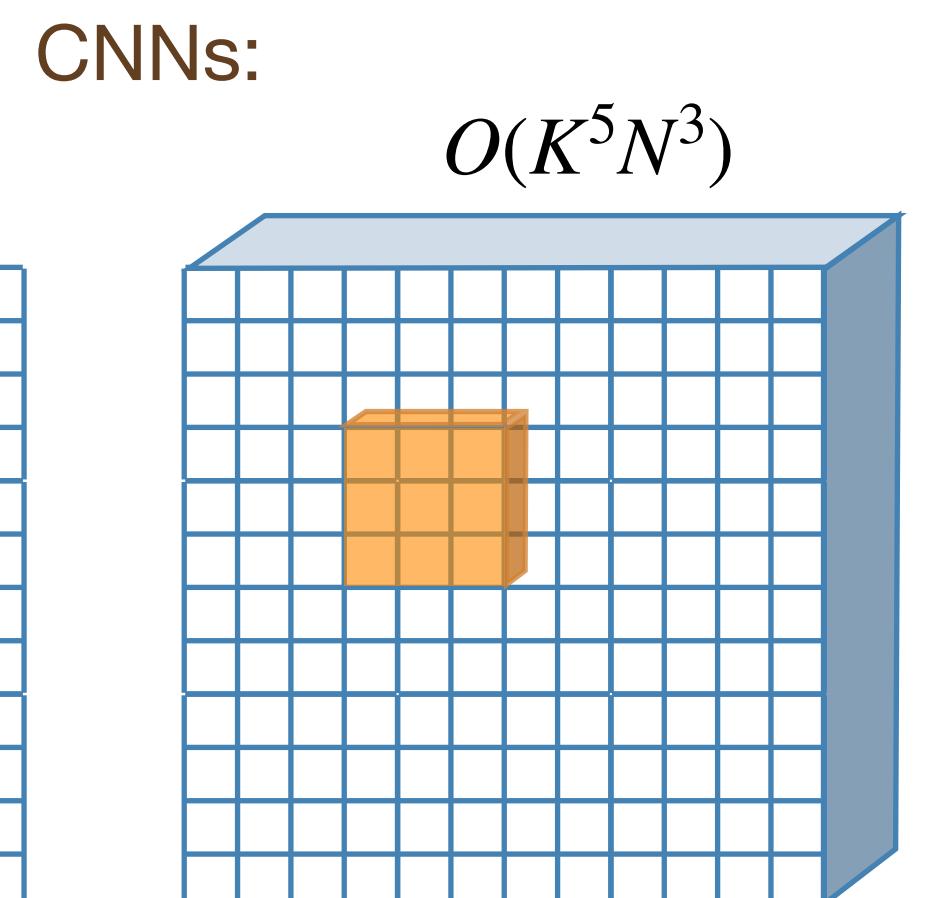
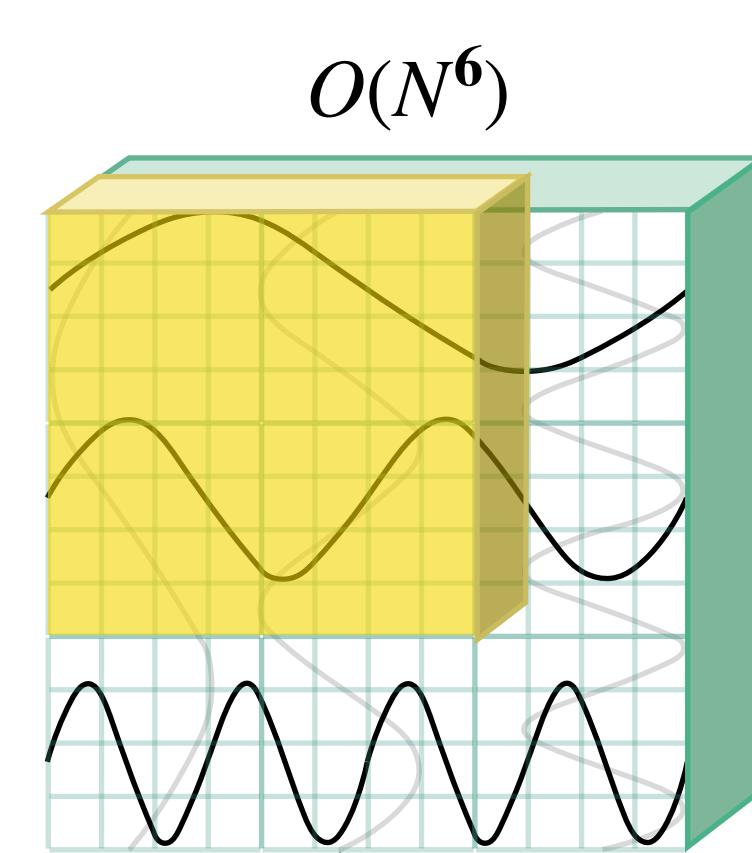
Receptive Fields

Spectral Representations

- Employ *Fourier transforms* (e.g., Fourier Neural Operators)
- Note on **operator perspective**: function transformation, “infinite dimensional”; in practice grids-based, truncated and discrete - not too different from CNN
- But: **global basis functions** from FFT; **suboptimal scaling** for 3D



MLP in Frequency Domain
("1" feature)



CNNs:
 $O(K^5 N^3)$
Convolutions in Spatial Domain
(kernel size K ; assuming feature dimension $O(K)$, not shown)

Receptive Fields

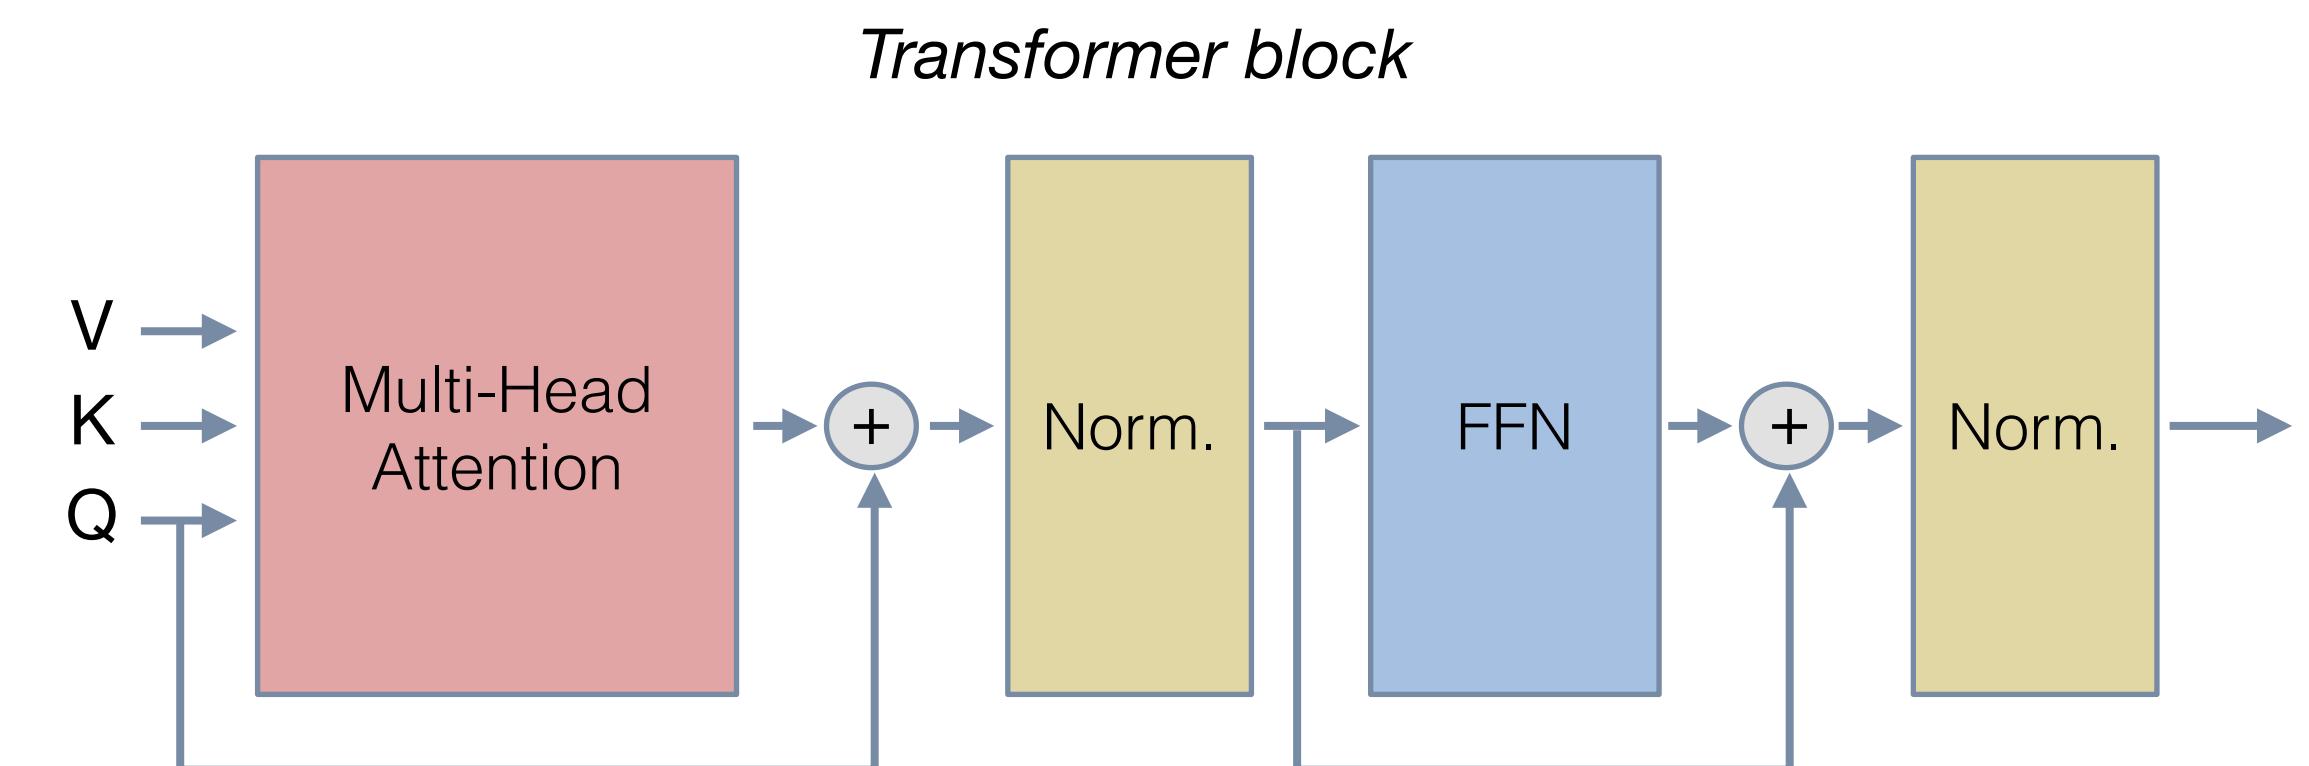
Summary Global vs Local

Details to follow!

	Grid	Unstructured	Points
Local	CNN , ResNet	GNN	CConv
Global			
- Hierarchy	U-Net, Dilation	Multi-scale GNN	Multi-scale CConv
- Spectral	FNO	Spectral GNN	(-)

Transformers

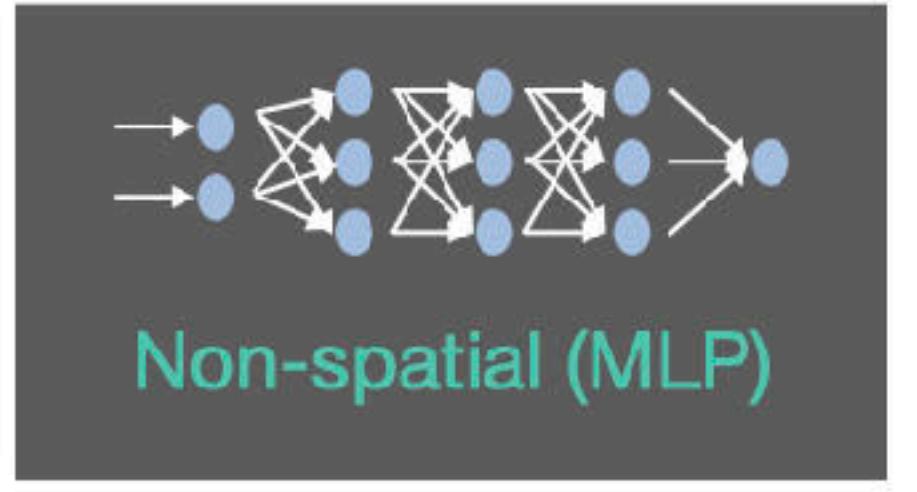
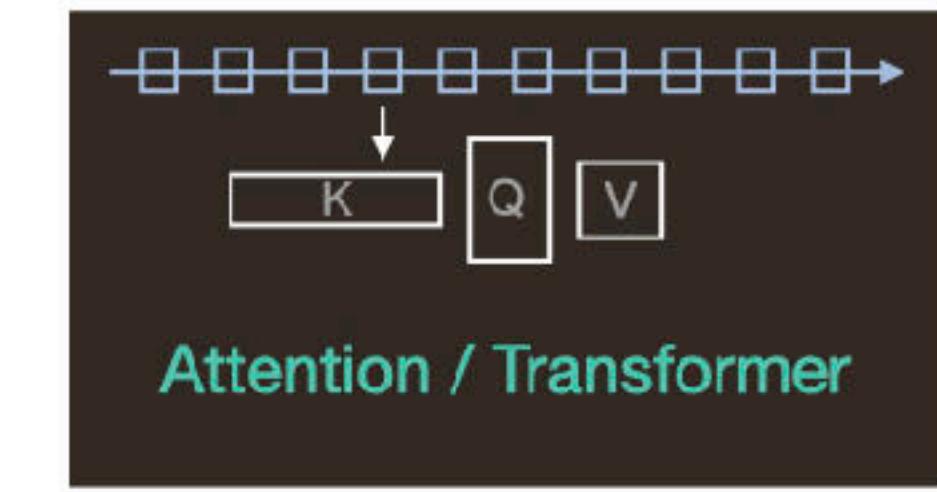
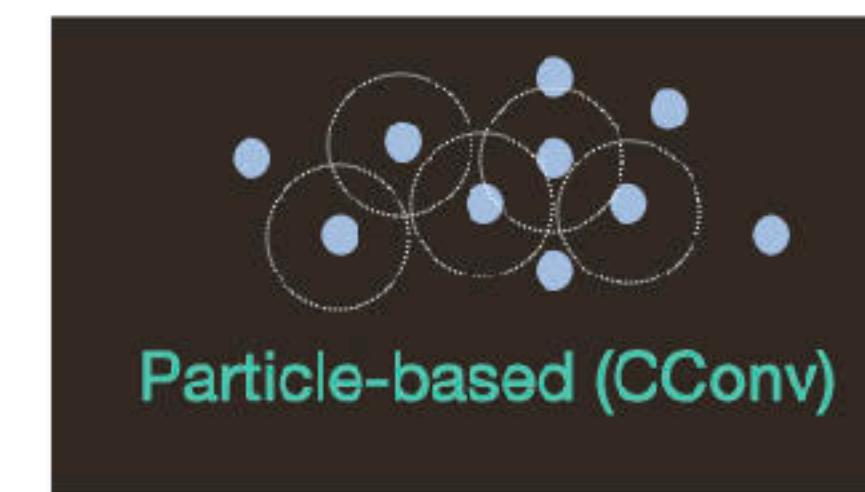
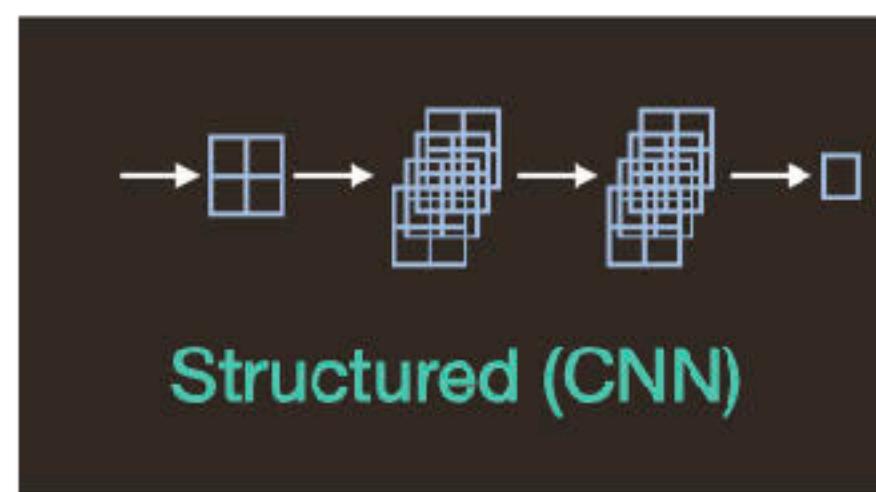
- Sequence-to-sequence architecture, token processing
- *Attention* yields “global” receptive field on token level
- Details out of scope...
- Great scaling (compute) , less ideal for memory (attention is quadratic)
- “The Future” !? 😊🤔



Neural Network Architectures

Summary

- Central consideration: local vs. global
- Grids or graphs: physics-concepts apply in the same way
- Transformers: likewise grids and graphs...





End