

Differentiable Simulations

DEEP LEARNING FROM AND WITH NUMERICAL PDE SOLVERS (PART 1)

Contents

Physical Loss Terms

Differentiable Physics Simulations

- Examples

Differentiable Physics Training

- Examples



Physical Loss Terms

Formulating a “Classic” Numerical Simulator

- Physical model PDE \mathcal{P} with phase space states $\mathbf{u}(\mathbf{p}, t)$, shortened to $\mathbf{u}(t)$
- Reformulate \mathcal{P} to compute new state or solution
- Time derivative $\mathbf{u}_t = \mathcal{F}(\mathbf{u}_x, \mathbf{u}_{xx}, \dots, \mathbf{u}_{xx\dots x})$
- Apply standard time integration; Euler step gives $\mathbf{u}(t + \Delta t) = \mathbf{u}(t) + \Delta t \mathbf{u}_t$
- More complex PDEs like Navier-Stokes $\frac{D\mathbf{u}}{Dt} = -\frac{1}{\rho} \nabla p + \nu \nabla^2 \mathbf{u} + \mathbf{g}$, $\nabla \cdot \mathbf{u} = 0$ require *operator splitting* or custom time integration schemes

Residual Equations

Example setup for time integration (many variations possible)

Solution given by NN $f(\mathbf{x}; \theta) \approx \mathbf{y}^* = \mathbf{u}(t + \Delta t)$ with $\mathbf{x} = \mathbf{u}(t)$, evaluate \mathcal{F} here

$R := \mathbf{u}(t + \Delta t) - \mathbf{u}(t) - \Delta t \mathcal{F}(\mathbf{u}_x, \mathbf{u}_{xx}, \dots, \mathbf{u}_{xx\dots x})$ for \mathcal{F} evaluated at time $(t + \Delta t)$

Minimize: $|f(\dots) - \mathbf{x} - \Delta t \mathcal{F}(\mathbf{u}_x, \mathbf{u}_{xx}, \dots)|^2$

⇒ Does not require pre-computed solutions of \mathcal{P} anymore

Integration into Neural Networks

Add residual as additional loss term


Including residual R in DL loss gives: $\arg \min_{\theta} \sum_i \alpha_0 (f(\mathbf{x}_i; \theta) - y_i^*)^2 + \alpha_1 R(\mathbf{x}_i)$

Relative weighting of terms via α_0, α_1 factors

Backpropagate through expressions to train NN

Physical Models and Residuals in NNs

Visual Outline



⇒ So far generic, 2 variants will follow...

Side Note: Steady-state Problems also fit in ...

- No explicit time dimension t , but typically still involve an *imaginary* time
 - **Iterative solvers** are used to compute converged solutions
 - Solver iterations can be treated as *imaginary / virtual time*
 - $\mathbf{u}(t)$ at initial time $t = 0$ typically zero
- Thus: **special case.** We'll focus on **time dependent problems**

Flexible NNs with Traditional Discretization - Variant 1


Work with **chosen spatial discretization** (grid, mesh, points ...)

E.g., \mathbf{u} represented by dense array in space

Discretize equations of R accordingly, and provide $(\partial R / \partial \theta)^T$

For $\mathbf{u} = f(\mathbf{x}; \theta)$ with $R(\mathbf{u})$: provide $(\partial R / \partial \mathbf{u})^T$, then rely on **backprop**

E.g., compute $\partial \mathbf{u} / \partial x$ via finite difference on grid



E.g.: value stored at
array location [2,3]

Example: Learning Divergence-Freeness


Specific example: Continuity equation $\nabla \cdot \mathbf{u} = 0$

Single time step of the form $\mathbf{u}(1) = \mathbf{u}(0) - \nabla p; p = \nabla^{-2}(\nabla \cdot \mathbf{u}(0))$


Learn to predict pressure field via NN: $p = f(\mathbf{u}(0); \theta)$

Goal is $\nabla \cdot \mathbf{u}(1) = 0$ and hence minimize $\nabla \cdot (\mathbf{u}(0) - \nabla f(\mathbf{u}(0); \theta))$

- Note: has to hold at all times, hence time step of 1
- Use Eulerian grid and CNN, discretize divergence and gradient operators via finite differences for all computational cells



Physical Models and Residuals in NNs (v1)



Classic Paper by Tompson et. al, reimplemented in:


<https://www.physicsbaseddeeplearning.org/physicalloss-div.html>

v1 Discussion

- ✓ Accurate and fast residual calculations, updates whole solutions \mathbf{u}
- ✓ No pre-computations and large data sets needed
- ✓ Good for multi-modal problems
- ✗ Not “adaptive”, discretization needs to be chosen carefully
- ✗ Slower than supervised approaches (evaluation and backprop of R)


Special Case with Fully connected NNs - Variant 2

Fully-connected NN to represent solution and for computing derivatives



“Neural field” representation

Value retrieved
by querying network
 $f(\mathbf{p}, \cdot; \theta) \quad \mathbf{p} = (2, 3)^T$




E.g.: value stored in
neural network
at location [2,3]

Physical Models and Residuals in NNs (v2)

Special Case with Fully connected NNs - Variant 2

Fully-connected NN to represent solution and for computing derivatives




Fully-connected NN

Spatial coordinates \mathbf{p} are a part of \mathbf{x} now, i.e. an input to f ; evaluate f to sample \mathbf{u}


Commonly known as “*physics-informed neural networks*” PINNs

Formulate residual via NN f , e.g., compute $\partial\mathbf{u}/\partial x$ via derivative of NN $\partial f/\partial x$



Physical Models and Residuals in NNs (v2)

Example from Raissi et al. “*Physics-informed neural networks*”



<https://www.physicsbaseddeeplearning.org/physicalloss-code.html>

v2 Discussion

- ✓ Likewise: No pre-computations and large data storages needed
- ✓ Likewise: Good for multi-modal problems
- ✓ Fully-connected NN “learns” discretization
- ✗ Slow derivatives (especially for higher orders)
- ✗ Localized updates, typically require supervised constraints
- ✗ Accuracy relies on current learned state

Physical Residuals v1 vs v2



Summary / Discussion

v1:

- ✓ / ✗ Chosen discretization
- ✓ Reliable and fast evaluation of residual
- ✓ Each learning step provides gradient for whole domain

v2:

- ✓ / ✗ Discretization adapts at training time, influences derivatives
- ✓ Easy to start with
- ✗ Incompatible with most existing numerical techniques
- ✗ Each learning step provides only “local” updates




Physical Residuals v1 vs v2

Summary / Discussion


Compare loss propagation at training time:

v1: ResNet/Unet based - learning signal for full domain, potentially influences full NN state

v2: MLP based - localized sampling points, only partial updates of NN state



Strongly influences convergence



Physical Residuals in General

✓ Uses physical model

✗ Soft constraints, no guarantees

Outlook “Neural Fields”:

- Hamiltonian / Lagrangian NNs , interesting variants of v2 residuals:
<https://arxiv.org/pdf/1906.01563.pdf> and <https://arxiv.org/pdf/2003.04630.pdf>

Next:

How can we use all of the “traditional knowledge” of physical simulations in DL?

Differentiable Physics Simulations

Same Starting Point: “Classic” Numerical Simulator

- Physical model PDE \mathcal{P} with phase space states \mathbf{u}
- \mathcal{P} to compute new state, time derivative $\mathbf{u}_t = \mathcal{F}(\mathbf{u}_x, \mathbf{u}_{xx}, \dots, \mathbf{u}_{xx\dots x})$
- Apply time integrator of choice
- \Rightarrow Augment \mathcal{P} with NN, train via simulation derivatives $(\partial \mathcal{P} / \partial \mathbf{u})^T$
- “Simply” broader picture than before: full simulator rather than (partial) residual
- ... before involving NNs, let’s look at how to compute and work with $(\partial \mathcal{P} / \partial \mathbf{u})^T$

Nomenclature

- Note: Naming schemes not (yet) unified
- Physics- constrained / based / augmented ...
- Closely related to “classical” *Adjoint methods*
- Equivalent: Reverse mode differentiation (== backpropagation)
- Distinguish physical simulations from financial, crowd, light simulations etc., hence “*differentiable physics*” (DP)

Jacobian Vector Products

- Solver as **sequence of m operations**, e.g.: $\mathbf{u}(t + \Delta t) = \mathcal{P}_m \circ \dots \mathcal{P}_2 \circ \mathcal{P}_1(\mathbf{u}(t), \nu)$, with parameter ν (e.g., viscosity)
- Jacobian for \mathbf{u} of solver operation i :
$$\frac{\partial \mathcal{P}_i}{\partial \mathbf{u}} = \begin{bmatrix} \frac{\partial \mathcal{P}_{i,1}}{\partial u_1} & \dots & \frac{\partial \mathcal{P}_{i,1}}{\partial u_d} \\ \vdots & & \vdots \\ \frac{\partial \mathcal{P}_{i,d}}{\partial u_1} & \dots & \frac{\partial \mathcal{P}_{i,d}}{\partial u_d} \end{bmatrix}$$
 (square in this case)
- We could compute $\frac{\partial \mathcal{P}_i}{\partial \nu}$, but Jacobians only needed for **optimized quantities**; here that's \mathbf{u} (later on θ)
- Evaluate chain rule, e.g., for two operators:
$$\left. \frac{\partial (\mathcal{P}_1 \circ \mathcal{P}_2)}{\partial \mathbf{u}} \right|_{\mathbf{u}^n} = \left. \frac{\partial \mathcal{P}_1}{\partial \mathbf{u}} \right|_{\mathcal{P}_2(\mathbf{u}^n)} \left. \frac{\partial \mathcal{P}_2}{\partial \mathbf{u}} \right|_{\mathbf{u}^n} \quad f(g(x)' = f'(g(x)) g'(x)$$
- Always **scalar loss as final operation** with 1-column $\left(\frac{\partial L}{\partial \mathbf{u}} \right)^T$, i.e. gradient vector
- Leads to sequence of **Jacobian-vector products**:
$$\left(\frac{\partial \mathcal{P}_1}{\partial \mathbf{u}(t)} \right)^T \left(\frac{\partial \mathcal{P}_2}{\partial \mathcal{P}_1} \right)^T \left(\frac{\partial L}{\partial \mathcal{P}_2} \right)^T$$

Jacobian Vector Products

- Solver as sequence of m operations, e.g.: $\mathbf{u}(t + \Delta t) = \mathcal{P}_m \circ \dots \mathcal{P}_2 \circ \mathcal{P}_1(\mathbf{u}(t), \nu)$, with parameter ν (e.g., viscosity)
- In this notation: \mathcal{P}_1 applied first, \mathcal{P}_m is last operation
- Solver derivative $\left(\frac{\partial \mathcal{P}}{\partial \mathbf{u}(t)} \right)^T$ Leads to sequence of m Jacobians:

$$\left(\frac{\partial \mathcal{P}_1}{\partial \mathbf{u}(t)} \right)^T \left(\frac{\partial \mathcal{P}_2}{\partial \mathcal{P}_1} \right)^T \dots \left(\frac{\partial \mathcal{P}_{m-1}}{\partial \mathcal{P}_{m-2}} \right)^T \left(\frac{\partial \mathcal{P}_m}{\partial \mathcal{P}_{m-1}} \right)^T$$

Differentiable Physics Simulations: Fluids as an Example

Re-cap Fluid Simulations

Traditional Numerical Solver

Many advanced techniques available

Two main steps via operator splitting:

- Advection - transport $D\mathbf{u}/Dt = 0$
- Pressure projection - enforce divergence-freeness $\nabla \cdot u = 0$
- Viscosity would require additional implicit solve, analogous to pressure solve

Navier-Stokes equations

$$\frac{D\mathbf{u}}{Dt} = -\frac{1}{\rho} \nabla p + \nu \nabla \cdot \nabla \mathbf{u}, \quad \nabla \cdot \mathbf{u} = 0$$

Example Problem with Explicit Update Step

Example: advection of passive scalar density $d(\mathbf{x}, t)$ in velocity field \mathbf{u}

Physical model \mathcal{P} performs time update $d(t + \Delta t) = \mathcal{P}(d(t), \mathbf{u}, t + \Delta t)$ via advection equation

$$\frac{\partial d}{\partial t} + \mathbf{u} \cdot \nabla d = 0$$

Discretize the advection step \mathcal{P} :

- For simplicity, with a simple upwinding scheme
- Semi-Lagrangian or higher order schemes analogous


Advection

Example Problem with Explicit Update Step

Find velocity such that density matches given target d^{target} at t^e , with $d(t^0) = d^0$



$$d(t^0) = d^0$$



$$d(t^e) = d^{target}$$

Advection

Example Problem with Explicit Update Step

Find velocity such that density matches given target d^{target} at t^e , with $d(t^0) = d^0$

Loss computed after one advection step: $\arg \min_{\mathbf{u}} |\mathcal{P}(d^0, \mathbf{u}, t^e) - d^{target}|^2$

Gradient for velocity $\Delta \mathbf{u} = \frac{\partial \mathcal{P}^T}{\partial \mathbf{u}} \frac{\partial L^T}{\partial \mathcal{P}} = \frac{\partial d^T}{\partial \mathbf{u}} \frac{\partial L^T}{\partial d}$

Loss Jacobian is simple enough: $\frac{\partial L}{\partial d} = \frac{\partial |\mathcal{P}(d^0, \mathbf{u}, t^e) - d^{target}|^2}{\partial d} = 2(d(t^e) - d^{target})$

Jacobian of advection $\frac{\partial \mathcal{P}}{\partial \mathbf{u}}$ from discretized advection operator

Advection

Example Problem with Explicit Update Step

First-order upwinding scheme:

$$d_i(t + \Delta t) = d_i - \Delta t [u_i^+(d_{i+1} - d_i) + u_i^-(d_i - d_{i-1})] \text{ with}$$

$$u_i^+ = \min(u_i/\Delta x, 0)$$

$$u_i^- = \max(u_i/\Delta x, 0)$$

I.e., either forward or backward finite difference:




$$\text{Negative velocity at } i \text{ gives } \mathcal{P}(d_i(t), \mathbf{u}(t), t + \Delta t) = (1 + \frac{u_i \Delta t}{\Delta x}) d_i - \frac{u_i \Delta t}{\Delta x} d_{i+1}$$

$$\text{Velocity derivative is simply } \frac{\partial \mathcal{P}}{\partial u_i} = \frac{\Delta t}{\Delta x} d_i - \frac{\Delta t}{\Delta x} d_{i+1}$$

Advection

Multiple Simulation Steps

Find velocity so that density matches given target d^{target} at t^e , with $d(t^0) = d^0$ after applying velocity multiple times



Multiple Simulation Steps


Simplified setup:

- Try to match target every step along the way, evaluate loss multiple times
- Keep single constant velocity and density target

Loss computed after multiple advection steps: $\arg \min_{\mathbf{u}} \sum_s |\mathcal{P}(d^0, \mathbf{u}, t^s) - d^{target}|^2$

Gives s loss terms each of which change \mathbf{u}


Multiple Simulation Steps - Schematic



Note: still no NN involved,
but \mathbf{u} used in each P !

Advection over Time

Term of velocity gradient only for s 'th step (*transpose of Jacobians omitted for brevity here!*)

$$\begin{aligned}\Delta \mathbf{u}_s &= \frac{\partial d(t^s)}{\partial \mathbf{u}} \frac{\partial L}{\partial d(t^s)} \\ &+ \frac{\partial d(t^s - \Delta t)}{\partial \mathbf{u}} \frac{\partial d(t^s)}{\partial d(t^s - \Delta t)} \frac{\partial L}{\partial d(t^s)} \\ &+ \dots \\ &+ \left(\frac{\partial d(t^0)}{\partial \mathbf{u}} \dots \frac{\partial d(t^s - \Delta t)}{\partial d(t^s - 2\Delta t)} \frac{\partial d(t^s)}{\partial d(t^s - \Delta t)} \frac{\partial L}{\partial d(t^s)} \right)\end{aligned}$$


Full gradient $\Delta \mathbf{u} = \sum_s \Delta \mathbf{u}_s$, e.g. for $s = 3$ we have $1 + 2 + 3 = 6$ terms...

⇒ Single gradient consists of $O(s^2)$ updates of \mathbf{u}

In line with NN training: treat each evaluation of \mathcal{P} separately to compute Jacobian-vector product

Re-cap Fluid Simulations



Traditional Numerical Solver

Two main steps via operator splitting:

- Advection - *done*
- Pressure projection - *next*
- Viscosity (either implicit / explicit, more of the same...)

Navier-Stokes equations

$$\frac{D\mathbf{u}}{Dt} = -\frac{1}{\rho} \nabla p + \nu \nabla \cdot \nabla \mathbf{u}, \quad \nabla \cdot \mathbf{u} = 0$$

Divergence Freeness Revisited

Example Problem with Implicit Update Step

Implicit step typically requires solving a linear system $\mathbf{u}^n = \mathbf{u} - \nabla p$

E.g., Poisson's equation for pressure $p = (\nabla^2)^{-1}b = A^{-1}b$

For fluids, right hand side is given by divergence: $b = \nabla \cdot \mathbf{u}$

Gradient for velocity, i.e. right hand side: $\frac{\partial p}{\partial b} \frac{\partial L}{\partial p} = A^{-1} \frac{\partial L}{\partial p}$, with $p = A^{-1}b$

Solver for the forward step can be re-used directly (applied to $\partial L / \partial p$)

Performance Considerations

Theoretically: Full iterative solver (e.g., Conjugate Gradient) could be implemented via **matrix-vector ops** of DL framework

Every op is tracked and intermediate state **registered and stored** (for backprop)

... slow and inefficient.

Rather: Use **efficient numerical solver** components for custom ops

(E.g., matrix inversion via fast CG solve in previous example.)

Note: We're essentially computing an *implicit derivative* here (cf. [Implicit Function Theorem](#))

Differentiable Physics Example



(Compare to Burgers Simulation via PINN)

<https://www.physicsbaseddeeplearning.org/diffphys-code-burgers.html>



End