

Simulation-based Inference and Diffusion Models

PROBABILISTIC INVERSE PROBLEMS, PART 2/2

Overview

- Normalizing Flows
- Score Matching
- Langevin Dynamics
- Denoising Diffusion
- Flow Matching
- Physics Constraints

Flow Matching & Rectified Flows

Flow Matching

- Revisit **score matching**: integrate flow (ODE) to transform probability densities
- But: Compute reference trajectories via **denoising framework** (instead of score)
- Main formulation: transform x from p_0 into a sample from p_1 by training integrating a simple ODE: $dx/dt = v_\theta(x, t)$
- Train v_θ with $\mathcal{L}_{\text{CFM}}(\phi) = \mathbb{E}_{q(z,t), p_t(x|z)} \| v_\phi(x, t) - u_t(x | z) \|^2$ with conditional likelihoods over the helper latent variable z

Mappings

- Which mapping from p_0 to p_1 to aim for?
- Straight, linear paths are ideal: correspond to optimal transport
- Slight complication: minimal noise amount σ_{\min} needed to ensure a continuous distribution (instead of discrete samples)

$$p_t(x|x_1) = \mathcal{N}(tx_1, (1 - (1 - \sigma_{\min})t)I)$$

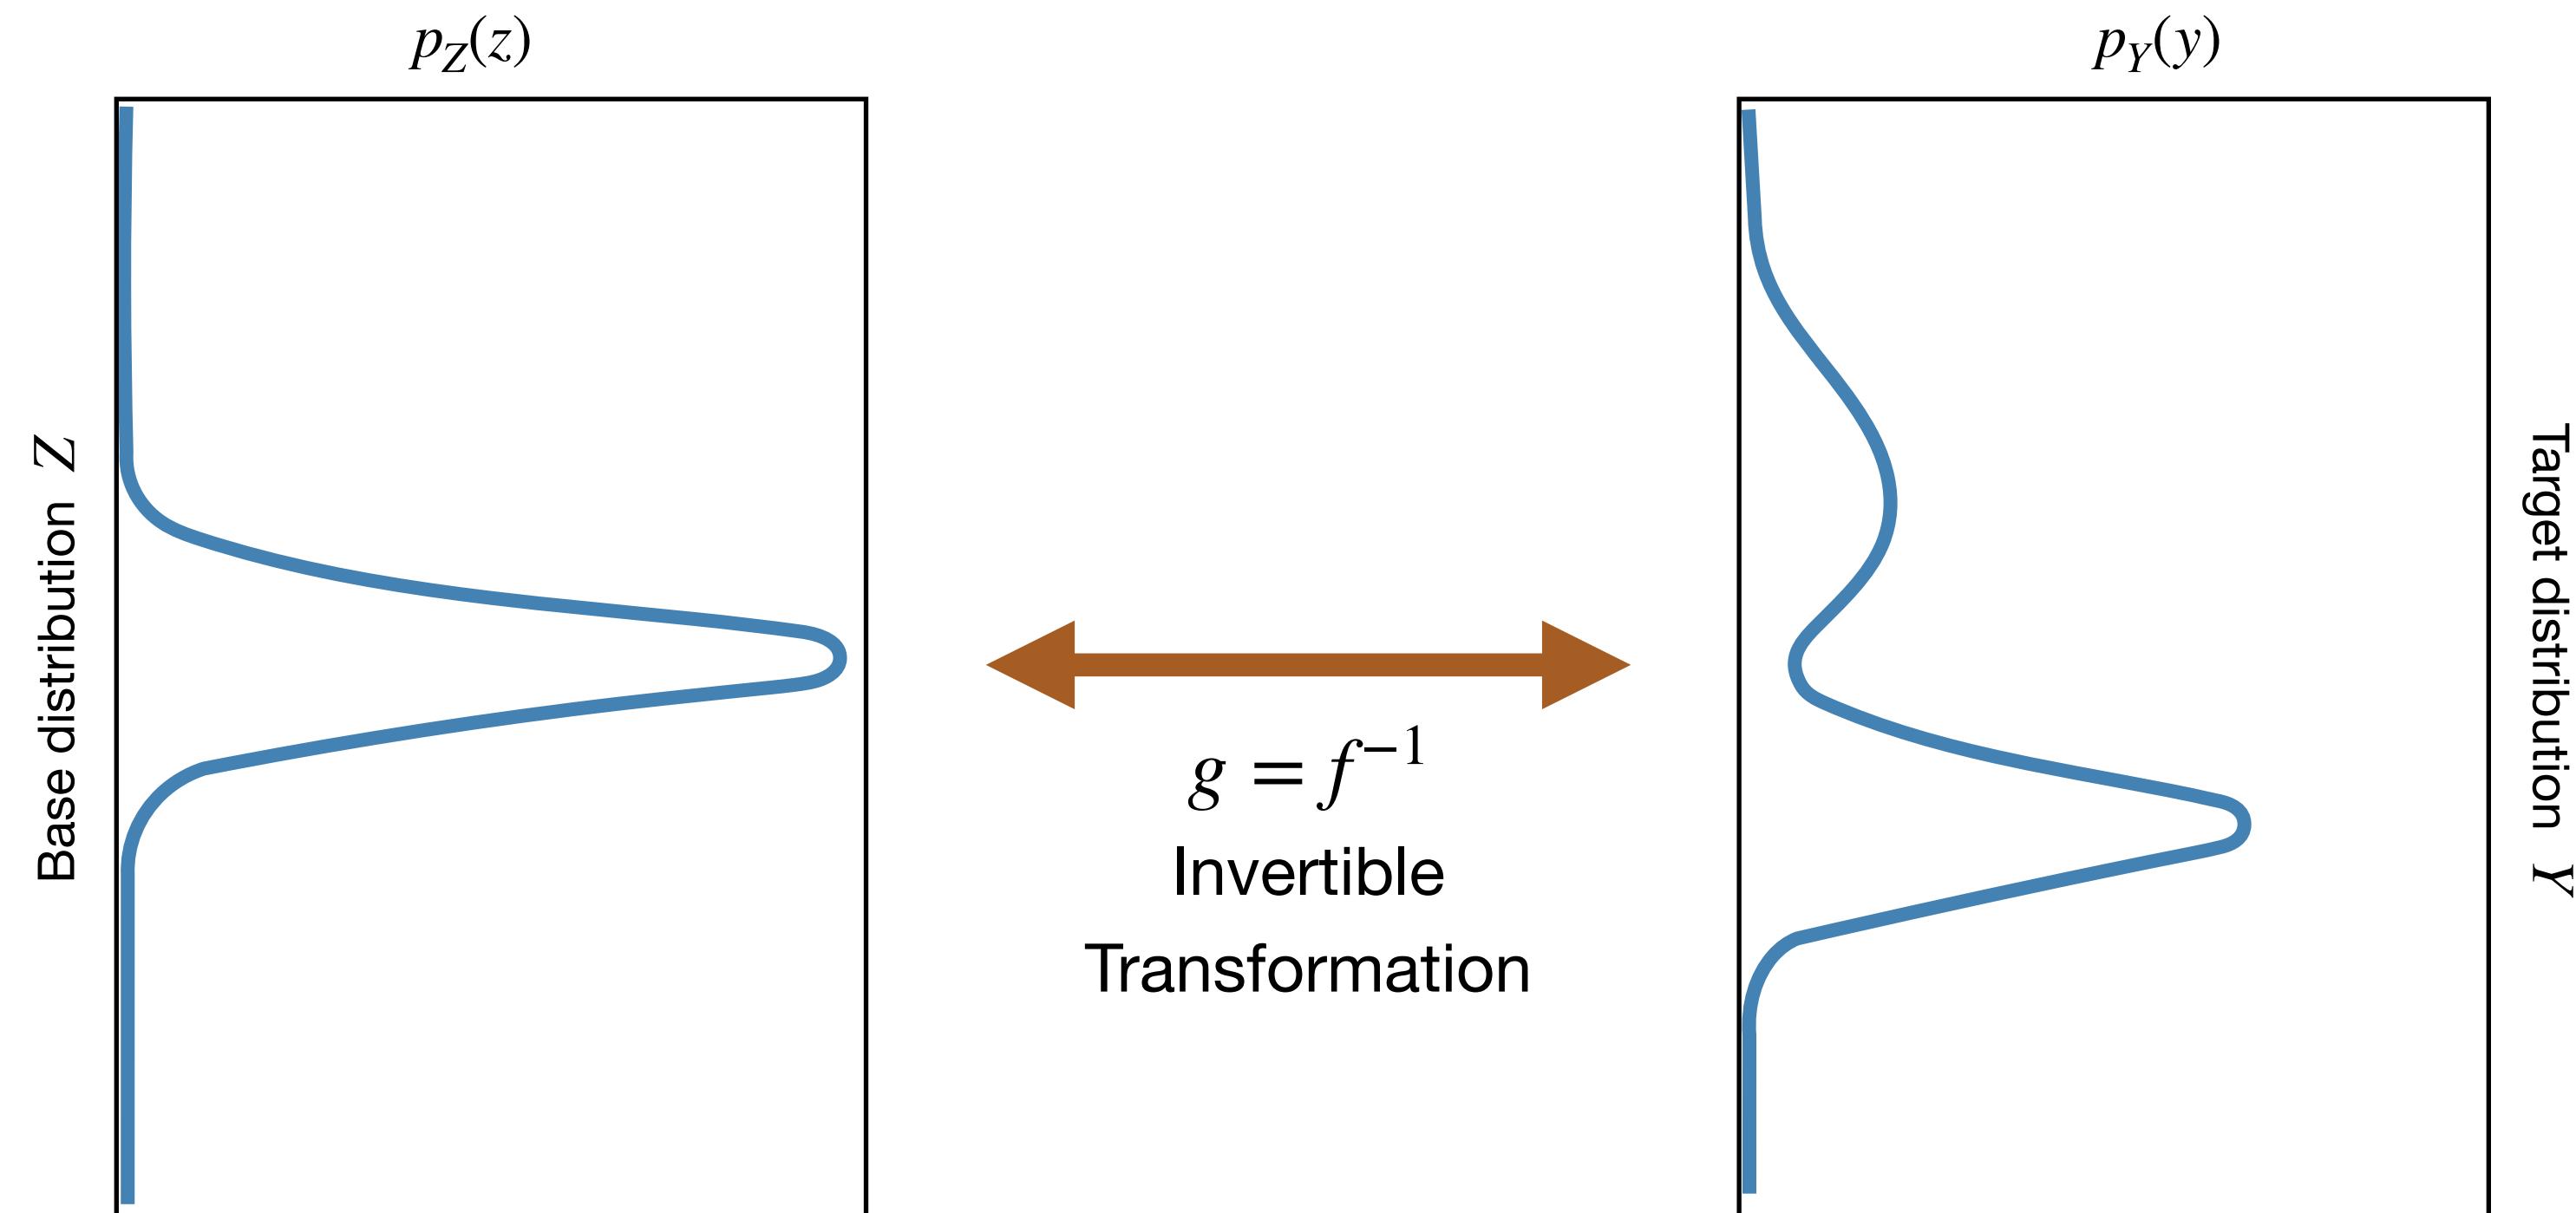
- Generating velocity given by:

$$u_t(x|x_1) = \frac{x_1 - (1 - \sigma_{\min})x}{1 - (1 - \sigma_{\min})t}$$

Note: without σ_{\min} is simply: $p_t(x|x_1) = \mathcal{N}(tx_1, (1 - t)I)$, $u_t(x|x_1) = (x_1 - x)/(1 - t)$

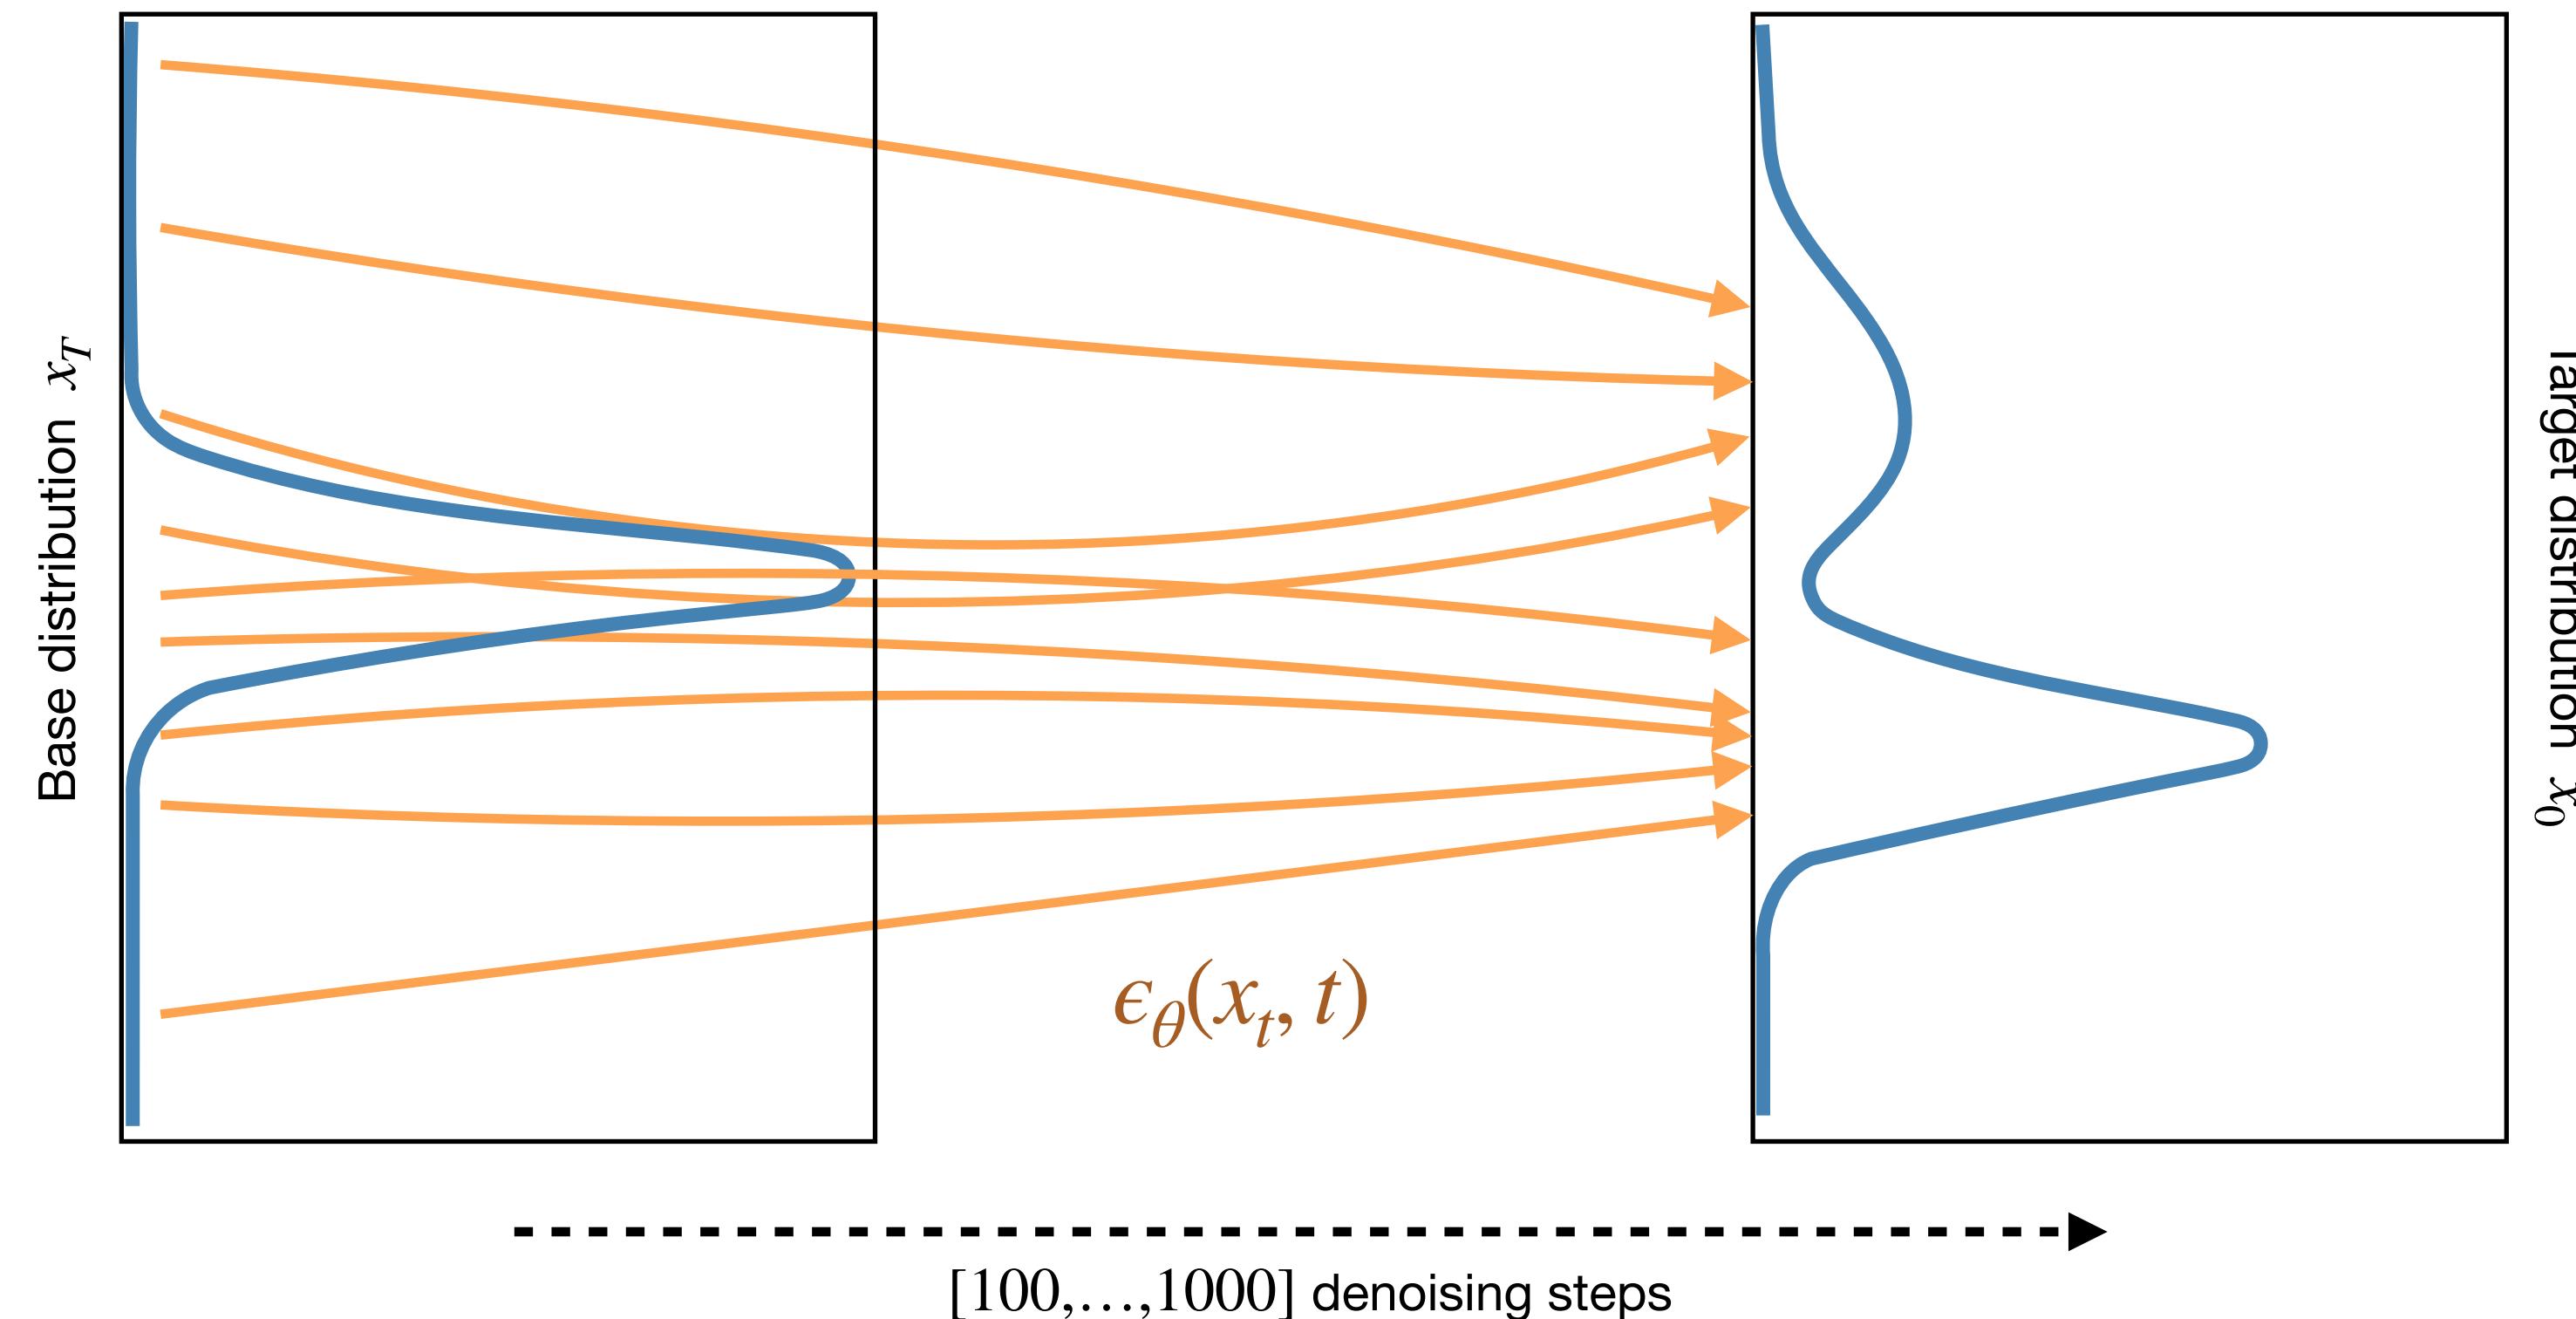
Denoising Diffusion Probabilistic Model

Previously: Normalizing Flows



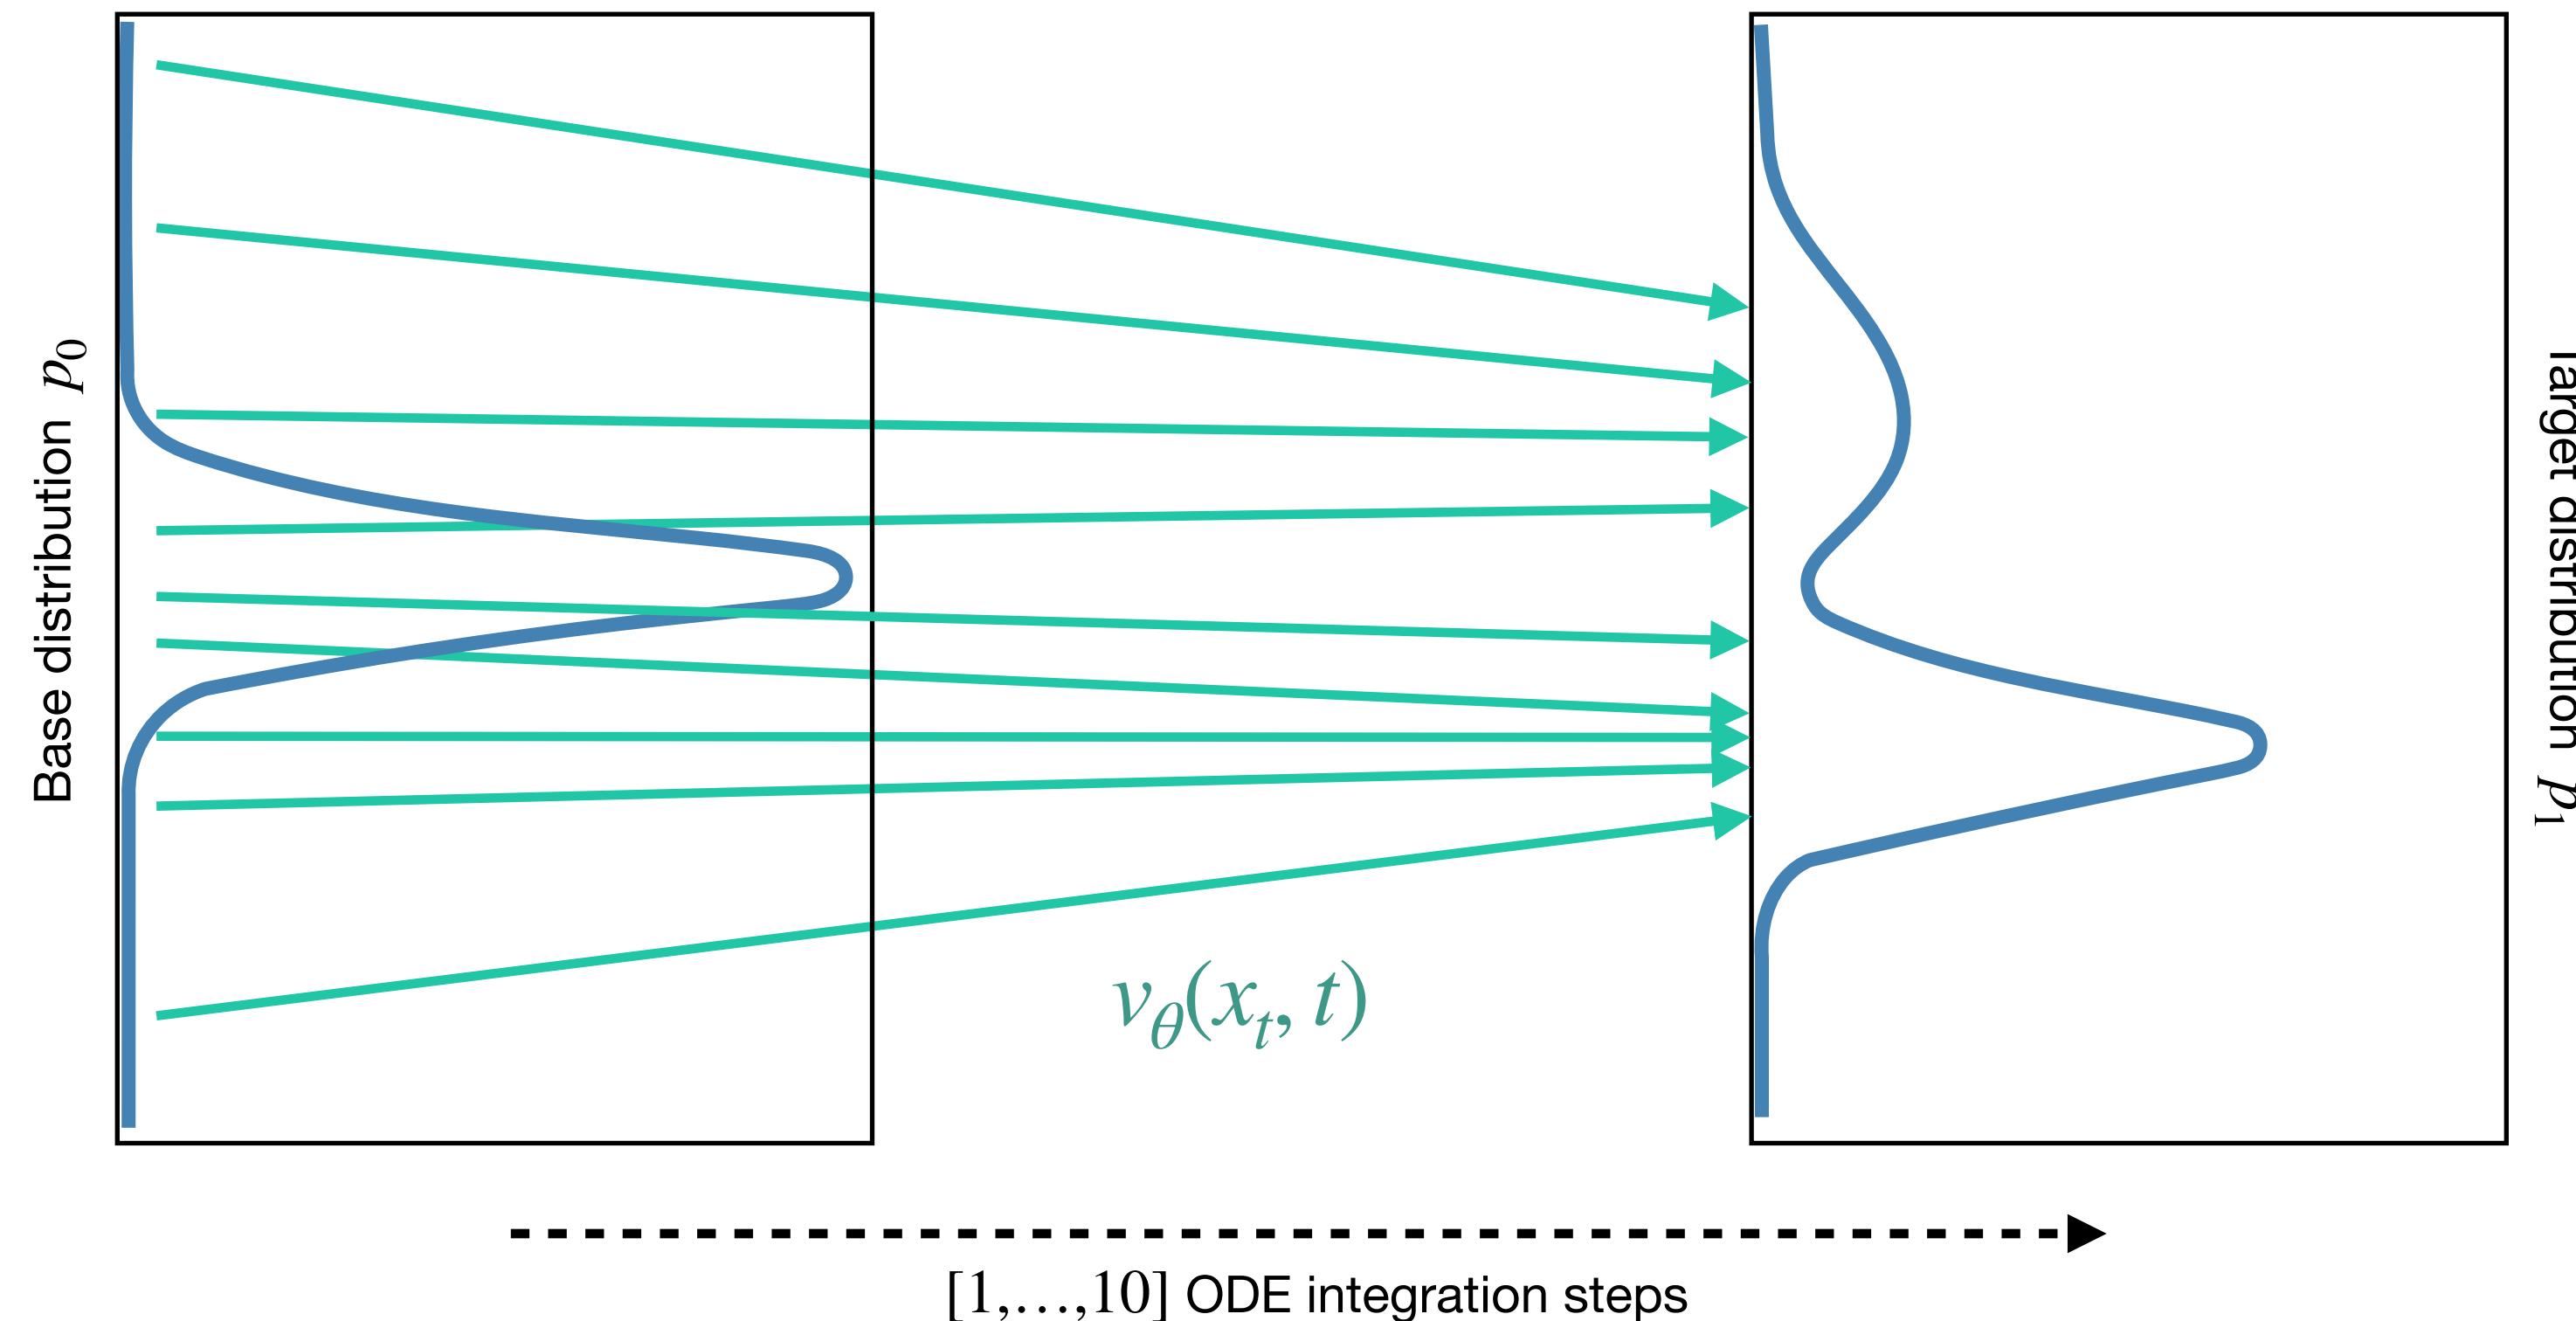
Flow Matching

Previously: DDPM



Flow Matching

Rectified Flows / Flow Matching



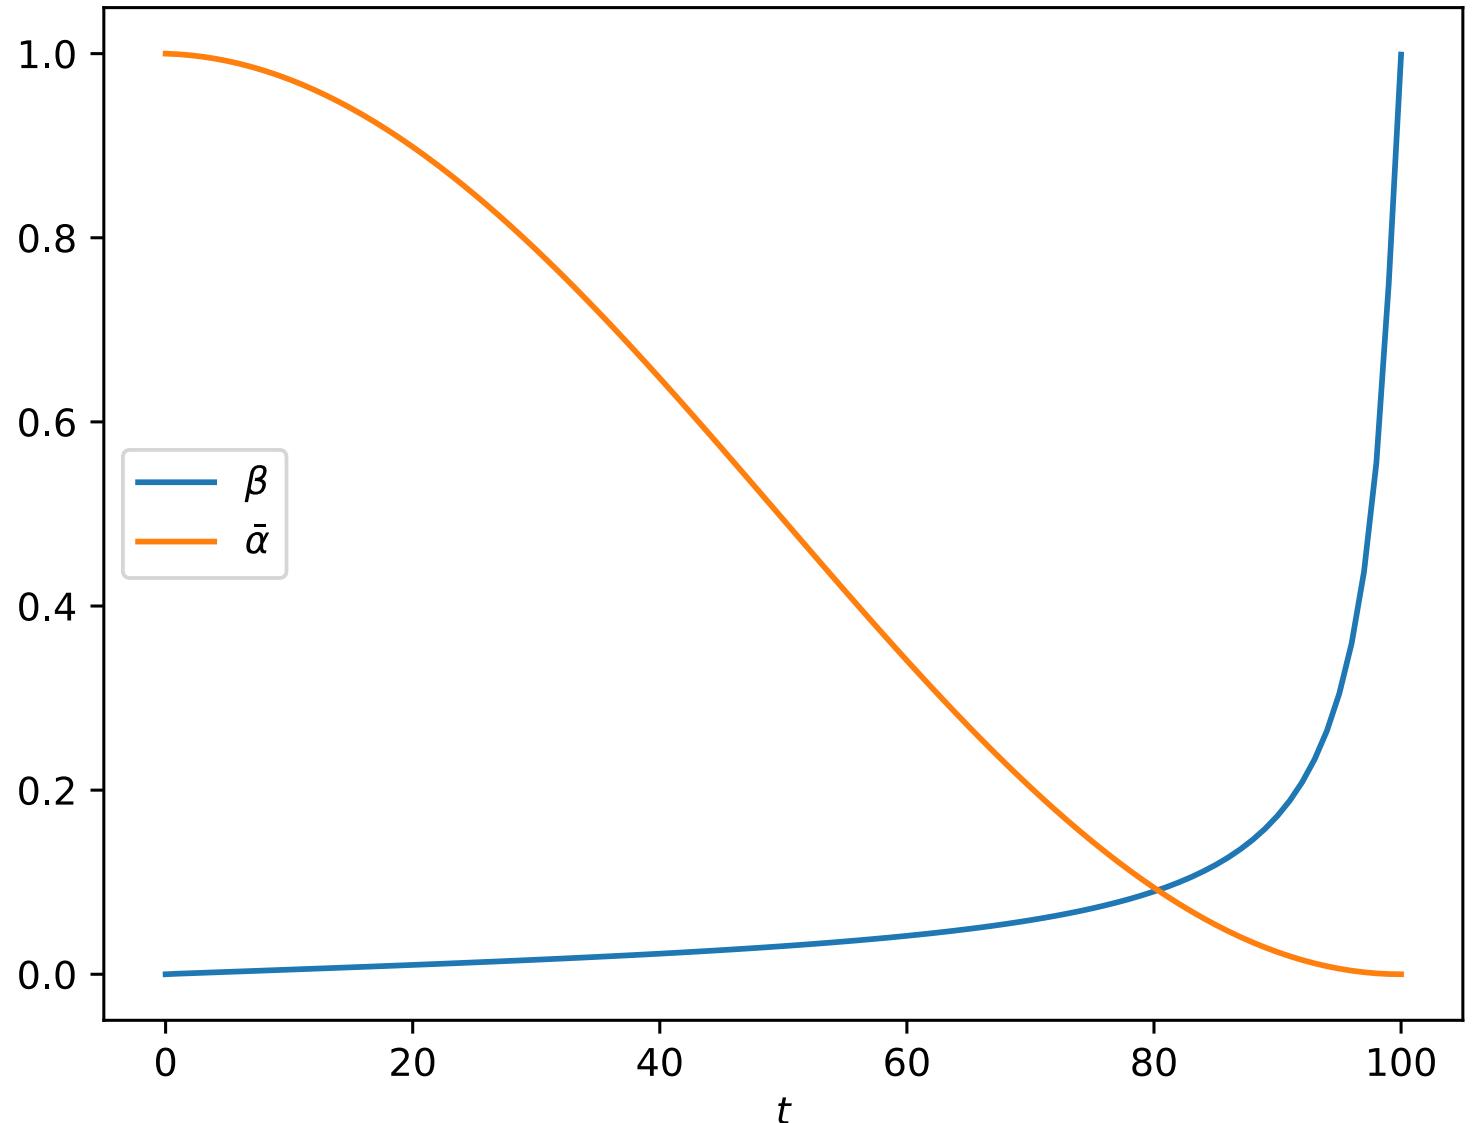
Ideally, converges
in a single step!

Flow Matching why is it better?

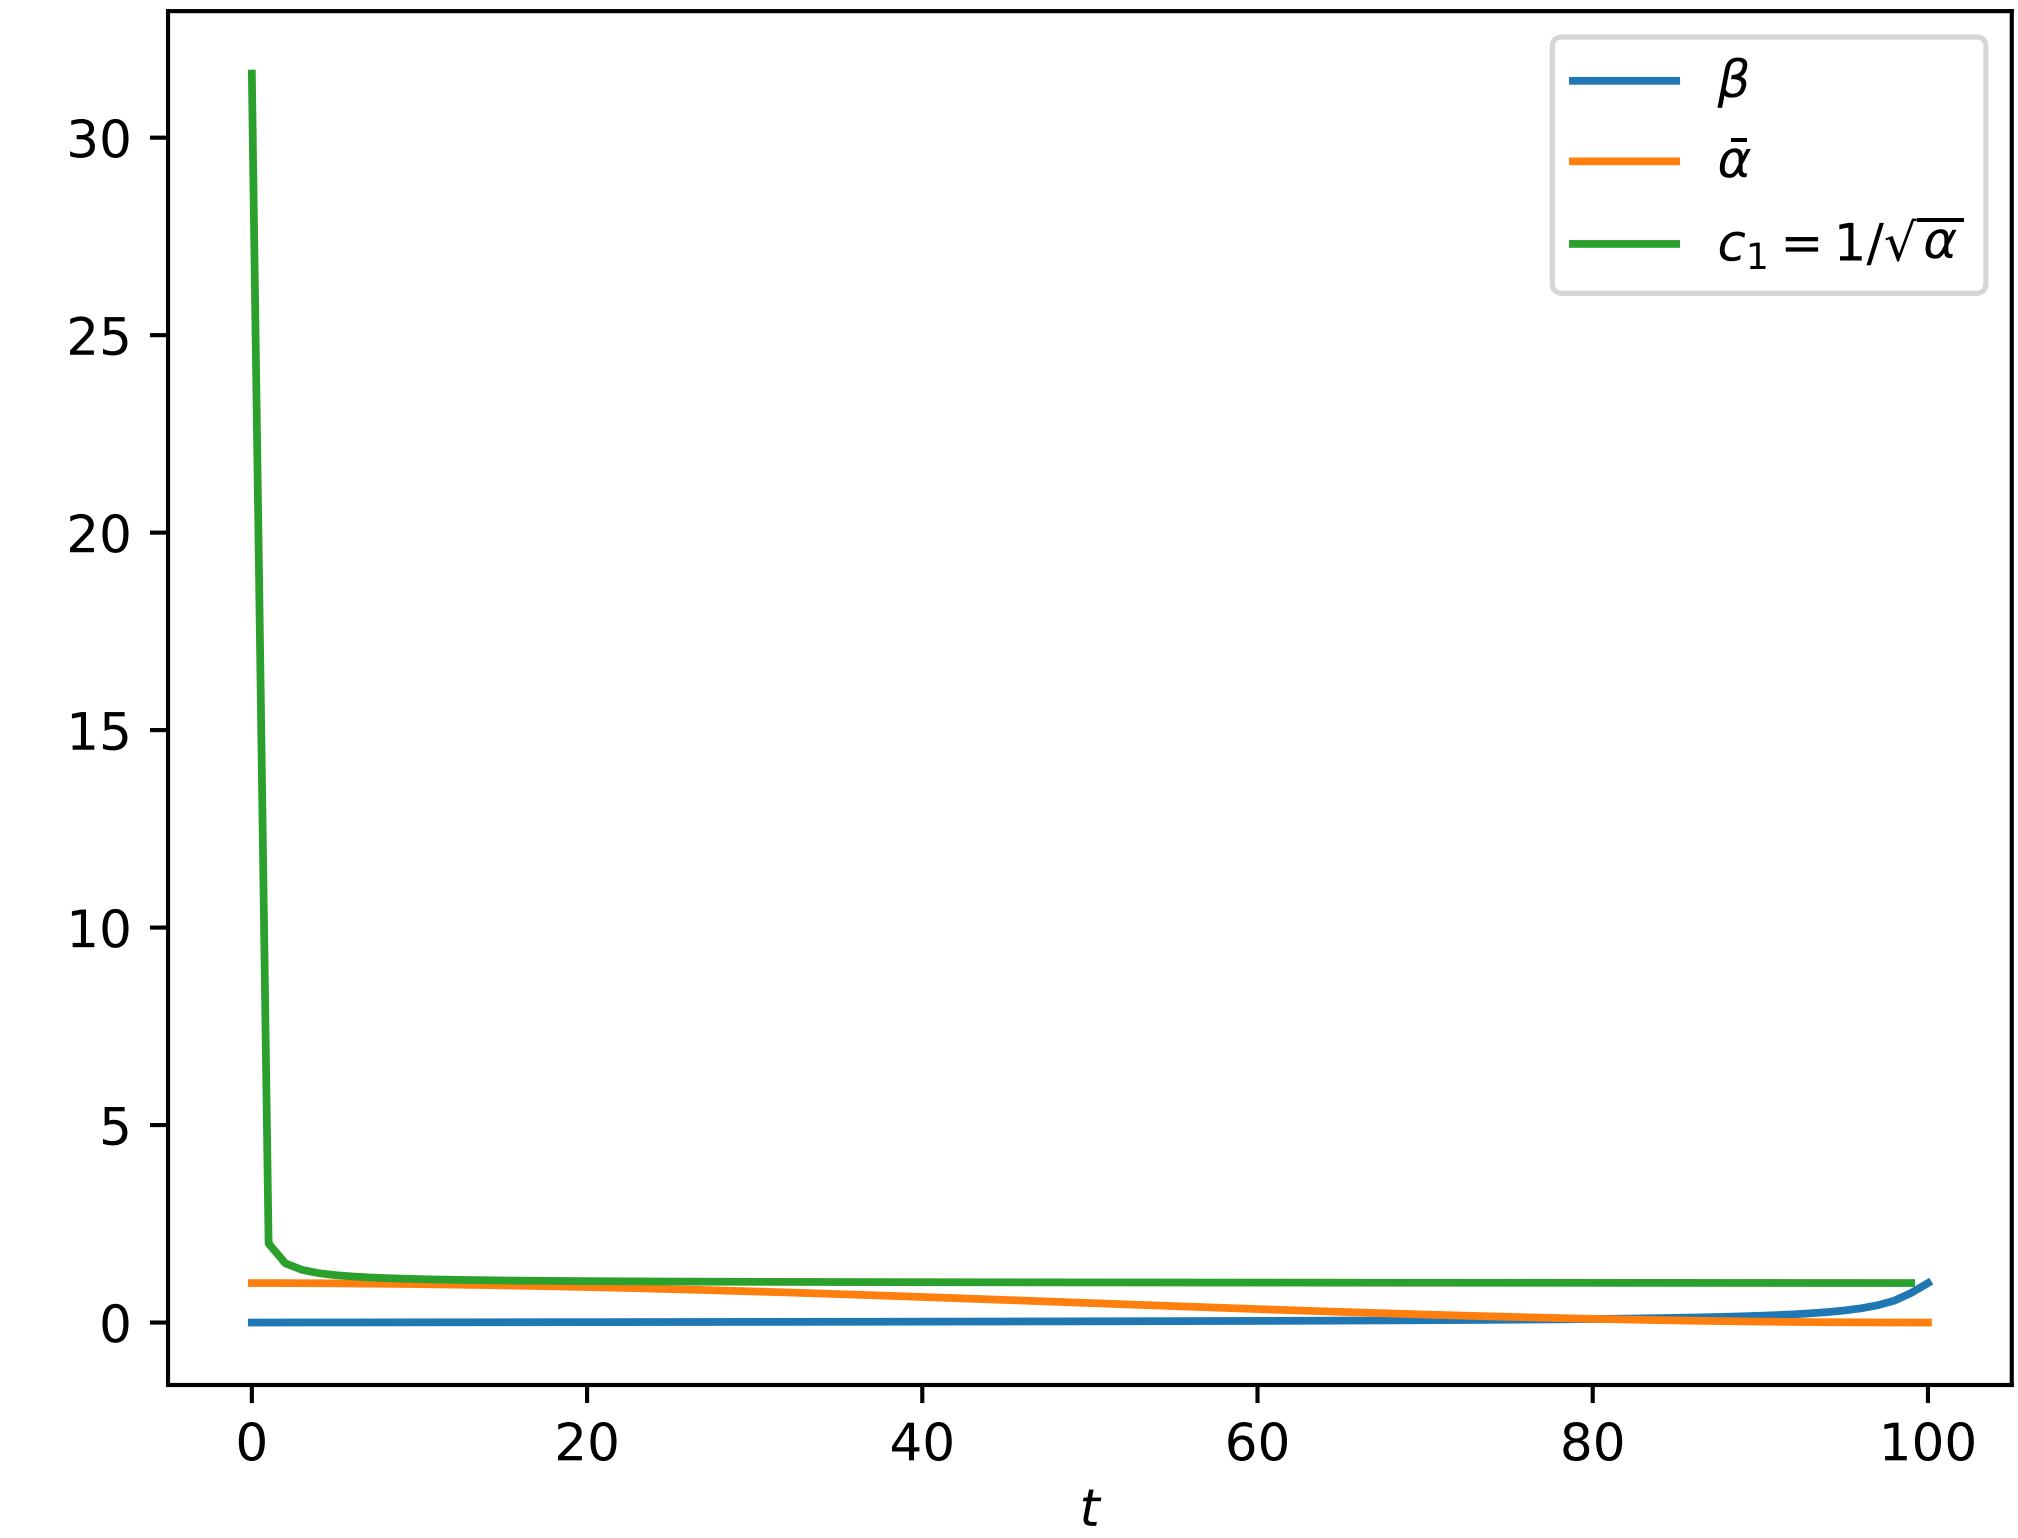
Reminder: Sampling with DDPM

Algorithm 2 Sampling

```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$ 
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\bar{\alpha}_t}} \left( \mathbf{x}_t - \frac{1-\bar{\alpha}_t}{\sqrt{1-\bar{\alpha}_t}} \epsilon_\theta(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$ 
5: end for
6: return  $\mathbf{x}_0$ 
```



$\bar{\alpha}, \beta$ are well behaved



$1/\sqrt{\bar{\alpha}}$ is not!

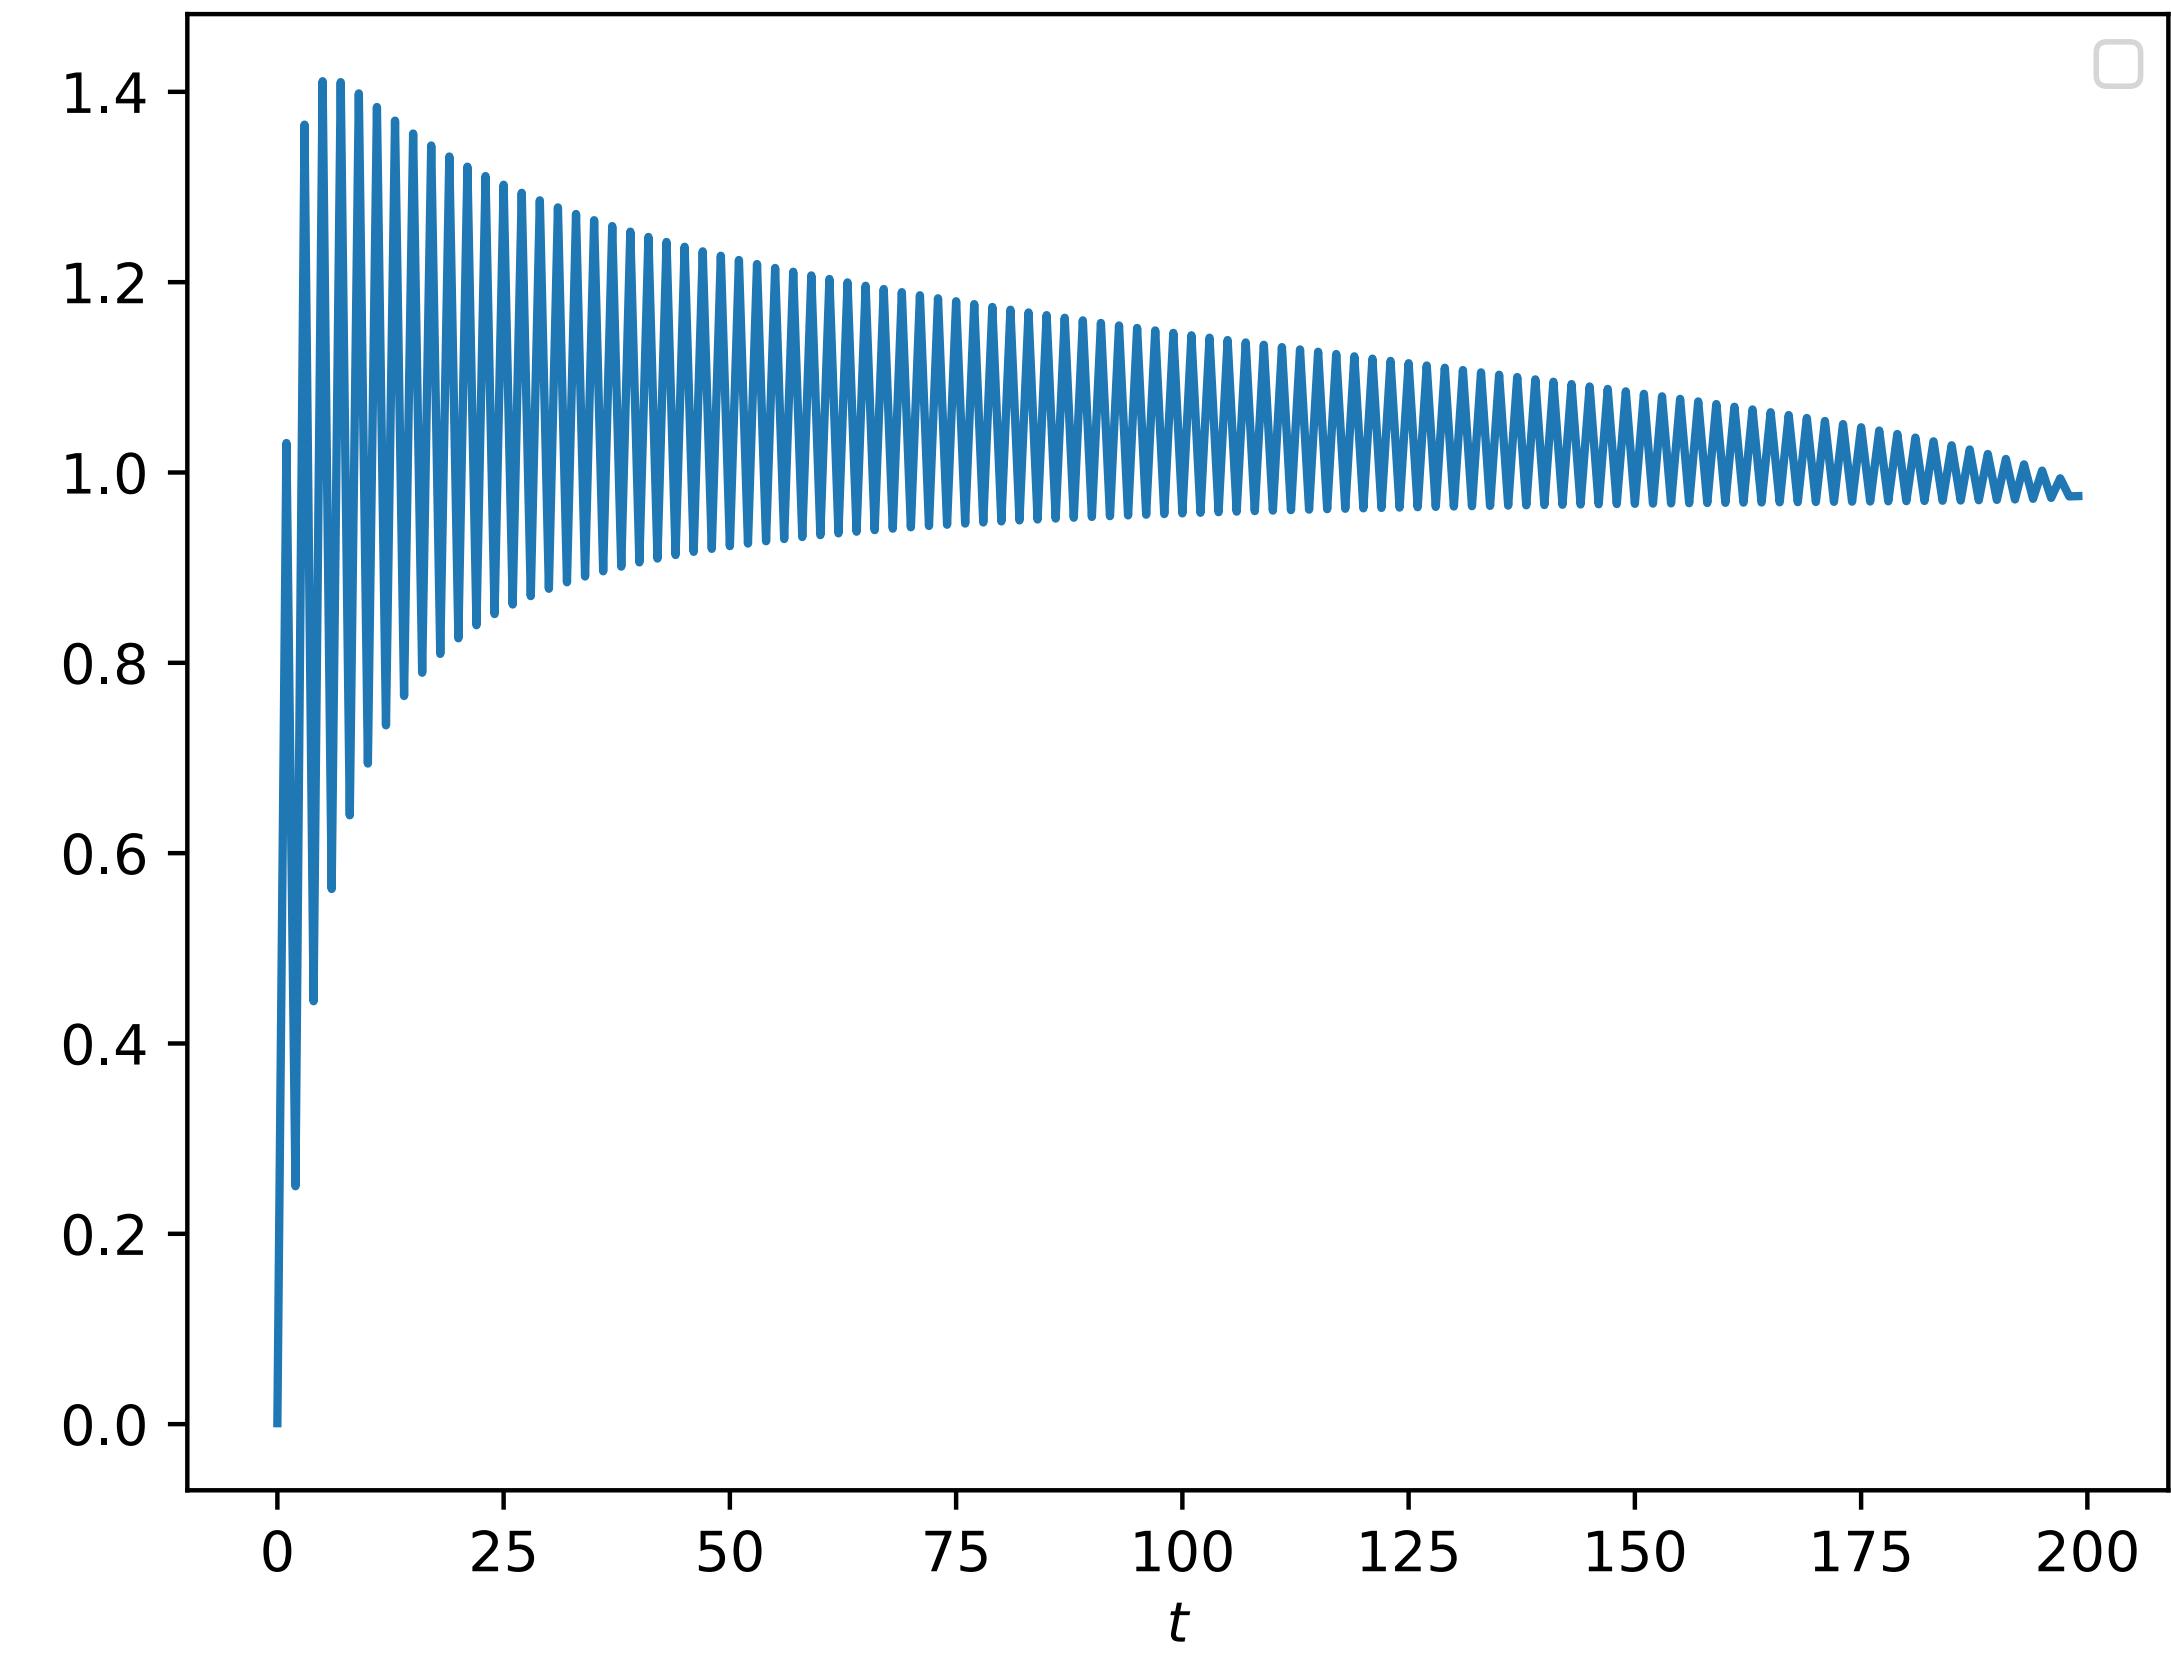
Flow Matching why is it better?

Reminder: Sampling with DDPM

Algorithm 2 Sampling

```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$  #1
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t}} \epsilon_\theta(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$  #2
5: end for
6: return  $\mathbf{x}_0$ 
```

- DDPM: Great training objective, but inference relies on high accuracy of NN
- Zig-zagging behavior (shown right)
- Flow matching: more stable, strictly linear!



Flow Matching

- Very simple implementation:

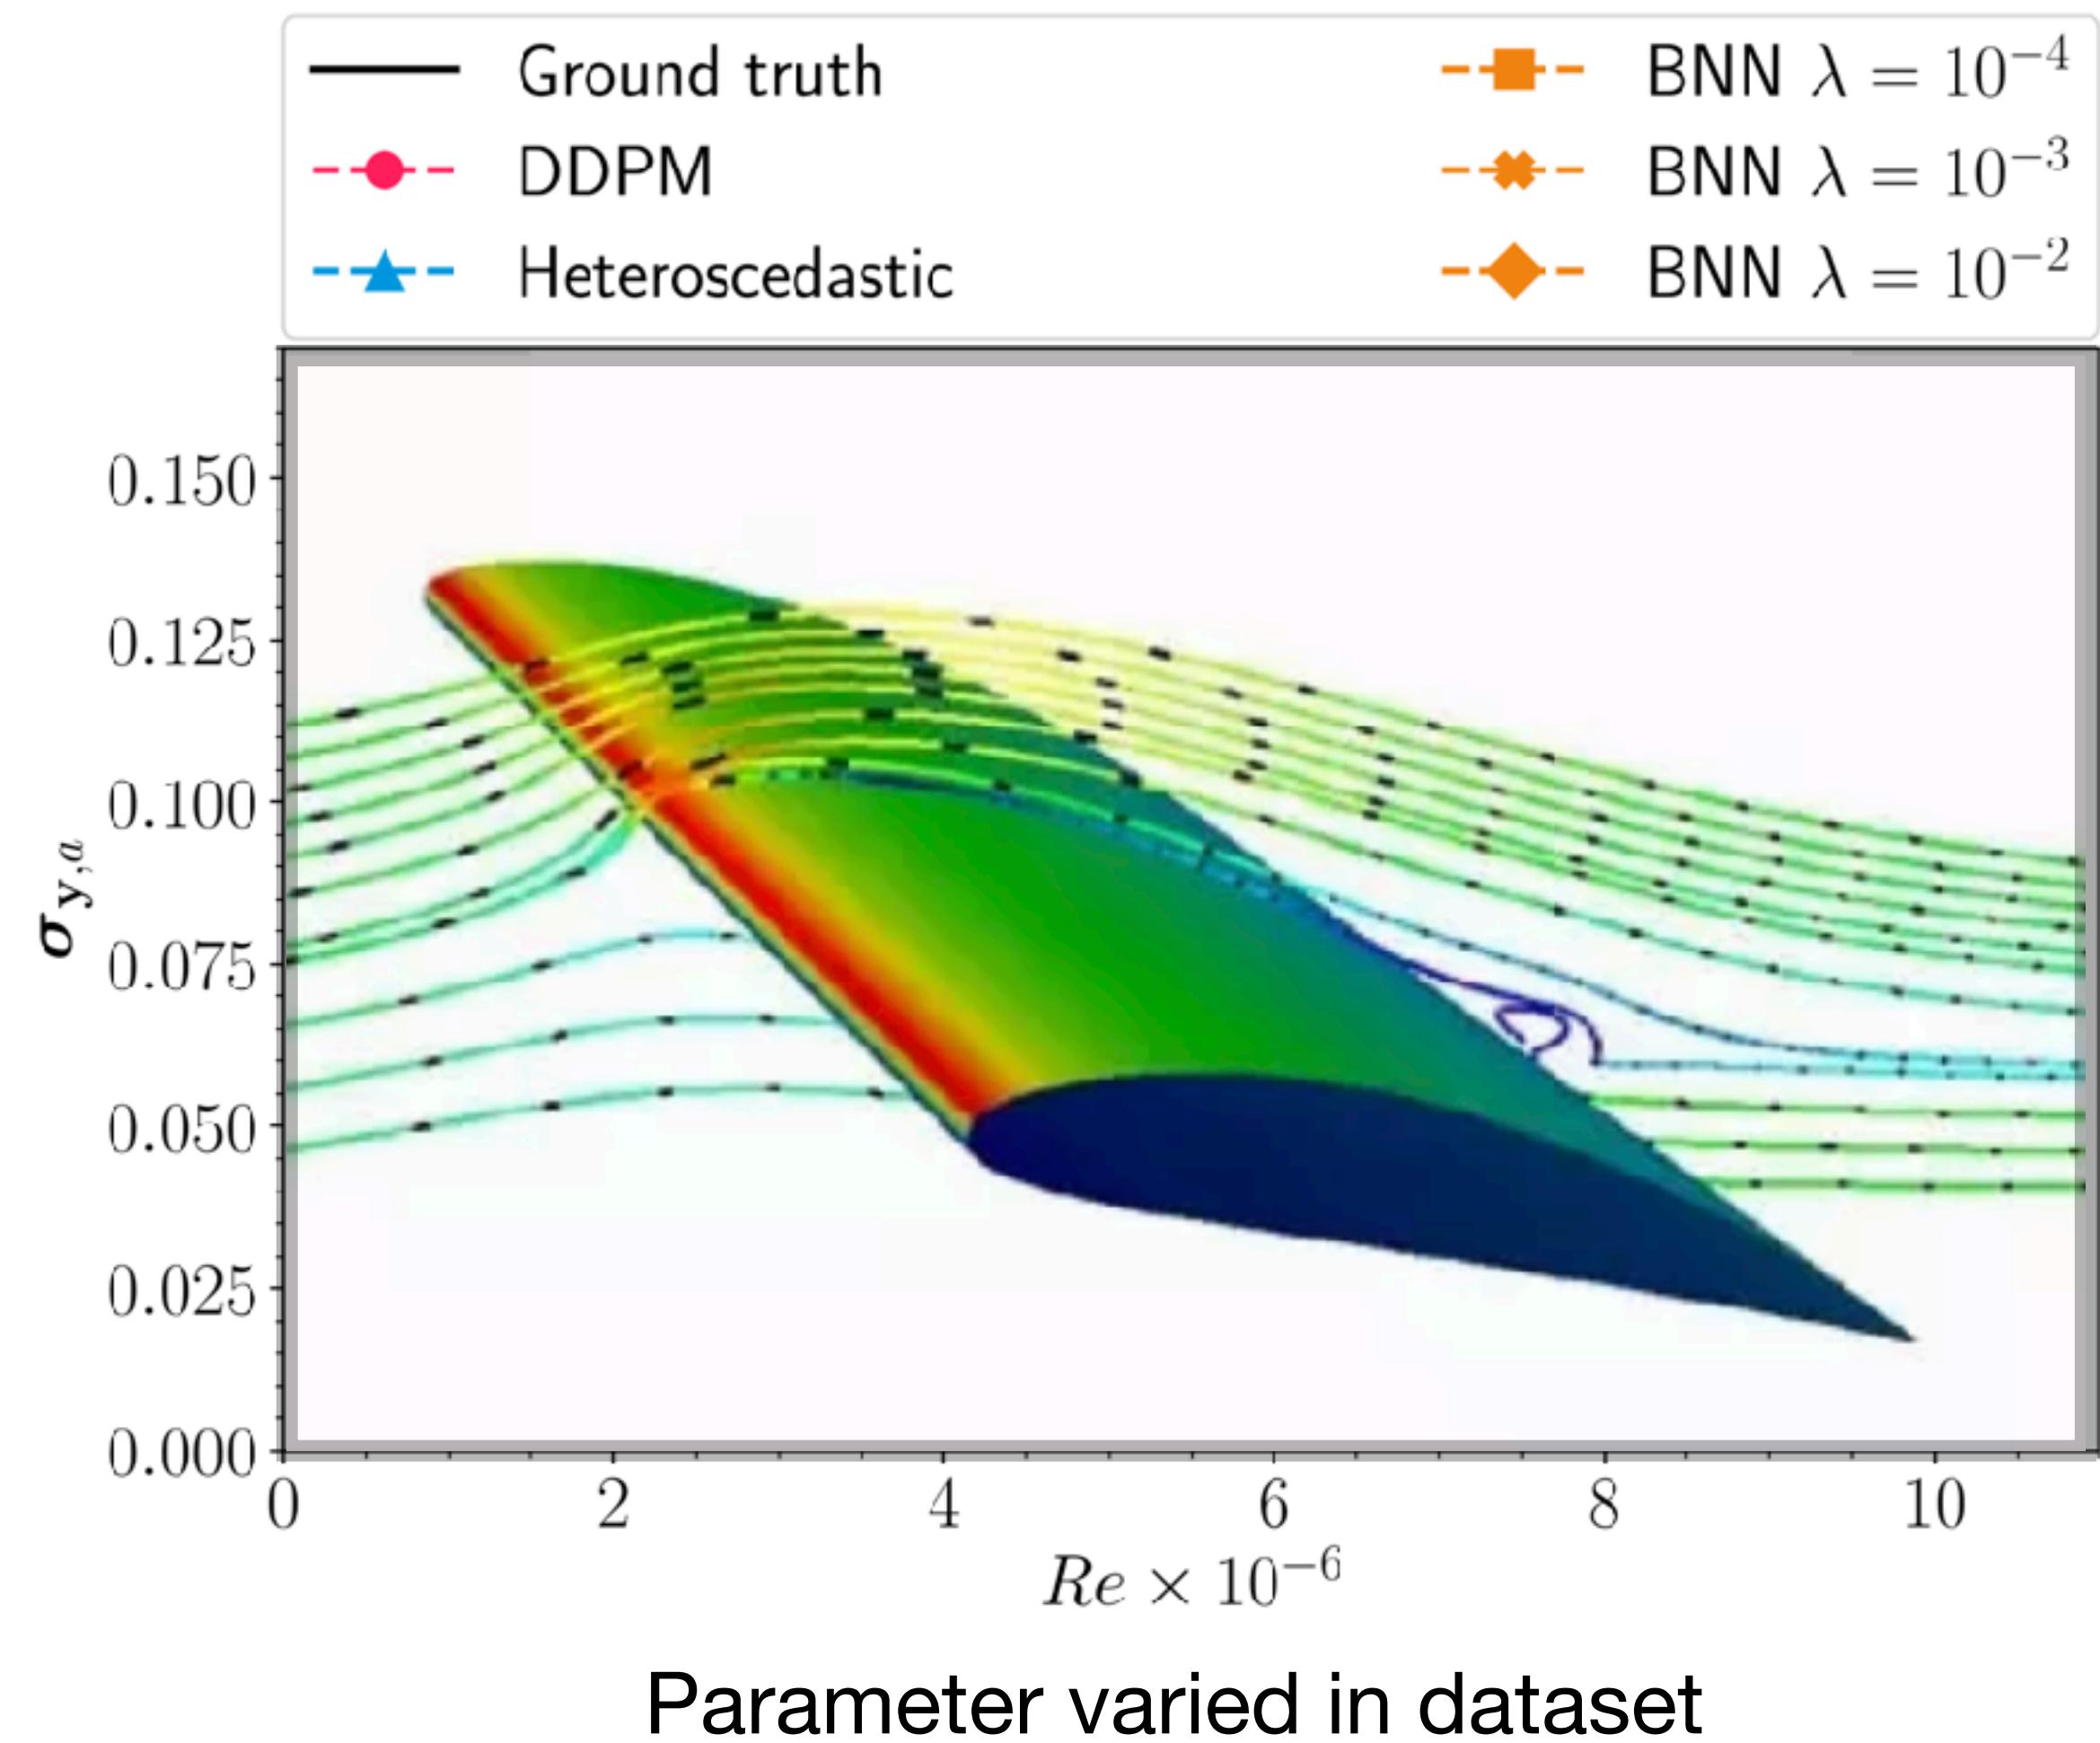
```
t = np.random.rand()  
x_t = (1 - (1 - self.sigma_min) * t) * x0 + t * xl  
u_t = (xl - x0)
```

- At inference time, simply integrate: `x += dt * model(x,t)`
- Fancier (higher-order) integration also possible...
- Resulting method: fast, and gives full access to posterior distribution!
- Try yourself: <https://.../probmodels-flowmatching.ipynb>

Flow Matching & Co. in Action

Known Ground Truth Distribution

Turbulent NS case with varying
Reynolds number:



Known Ground Truth Distribution

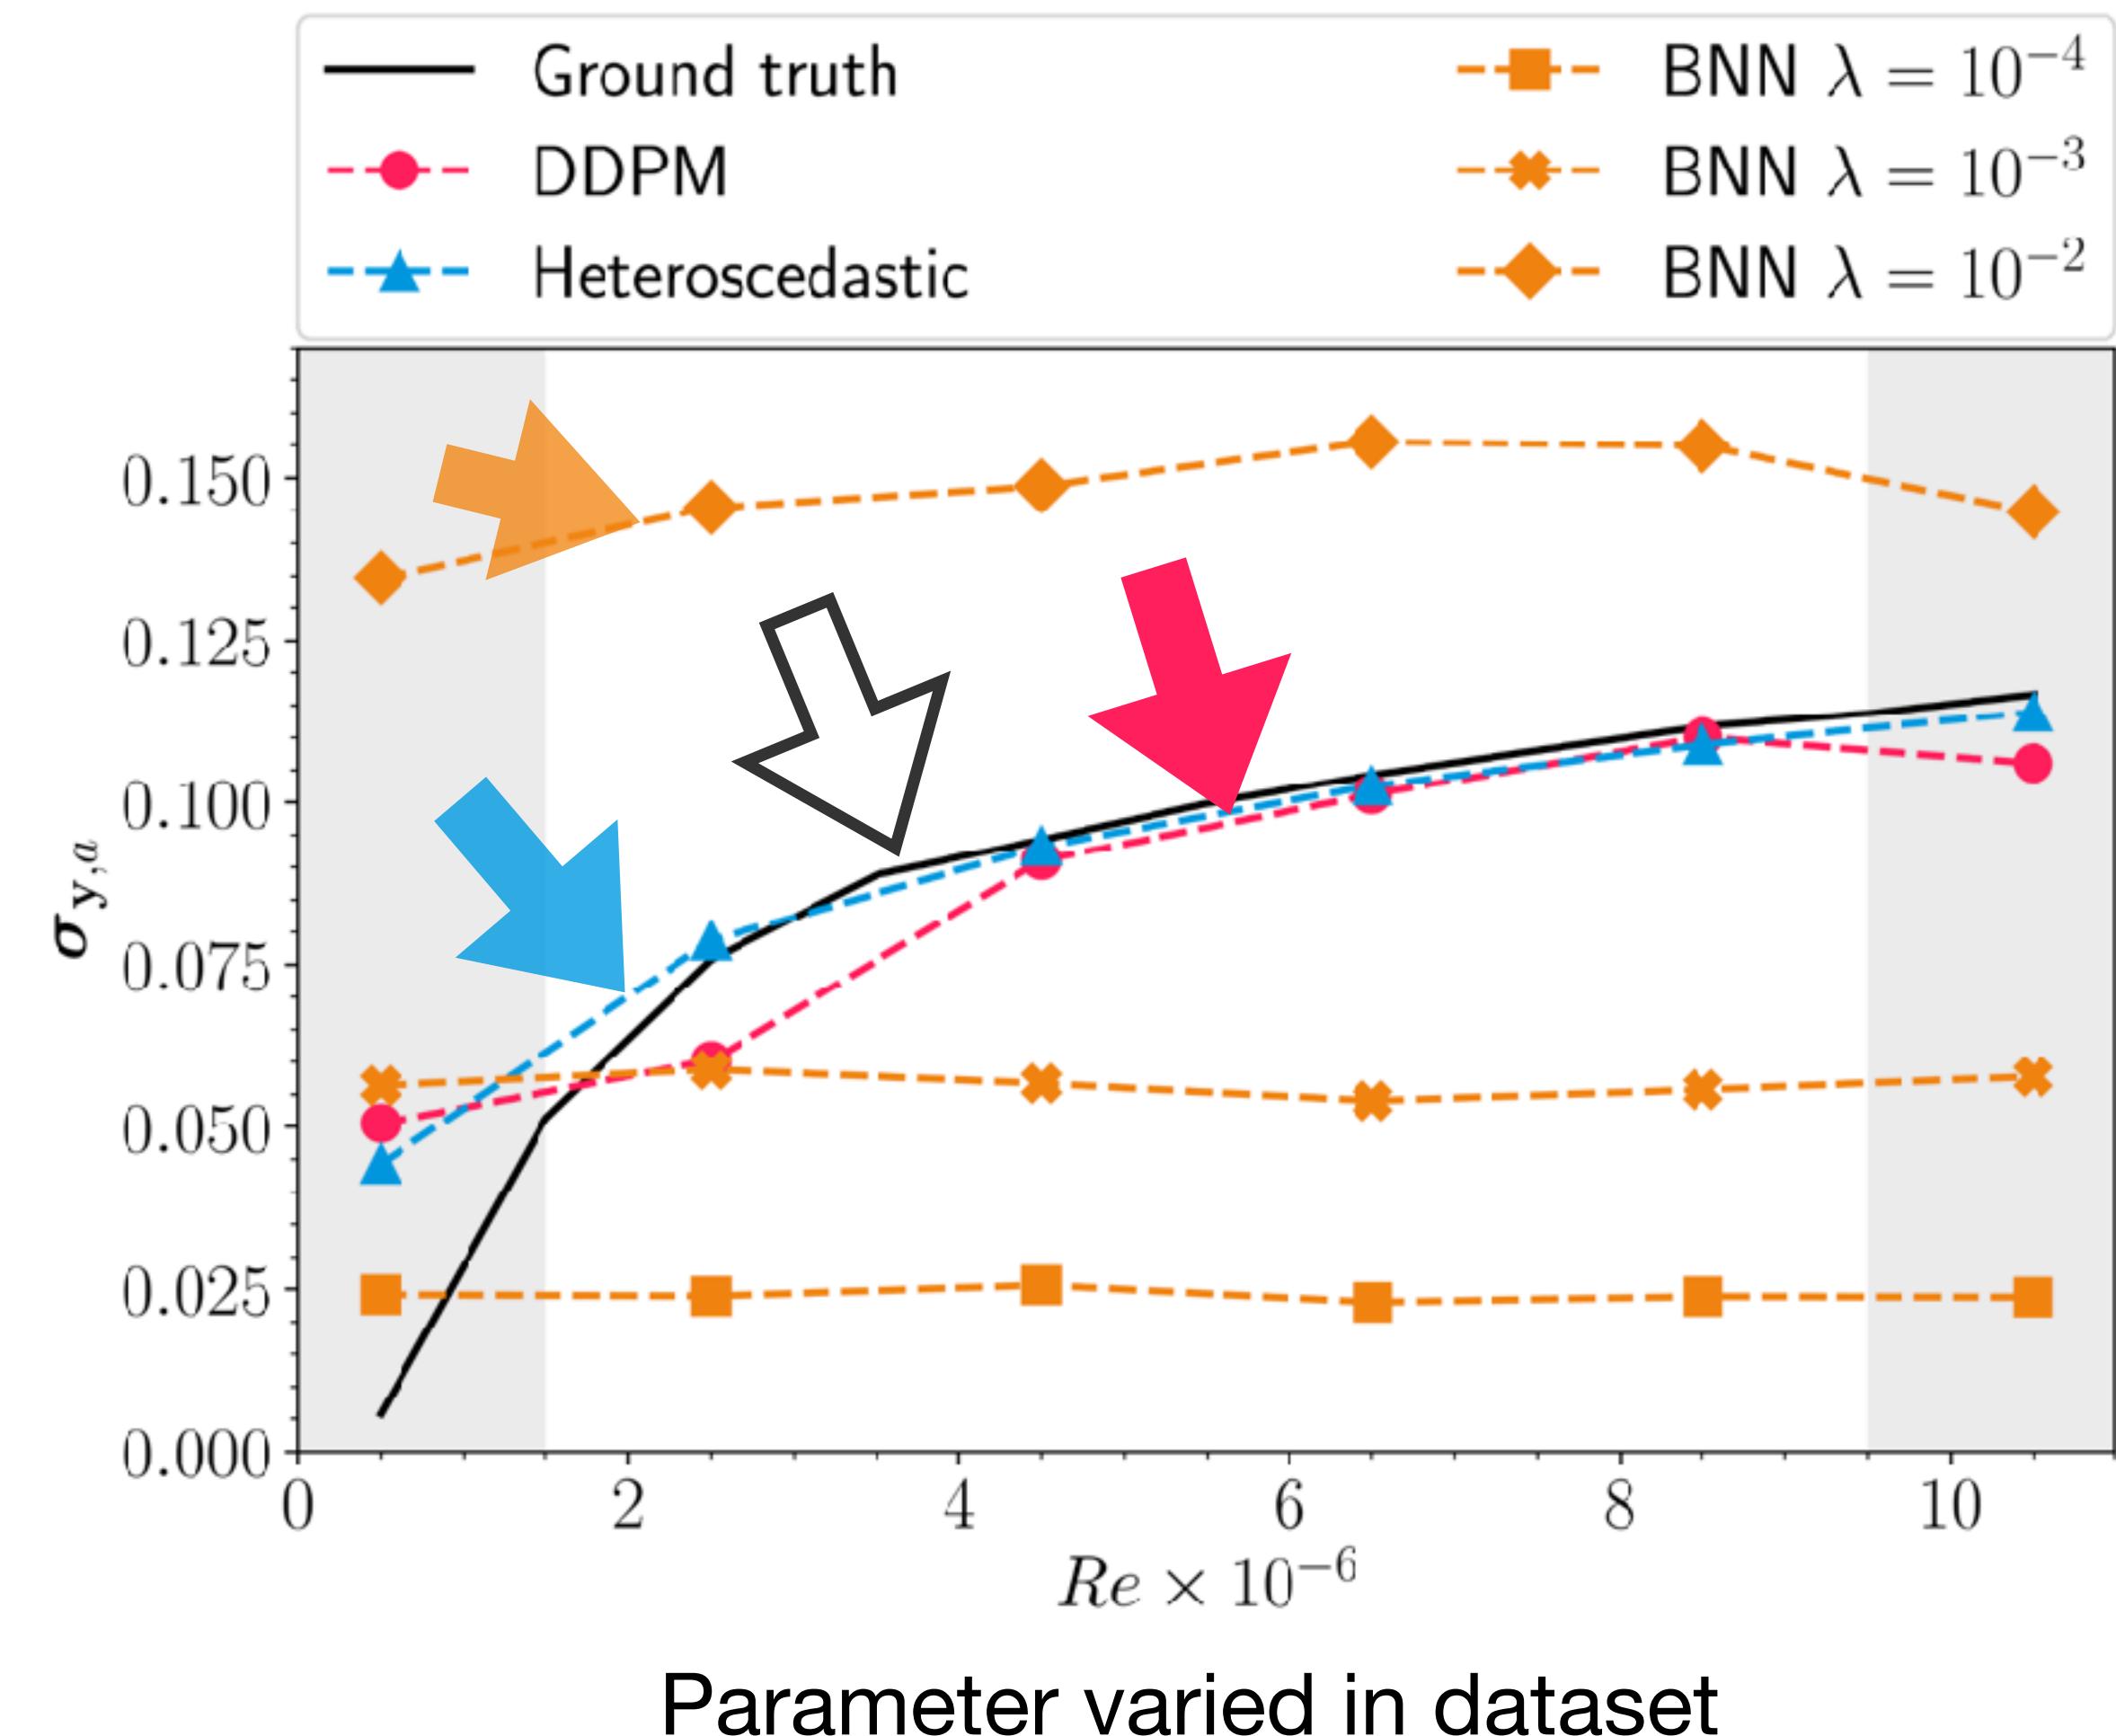
Turbulent NS case with varying
Reynolds number:

Ground Truth

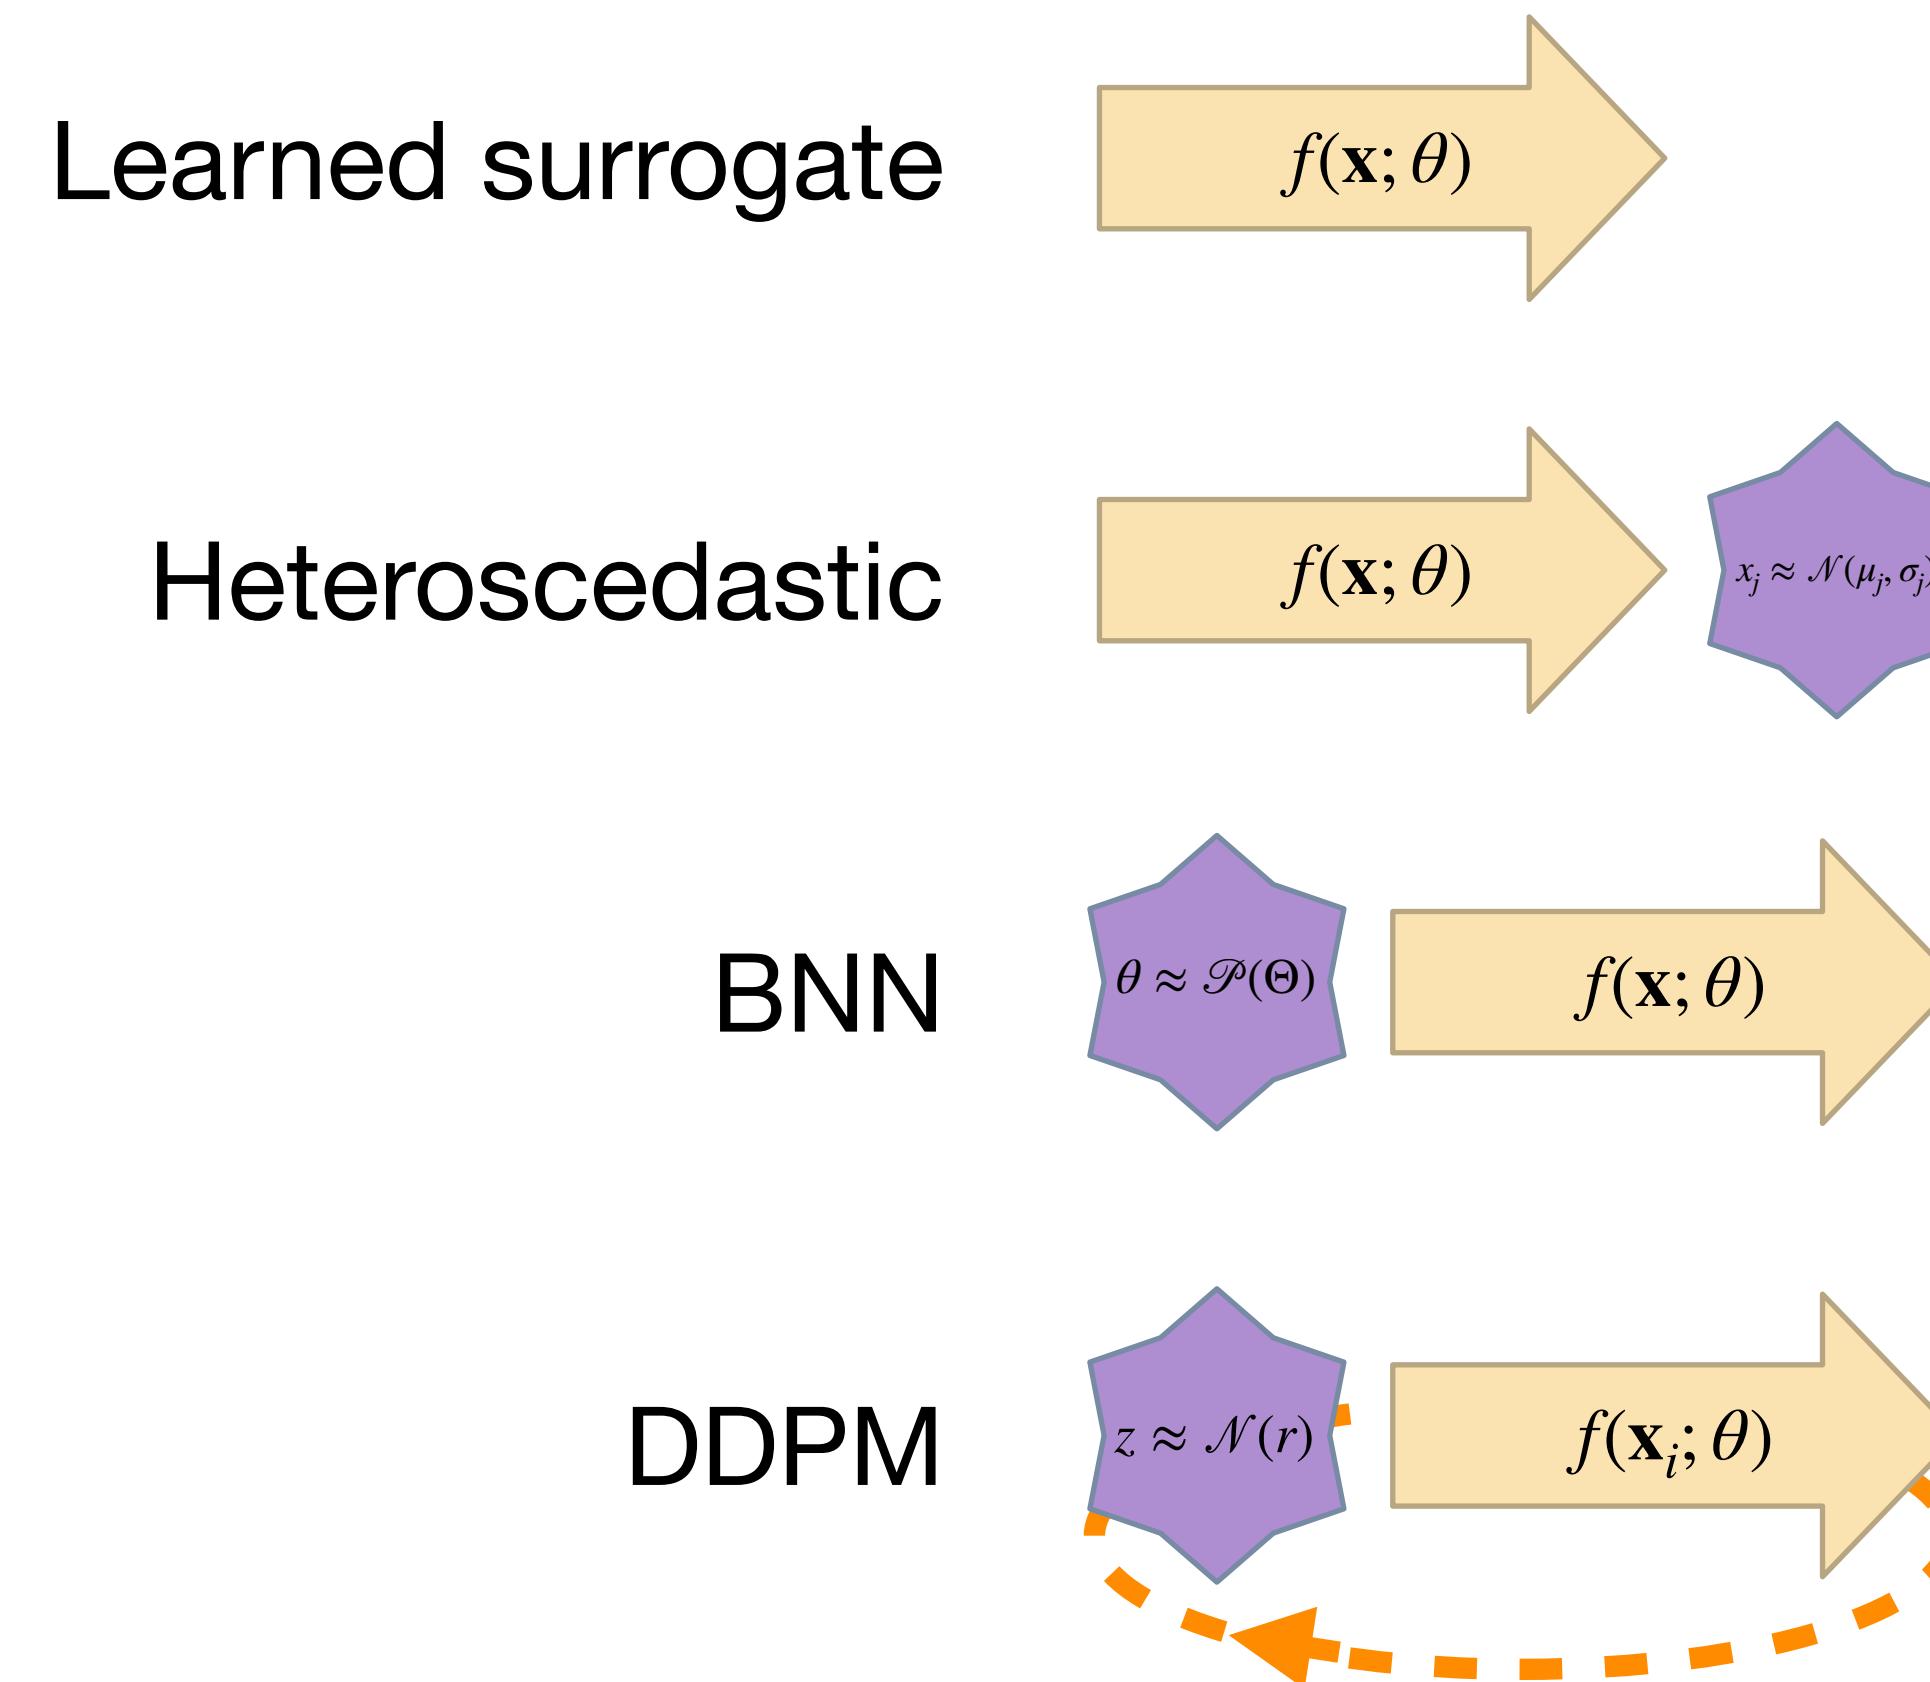
Bayesian NN

Heteroscedastic

Diffusion Model



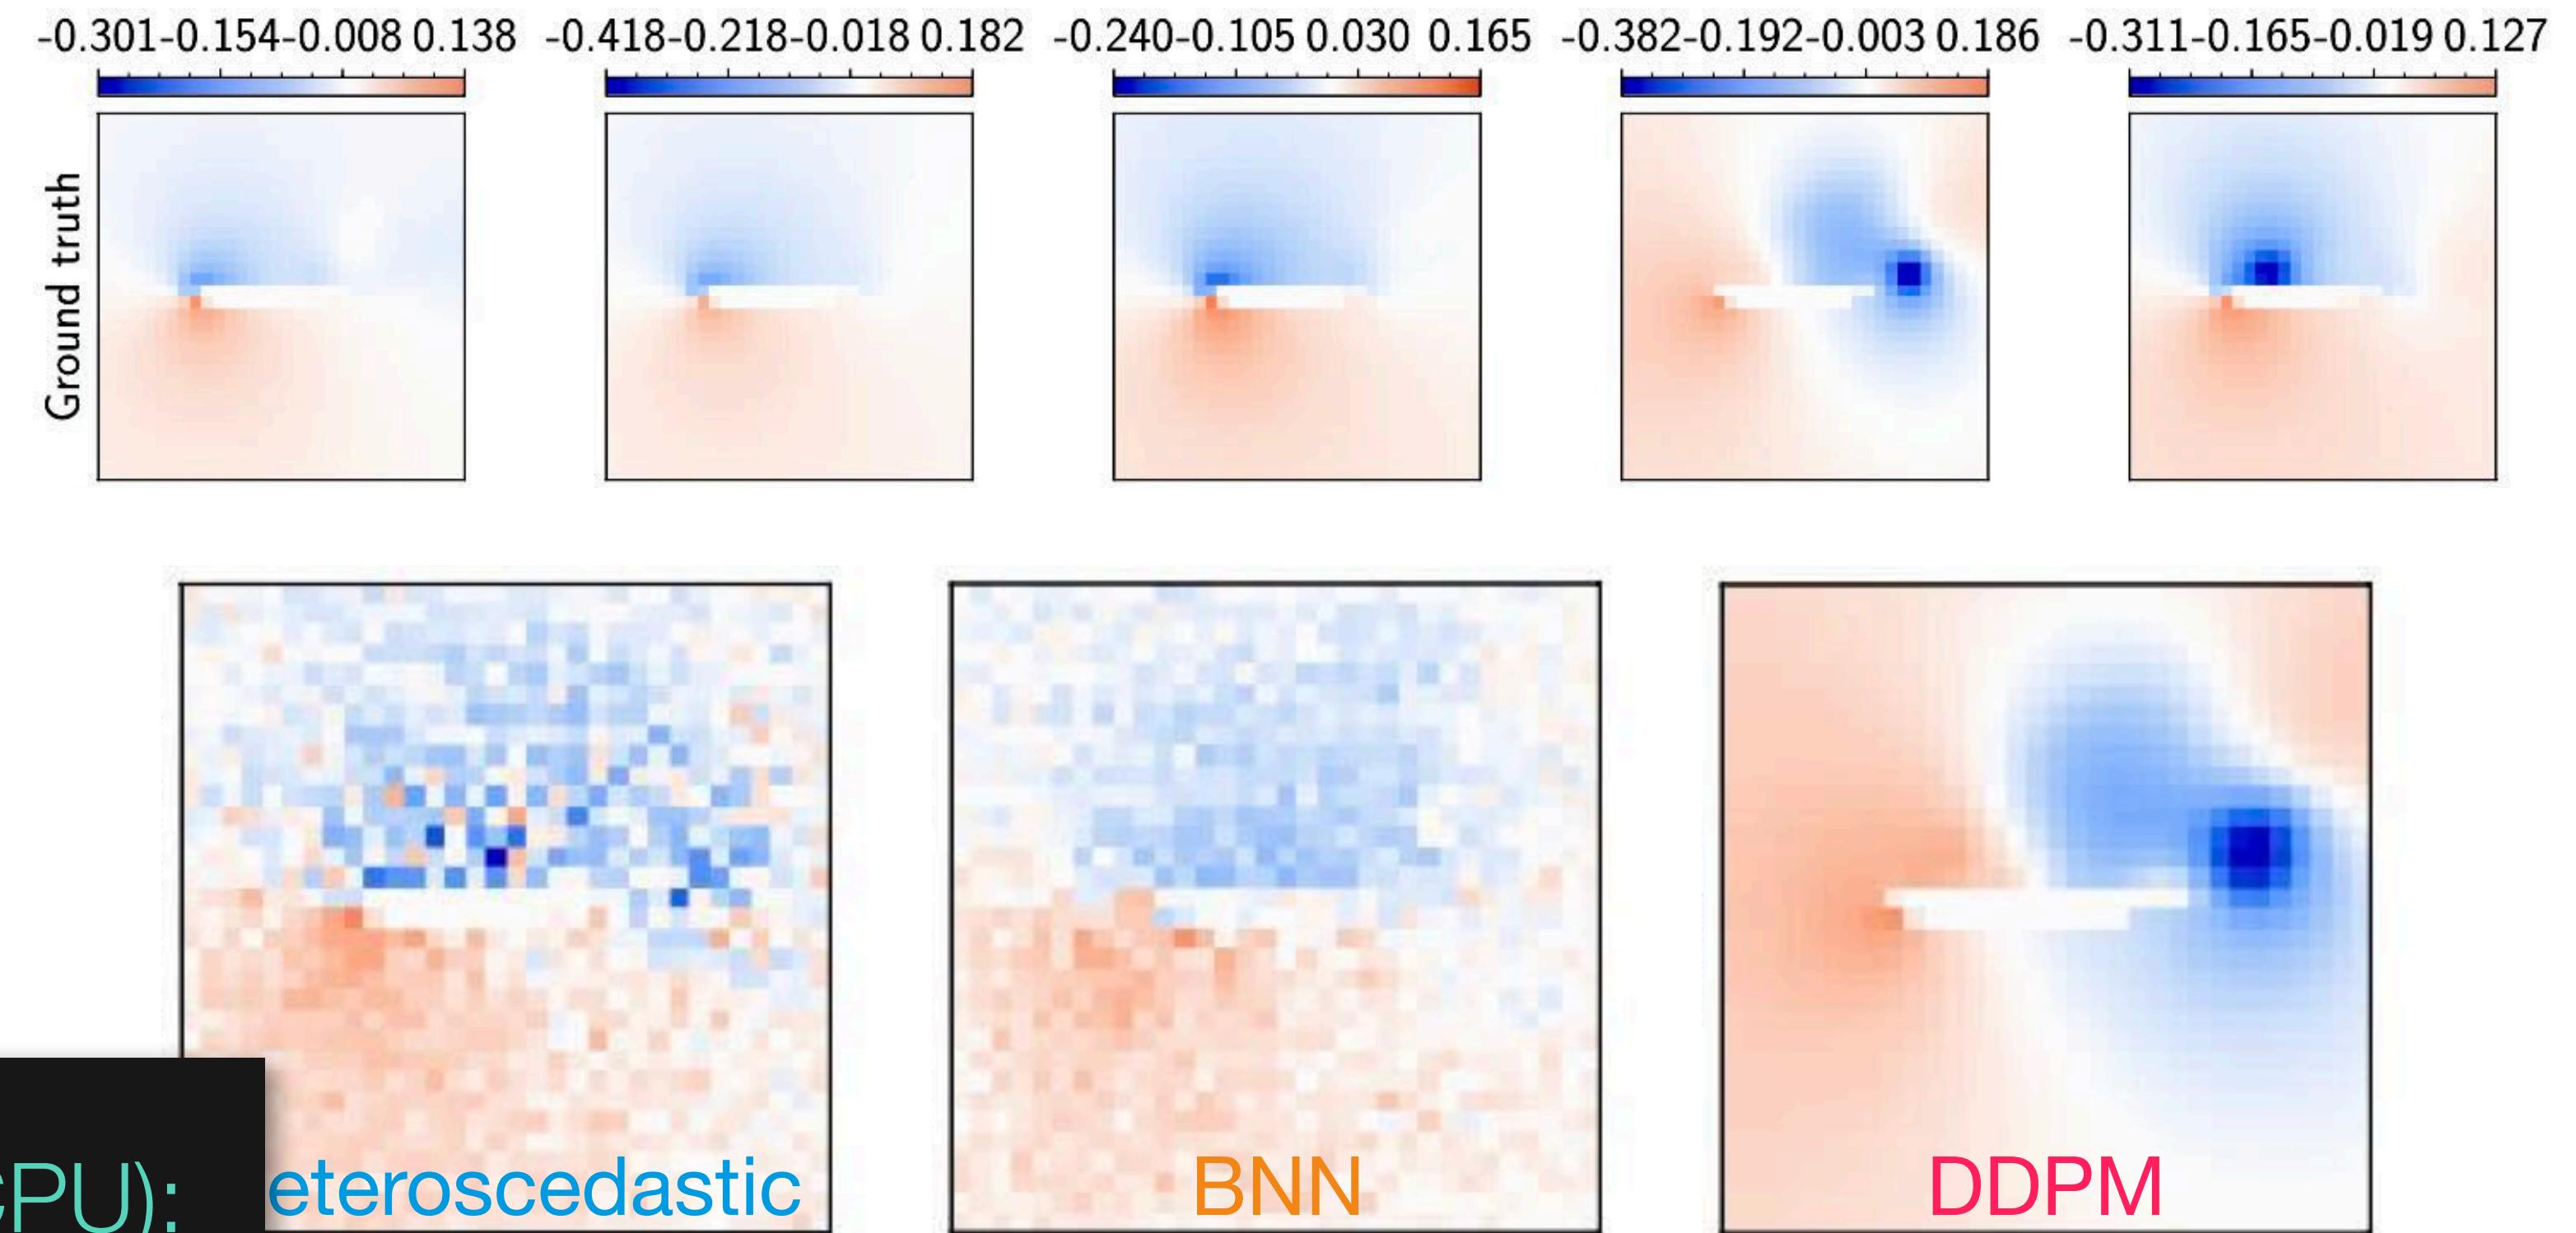
Known Ground Truth Distribution



Known Ground Truth Distribution

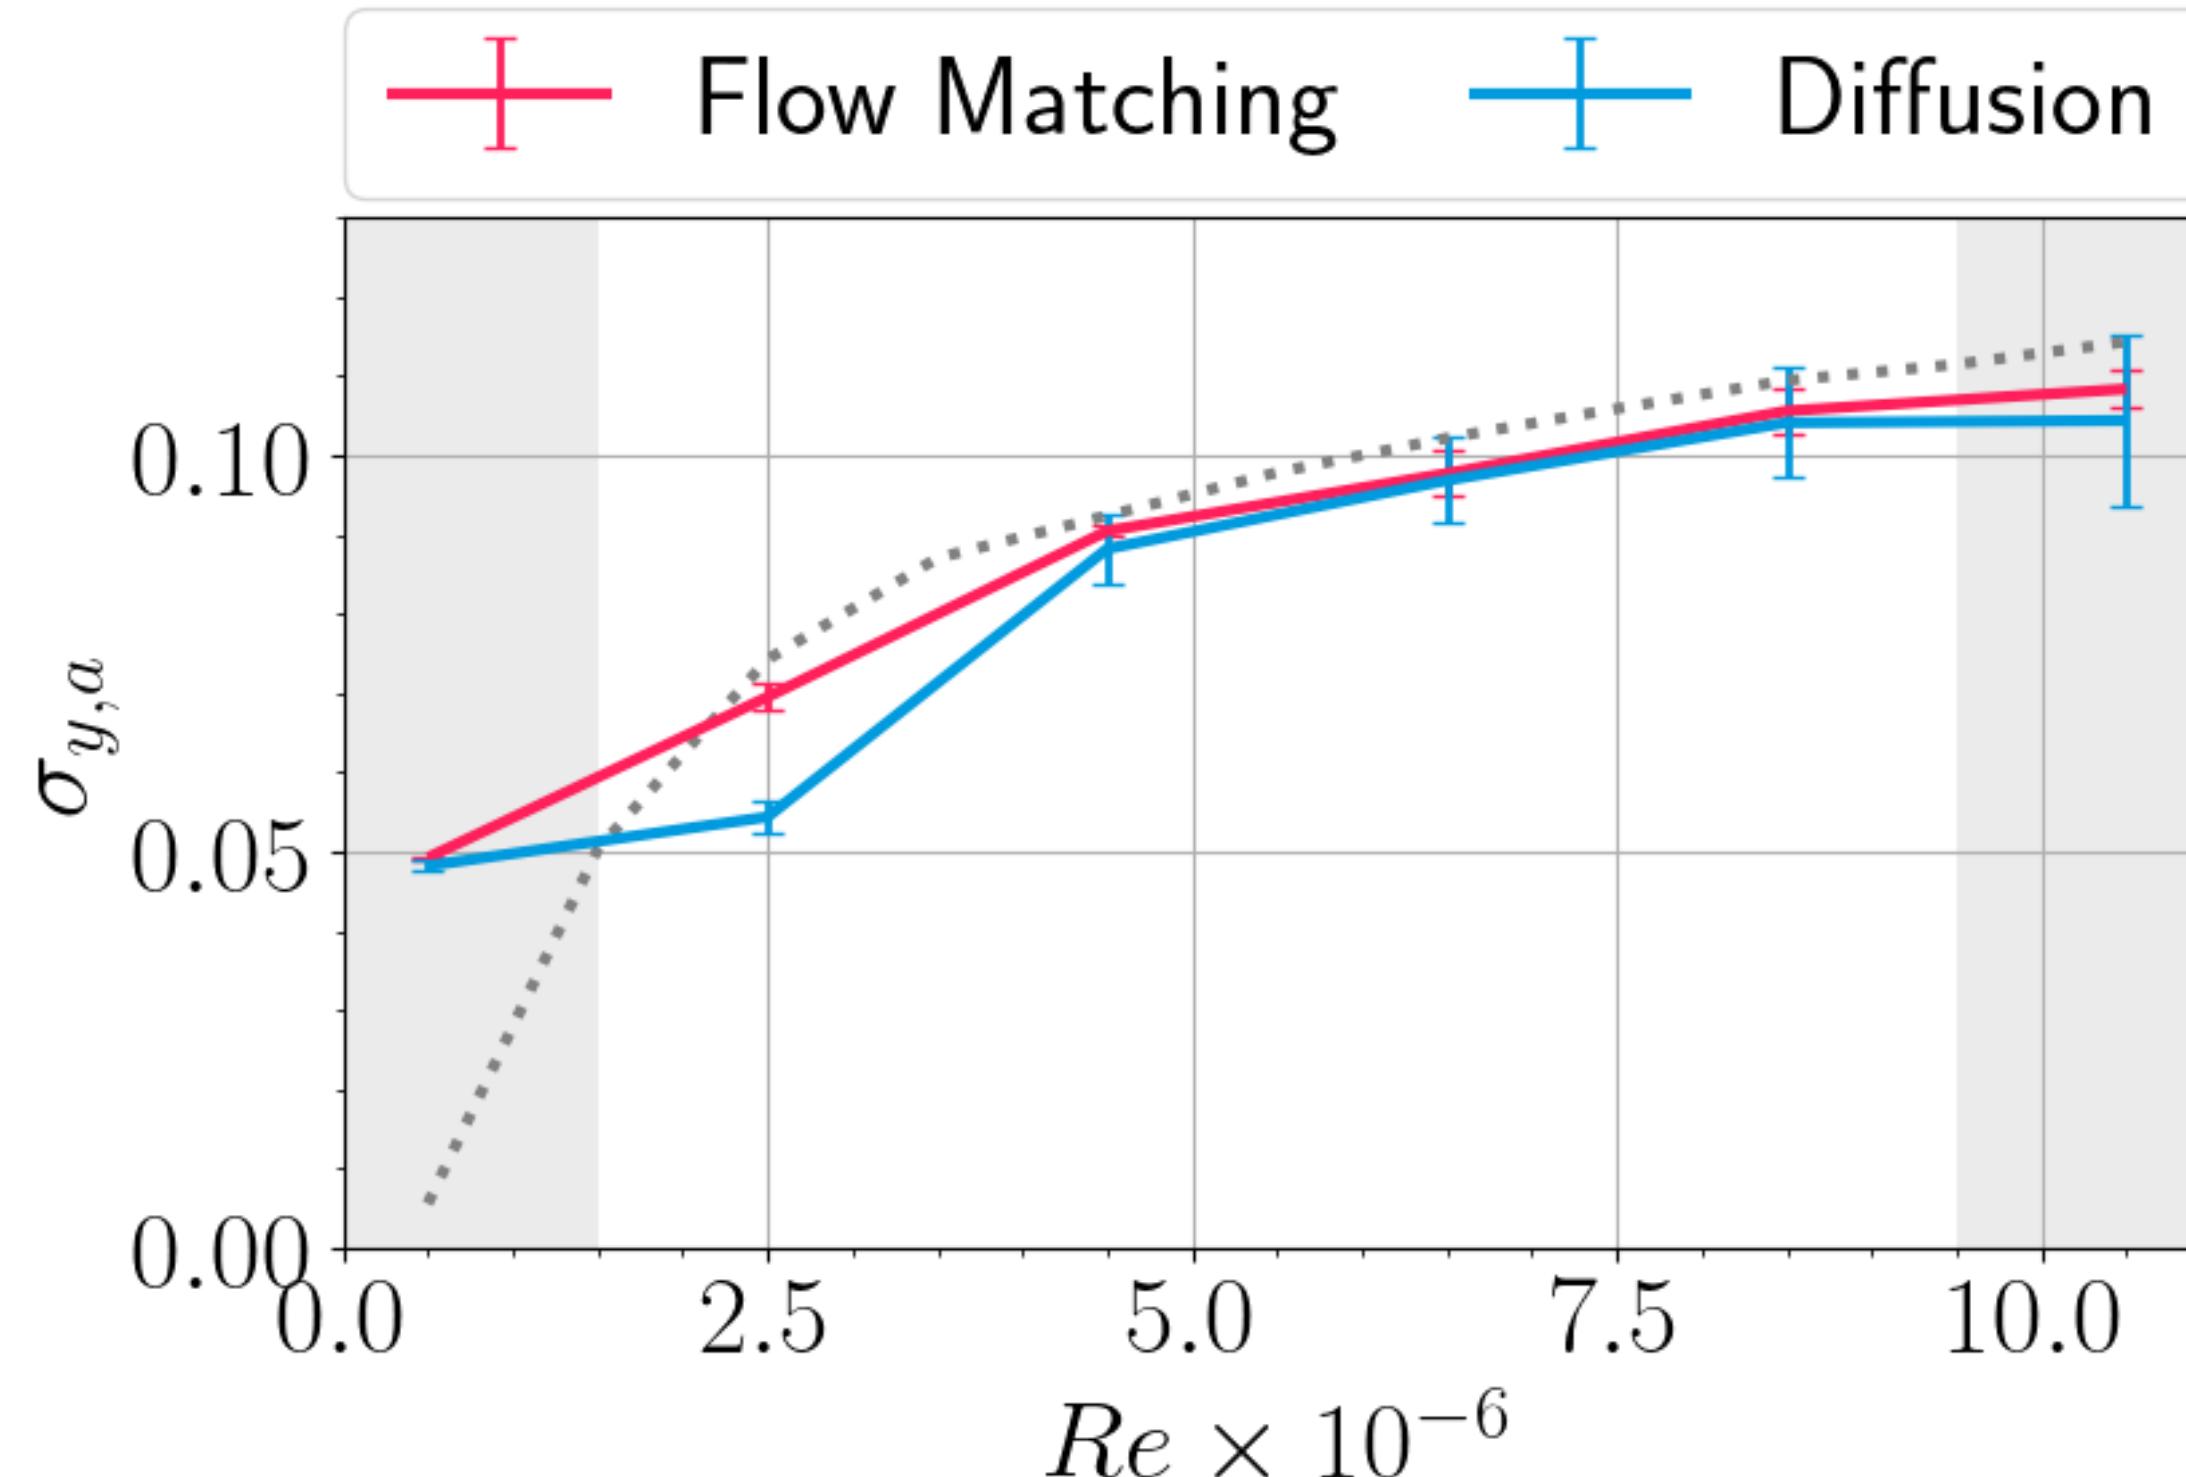
Turbulent NS case with varying
Reynolds number:

- *Heteroscedastic*
- *Bayesian NNs*
- *DDPM*



Fair comparison with solver (CPU):
DDPM ca. **9.5x** faster

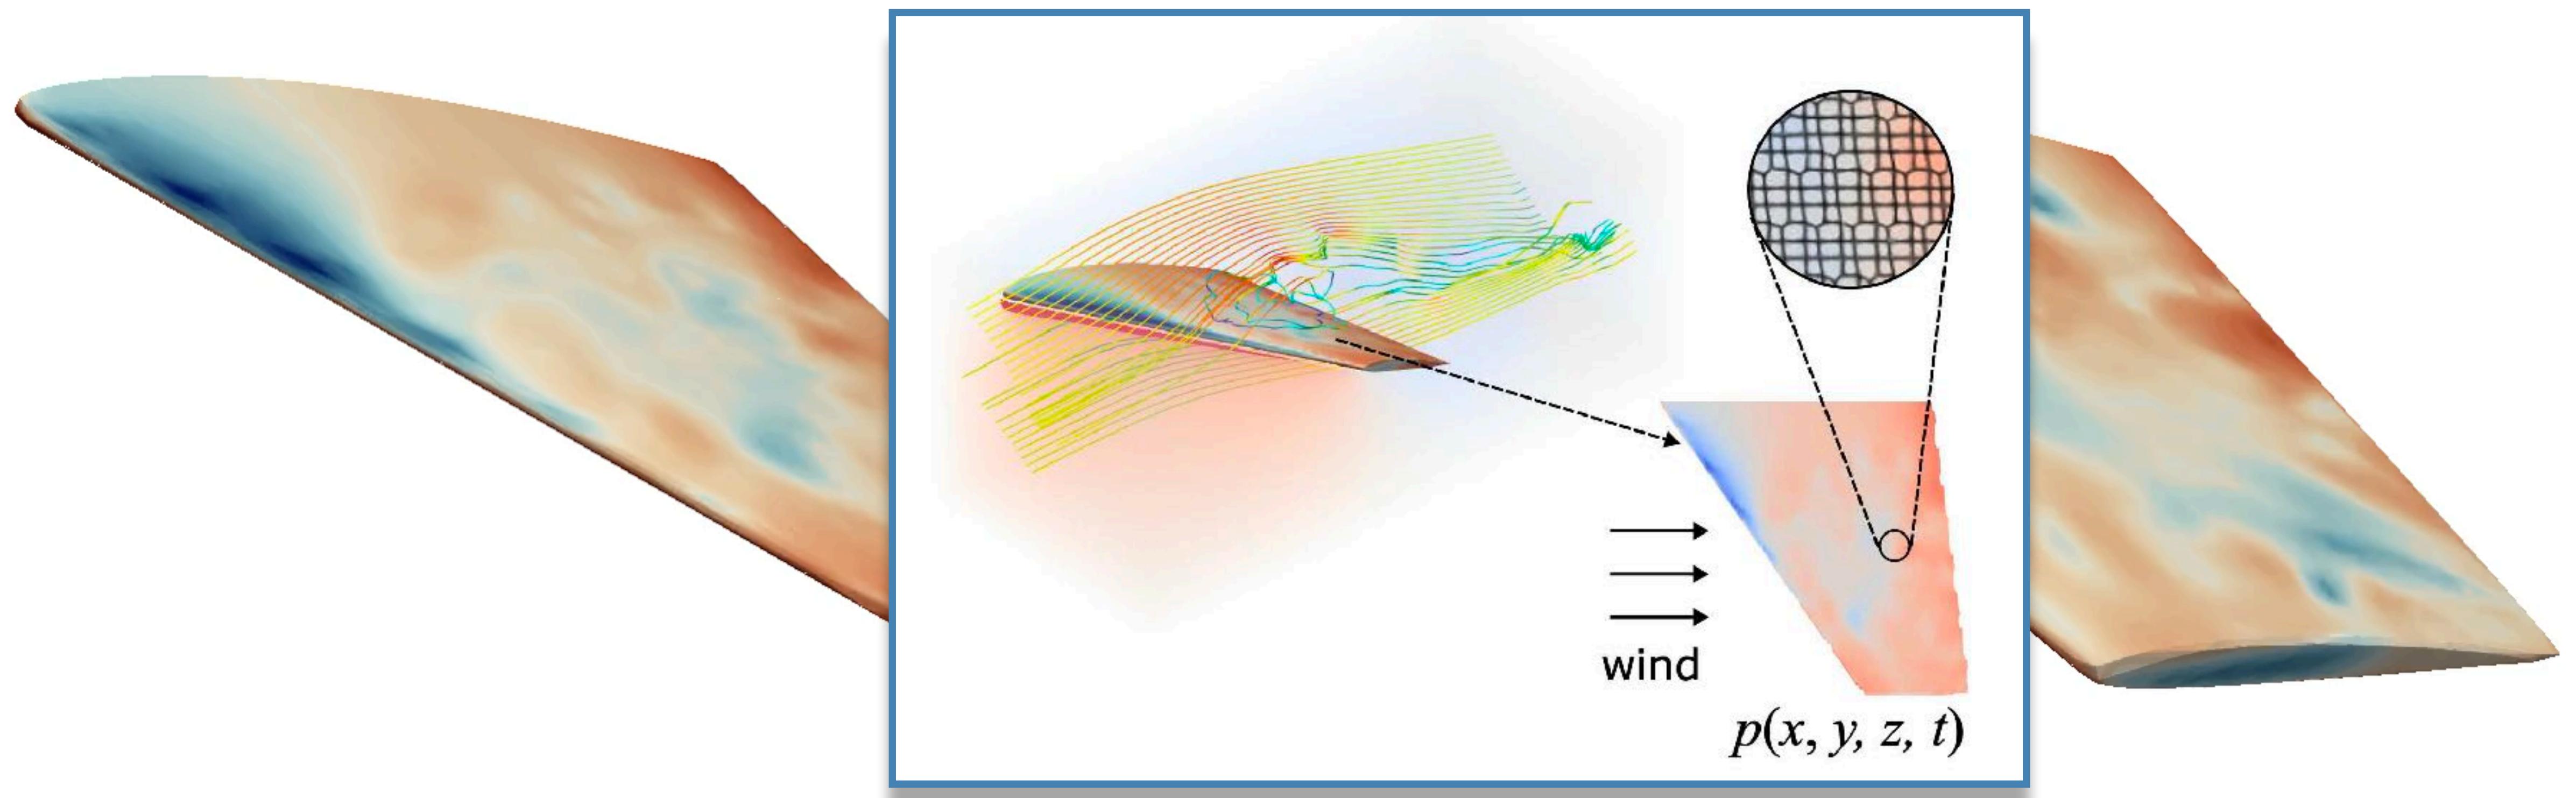
DDPM vs Flow Matching



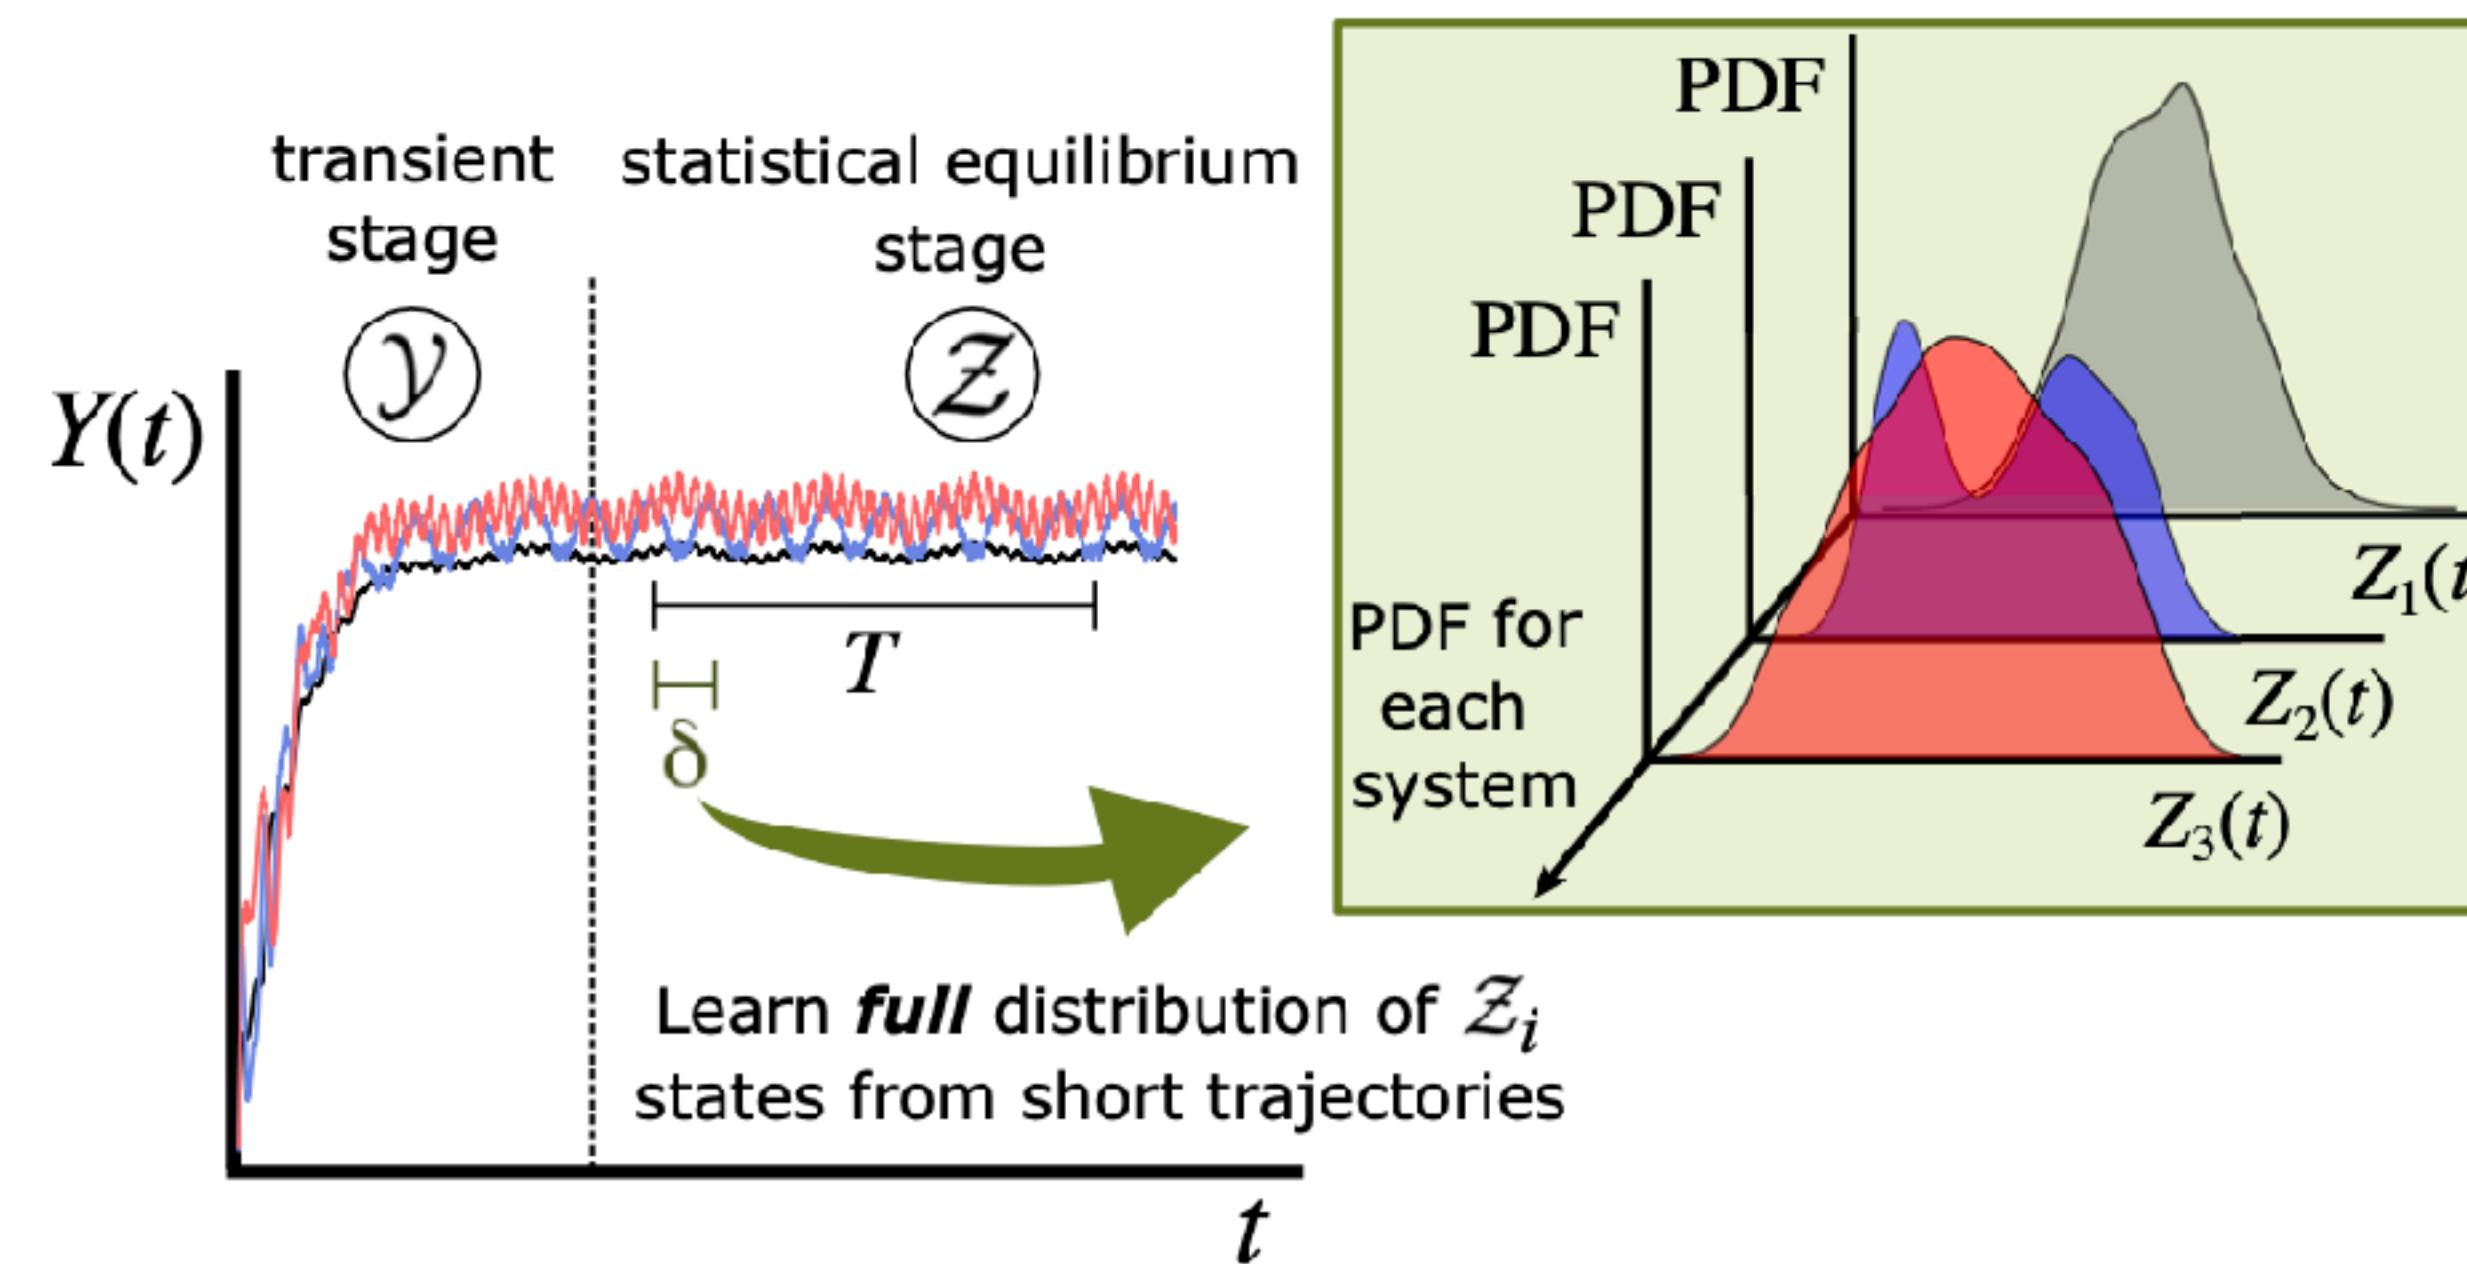
- 1D Airfoil Case again
- Diffusion iterations **n=10**
- Flow matching with n=10 roughly on-par with DDPM at n=200
- ... plus improved training stability

Flow Prediction for Unstructured Meshes

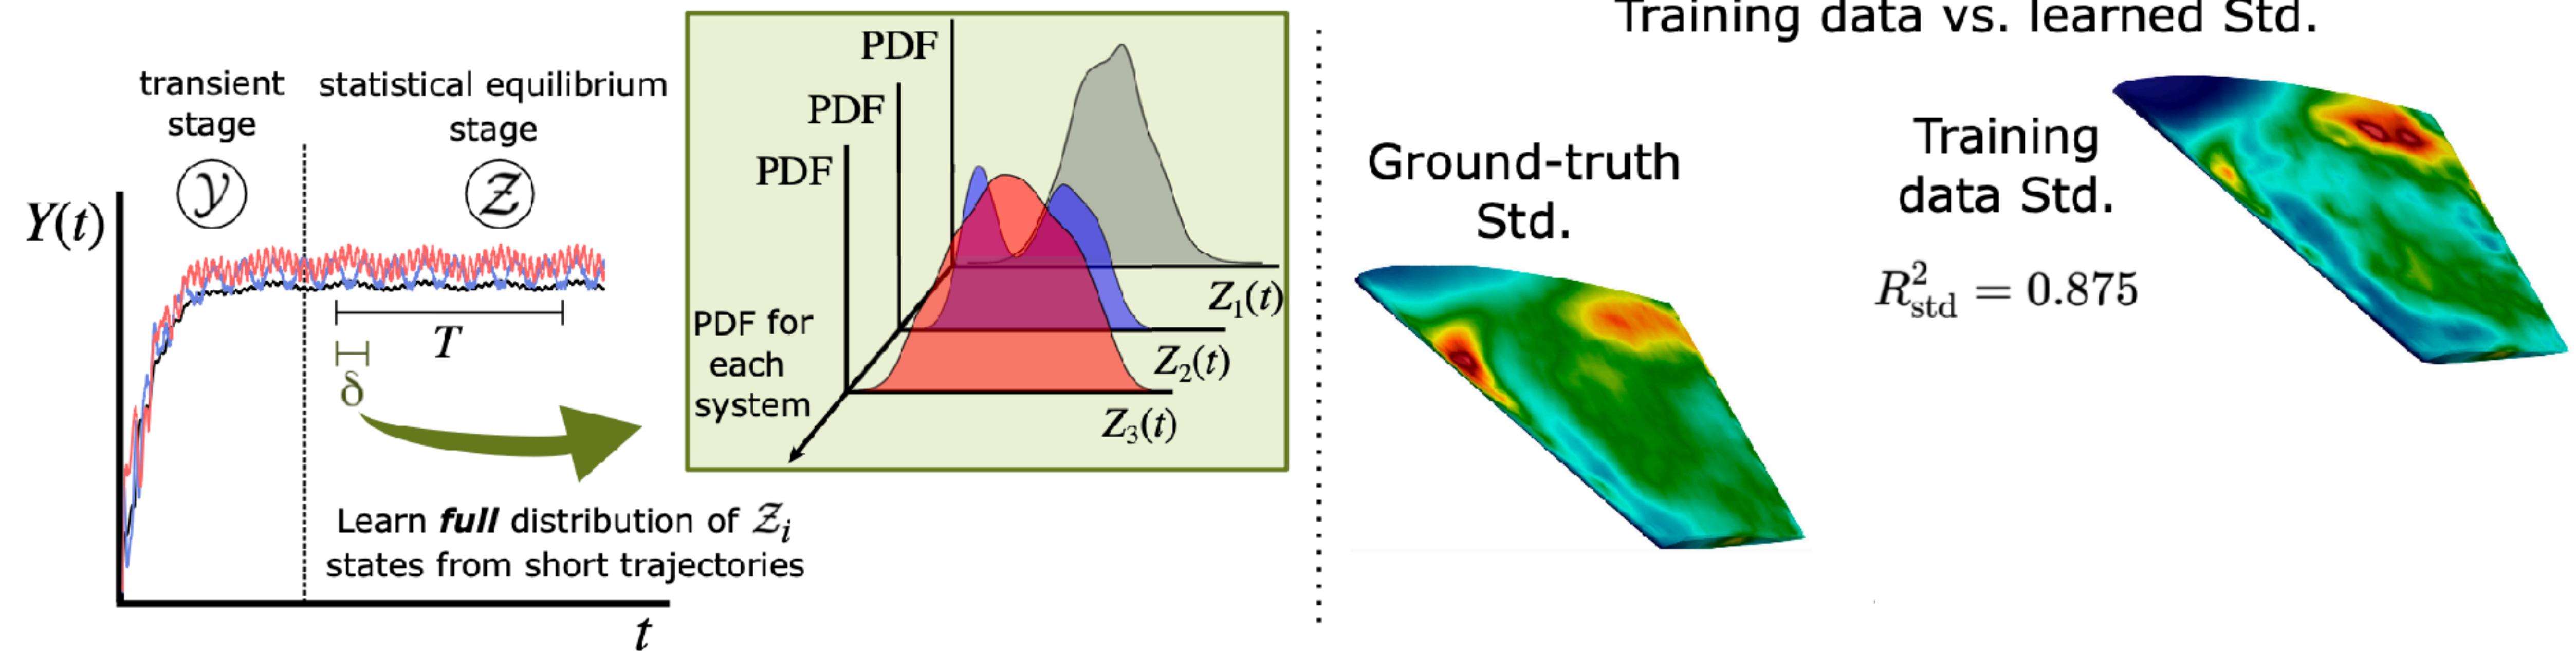
Targeting complex unsteady dynamics (no single mean solution)



Flow Prediction for Unstructured Meshes

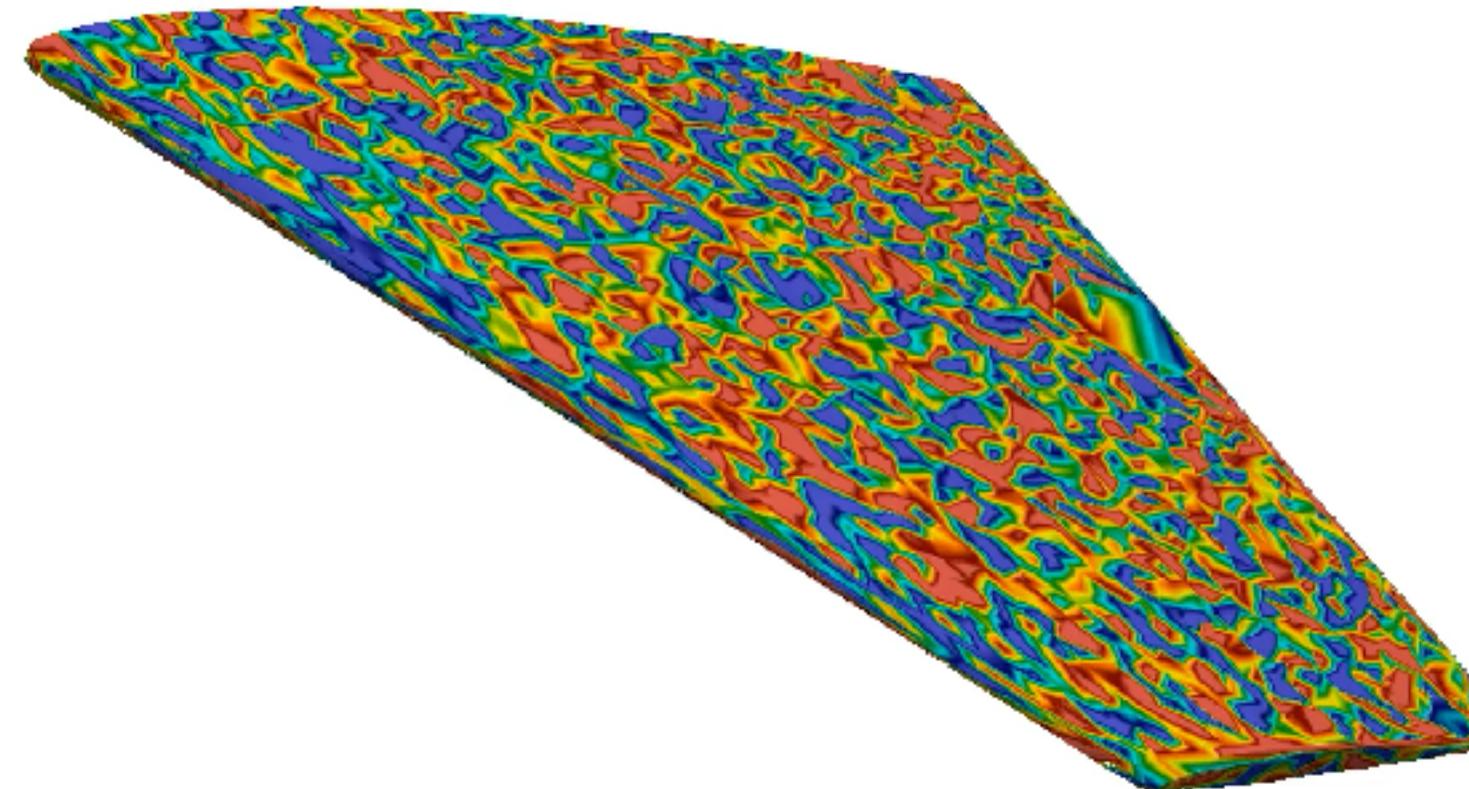


Flow Prediction for Unstructured Meshes

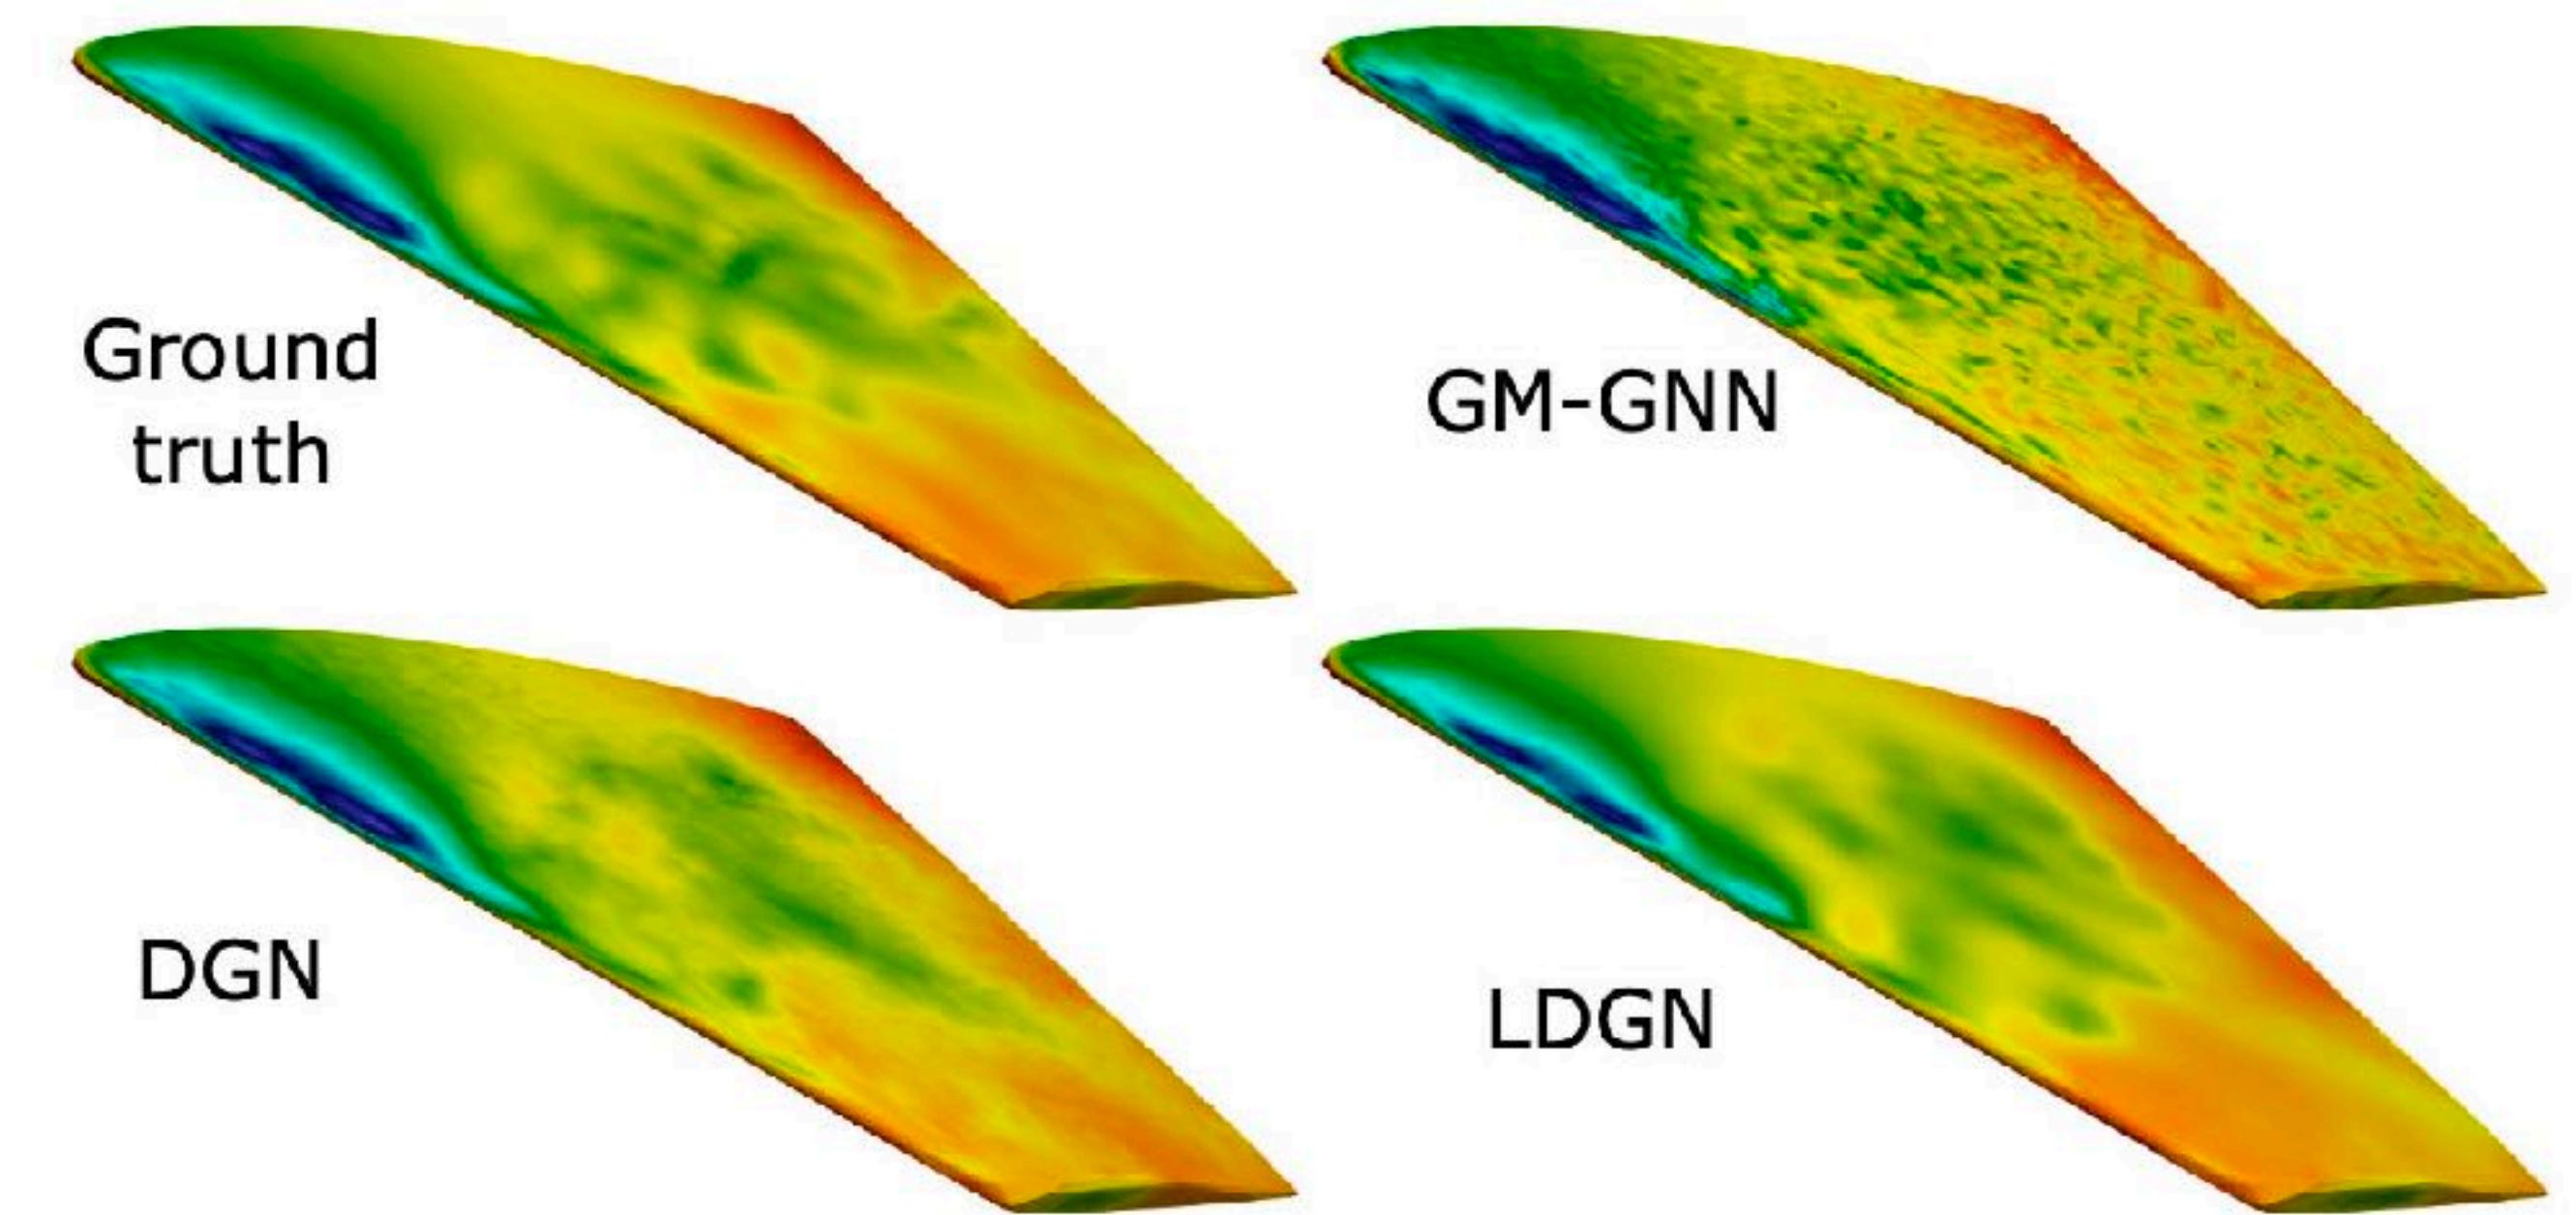


Flow Prediction for Unstructured Meshes

Denoising Process

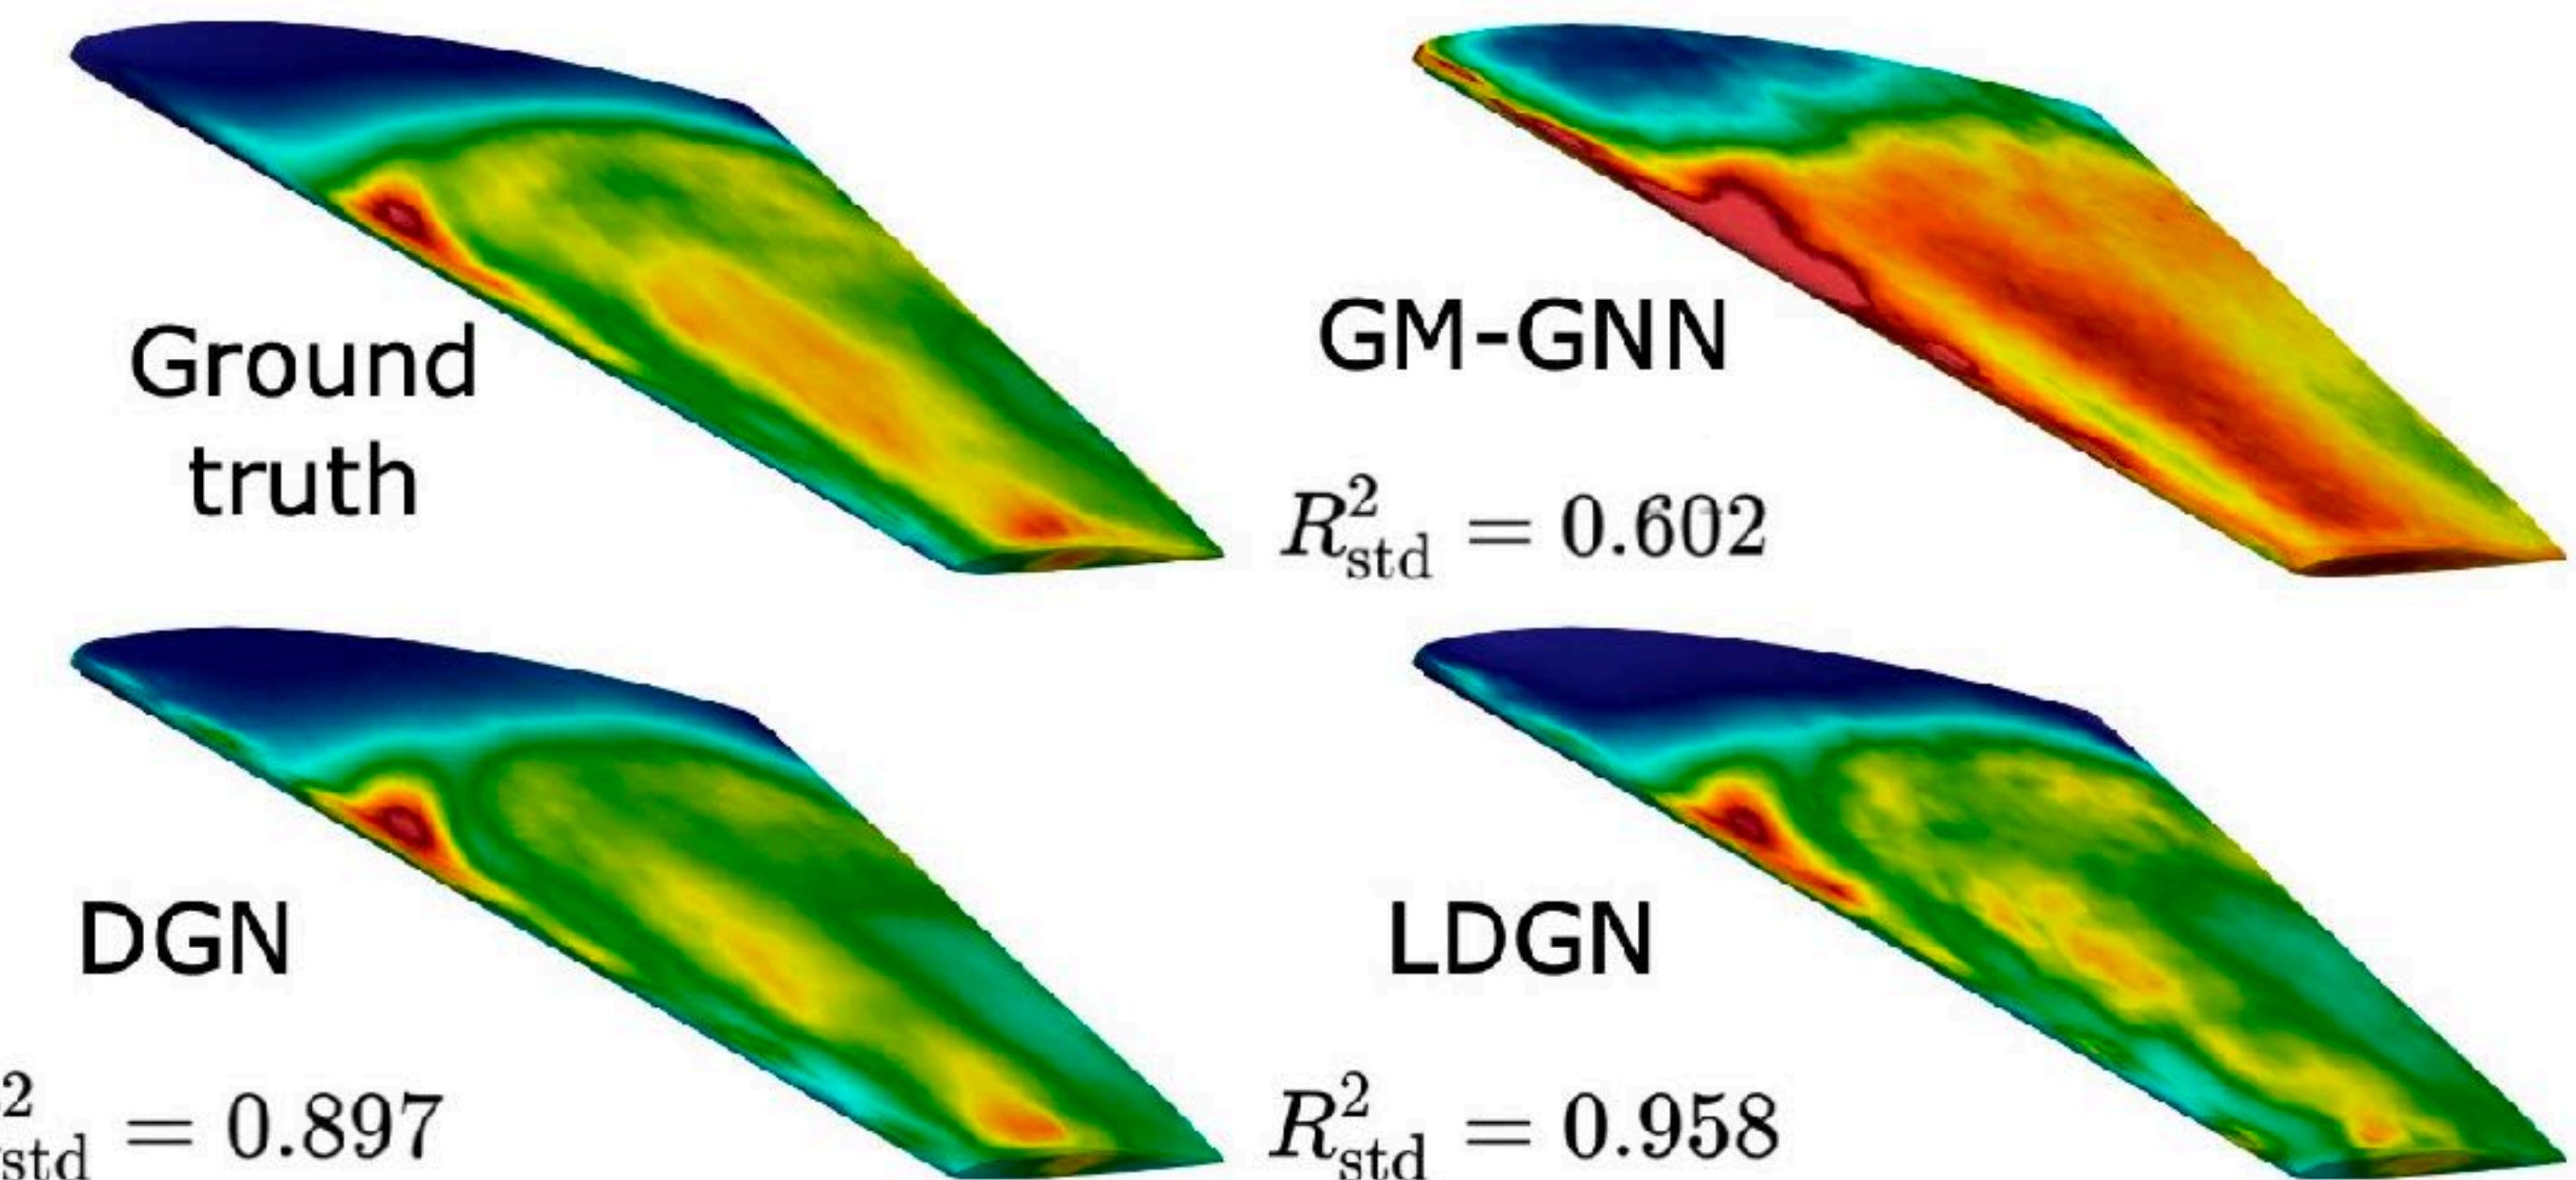


Produced Samples



Flow Prediction for Unstructured Meshes

Excellent Distributional Accuracy:
(Inferred std. dev. on test case)



Model	CPU s/sample	CPU min/distribution (speedup)	GPU s/sample	GPU min/distribution (speedup)
DGN	6.81	340 ($\times 8.8$)	0.59	19 ($\times 152$)
LDGN	0.98	49 ($\times 61$)	0.20	2.43 ($\times 1235$)

Probabilistic Models & Physical Constraints

Physics Constraints

- So far, purely focused on learning complex distributions. Powerful tool, but (akin to supervised learning) no physics priors involved
- Back to differentiable simulations & PDEs: how to include prior knowledge?
- Main options: invoke \mathcal{P} at **training or inference time**
- 1) **Inference time:** train diffusion model (DM) as usual; follow physics gradient when de-noising
- 2) **Training time:** add physics gradient at training time; inference unmodified

1) Physics Constraints at Inference Time

- (From *Physics-informed Diffusion Models*)
[Bastek et al. '25]
- For reference: regular DDPM on the right
- After step 4 add:

```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$ 
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t}} \epsilon_\theta(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$ 
5: end for
6: return  $\mathbf{x}_0$ 
```

- Predict approximate target \tilde{x}_0 (e.g., Tweedie's or PIDM) , the less noisy the better
- Evaluate physics residual and make GD step with $\lambda_{\mathcal{P}}$ as step size: $x_{t-1} - = \lambda_{\mathcal{P}} \nabla_{\mathcal{P}}(\tilde{x}_0)$
- Iterate...

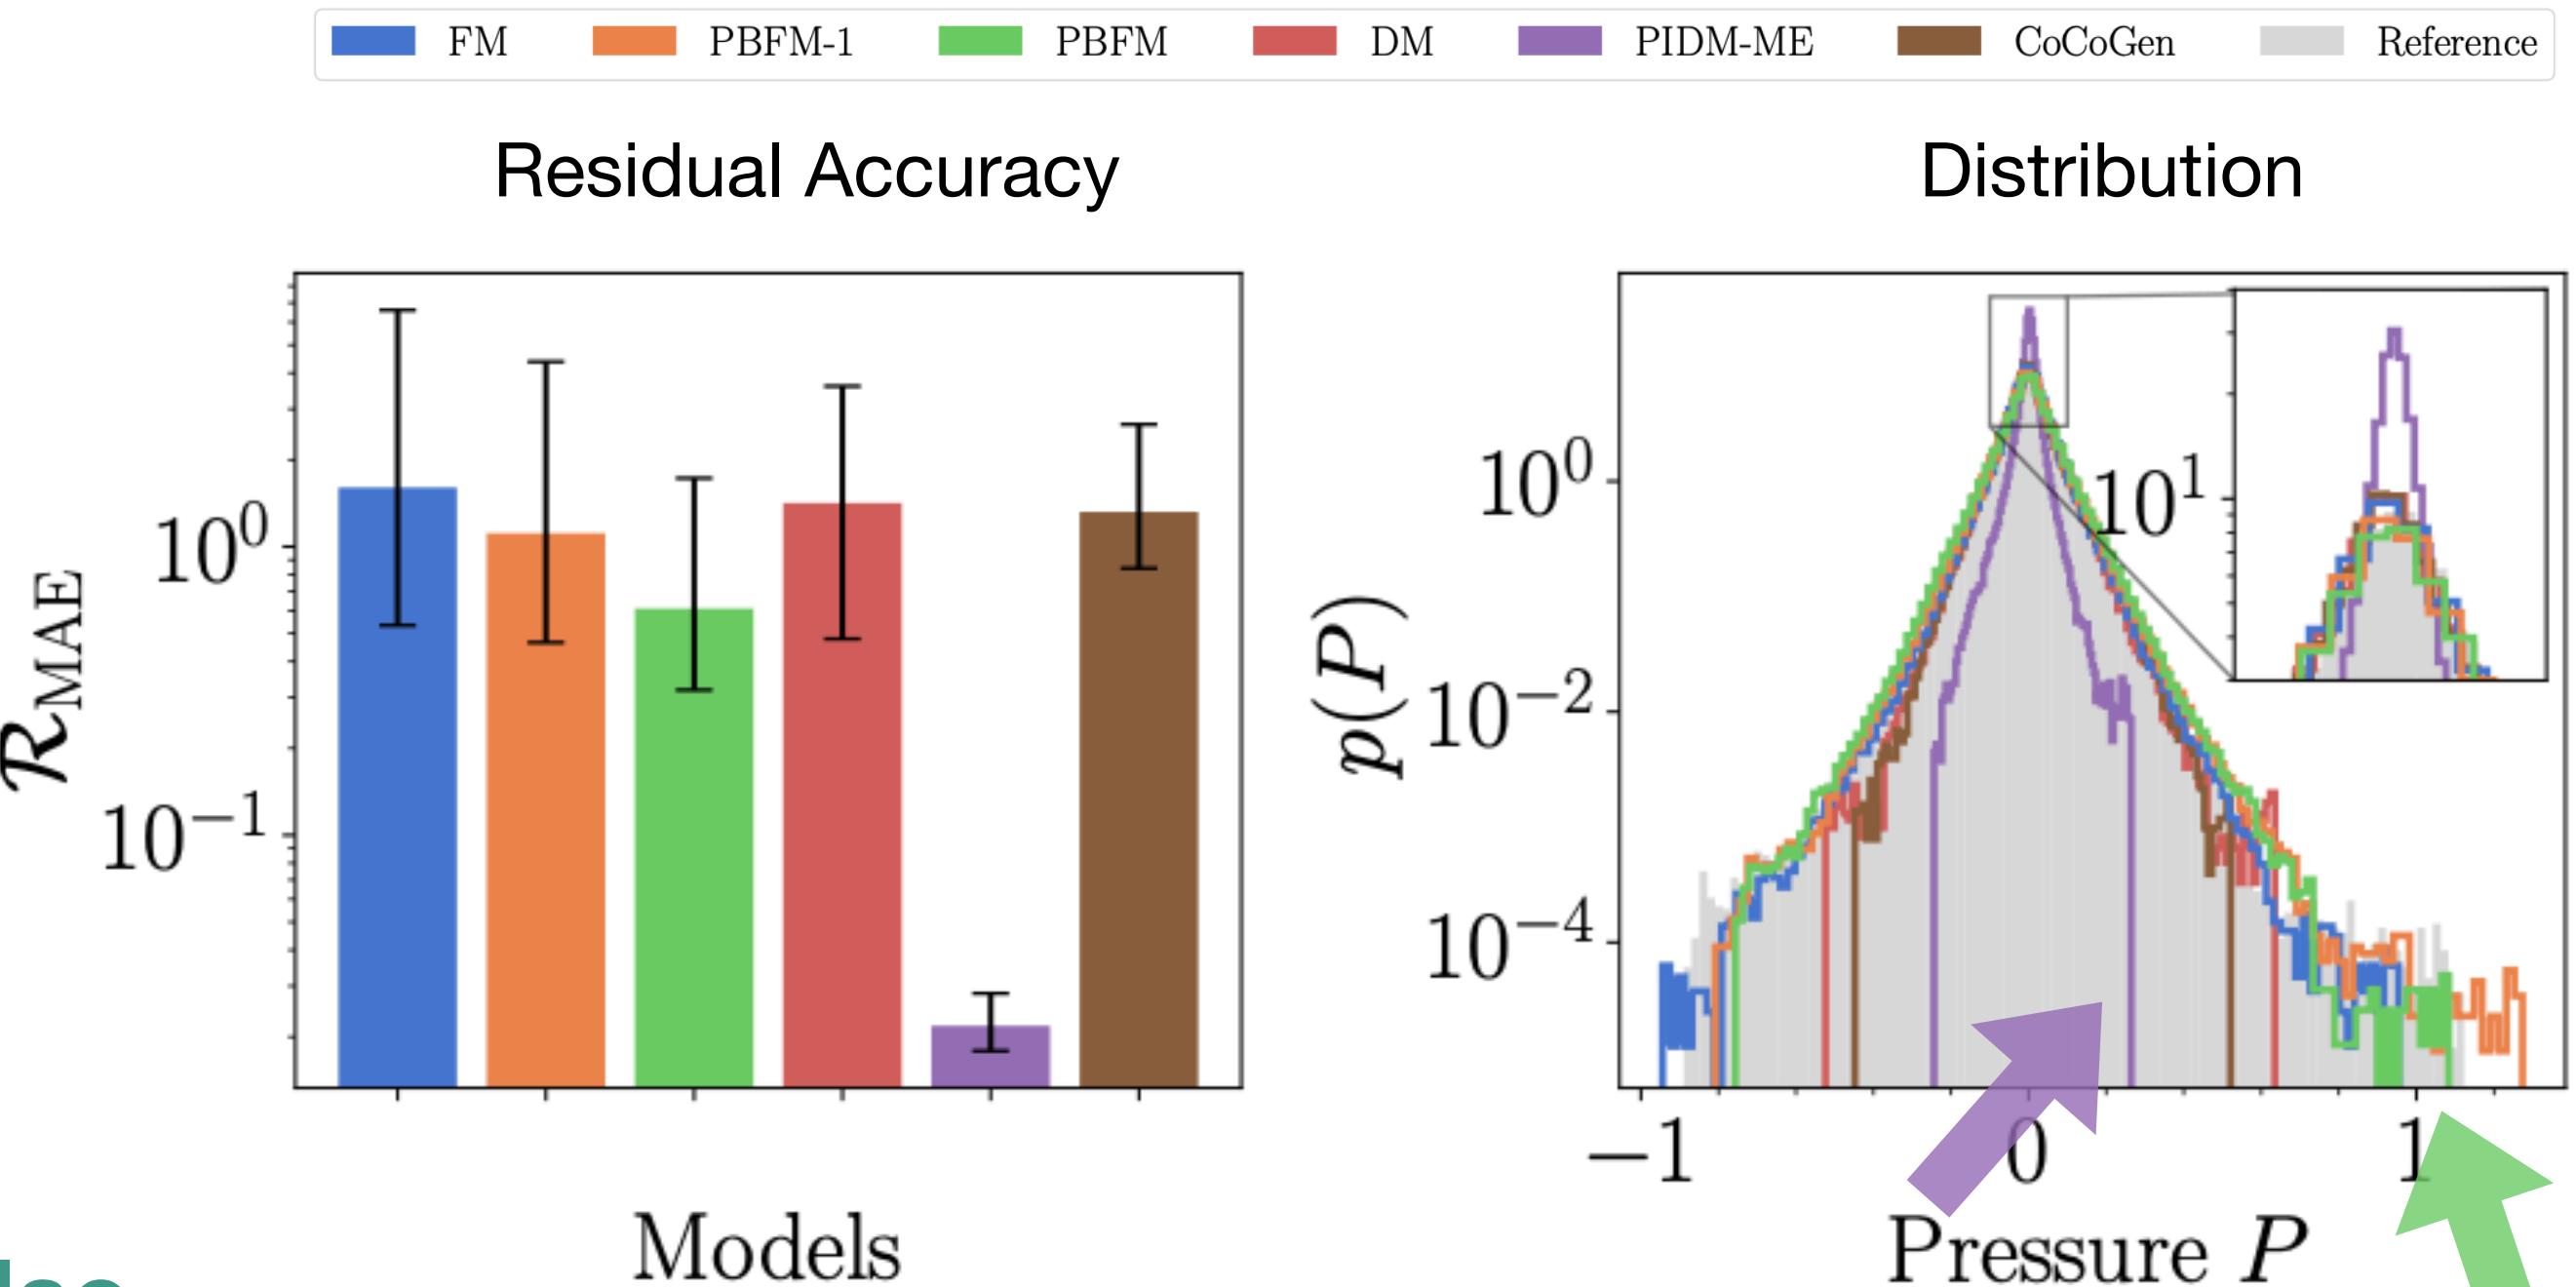
2) Physics Constraints at Training Time

- (From *Physics-based Flow Matching*)
[Baldan et al. '25]
- Similar, but evaluate $\nabla_{\mathcal{P}}$ at for each batch
- Unrolling to better predict target \tilde{x}_1 for FM
- Evaluate physics residual to make “conflict-free” step
- Advantage: no additional cost at inference time.
Disadvantage: slower training

```
n ← number of unrolling steps
dt ← (1 - t)/n
t̃ ← t
u_t^θ ← model(x_t, t)
x̃_1 ← x_t + dt · u_t^θ
for i = 1, i < n do
    t̃ = t̃ + dt
    ũ_t^θ ← model(x̃_1, t̃)
    x̃_1 ← x̃_1 + dt · ũ_t^θ
end for
R ← compute residual(x̃_1)
L_R ← ||t^p · R||_2
L_FM ← ||u_t^θ - u_t||_2
∇_θ ← compute g_update via Eq. 3
AdamW optimizer step with ∇_θ
```

2) PBFM: Model Comparison

- Standard test case: Darcy flow (porous medium)
- PIDM (purple) has very low residual error, huge errors in distribution
- PBFM (green) better residual error than e.g. regular FM, still good distribution at same comp. cost
- Main conclusion: differentiable solvers also improve probabilistic learning!



Summary & Outlook

Probabilistic Learning

- ✓ Powerful methods for learning complex distributions
- ✓ High training stability, and can handle large dimensionalities
- ✓ Important downstream tasks enabled, e.g., SBI and uncertainty quantification
- ✗ Increased training and inference cost
 - (✗ Unnecessarily slow for learning unique mappings)

Outlook

- Diffusion models can be transformed into continuous normalizing flows via the probability flow ODE → likelihood evaluation + deterministic sampling ([Song et al. 2021](#), [Lipman et al. 2023](#))
- Reduce the model size and cost by modeling high-dimensional data in a latent space ([Rombach et al. 2021](#))
- How to condition models on different inputs, e.g. text ([Rombach et al. 2019](#))
- Reduce the number of inference steps: network distillation and rectified flows ([Ho et al. 2022](#), [Liu et al. 2023](#))
- How to use diffusion models as priors for inverse problems? ([Kawar et al. 2022](#), [Chung et al. 2022](#))
- For simulation-based inference: include physics-based controls as self-conditioning to improve quality of samples ([Holzschuh et al., 2025](#))



End