


# Reinforcement Learning

## Physics and Reinforcement Learning (RL)

Common ground:

- Both use simulators to train models.
- Both treat multi-step problems.
- Both assume the Markov property.

# Part 1: RL Introduction




- RL Framework
- Value Functions
- RL Algorithms

Part 2: Ideas behind RL algorithms

# RL Framework

# RL Setup



**Agent:** learning and choosing actions.

**Environment:** responding by giving a reward and transitioning to a new state.

**Episode:**  $S_0, A_0, R_1, S_1, A_1, R_2, S_2, A_2, R_3, S_3, A_3, \dots$

**Policy:** the rule followed by the agent to choose actions     $\pi : S \rightarrow A$

**Goal:** Finding the policy that maximizes rewards

# Rewards



Immediate Reward

+1 for not falling over (each step)

+d for the amount moved along the x-axis  
(each step)



Delayed Reward


+1/0/-1 for win/draw/lose  
(at the end)

# Physics Tasks in the RL Framework

## Fluid Control

Define a loss to measure if the object is moving through the left gate.


Reward is the negative loss.



## Error Correction

Actions are the state change applied after each step

Reward is based on the similarity to the reference trajectory.




# Physics Tasks in the RL Framework II

## Learned Turbulence Models


Long term training signal via flow statistics

Fundamental topic: approximate influence of unresolved scales


Low-fidelity




SGS (Smagorinsky)



Learned




Reference




# The Idea of a Critic

A central feature of RL:

Value functions instead of rollout-based estimates



# Categories of RL methods



- Policies are learned
- Similar to most physics setups.
- Values are learned
- Policies are learned.
- Values are learned
- Policies are derived implicitly

# Value Functions

# Markov Property


The future depends only on the present, not the past.

$$p(s_1, r_1 \mid s_0, a_0)$$




# Visualization of Markov Processes


## Markov Decision Process (MDP)



## Markov Reward Process (MRP)



## Markov Process (MP)



# Value Functions


Value function:  $v_\pi(s) = \mathbb{E}[\sum_{i>t} R_i | S_t = s]$

Bellman equation: consistency equation fulfilled by the true value function


(deterministic)  $v_\pi(s_i) = v_\pi(s_{i+1}) + r$

(stochastic)  $v_\pi(s_i) = \sum_{s_{i+1}, r, a_i} p(s_{i+1}, r | s_i, a_i) \cdot \pi(a_i | s_i) \cdot (v_\pi(s_{i+1}) + r)$


# Example




Policy 1:  $\pi(s_0) = a_1, \pi(s_1) = a_2$



Policy 2:  
 $\pi(s_0) = a_1, \pi(s_1) = a_2$



# V- and Q-Values



V-Values (Policy 1):  $v(s_0) = 3.$   $v(s_1) = 1$

Q-Values:  $q_{\pi}(s, a) = \mathbb{E}[\sum_{i>t} R_i | S_t = s, A_t = a]$

$$q(s_0, a_0) = 3 \quad q(s_0, a_1) = 2 \quad q(s_1, a_2) = 1$$

# RL Algorithms

# Value Iteration

How to estimate value functions?

$$q(s, a) = \mathbb{E}[\sum_{i>t} R_i | S_t = s, A_t = a]$$

Monte-Carlo Loss

Generate a complete episode starting from state  $s_0$  and action  $a_0$  and collect all rewards  $r_i$

$$L_{MC} = (q(s_0, a_0) - \sum_i r_i)^2$$

(Supervised approach)

Temporal Difference Loss

Generate a transition from state  $s_0$  and action  $a_0$  and collect reward  $r$ , next state  $s_1$  and next action  $a_1$

$$L_{TD} = (q(s_0, a_0) - q(s_1, a_1) - r)^2$$

(Bellmann-residual approach)

How to improve policies?

Given q-values, update policy by setting:

$$\pi(s) = \operatorname{argmax}_a q(s, a)$$

Value-based RL algorithms iterate between  
value estimation and policy improvement

# Monte Carlo Learning

*Repeat:*

*Generate episode  $s_0, a_0, r_1, s_1, a_1, \dots, s_{T-1}, a_{T-1}, r_T$  following policy  $\pi$*

$$g \leftarrow 0$$

*For  $t = T - 1, \dots, 0$  repeat:*

$$g \leftarrow g + r_{t+1}$$

*Append  $g$  to Returns( $s_t, a_t$ )*

$$Q(s_t, a_t) \leftarrow \text{average}(\text{Returns}(s_t, a_t))$$

$$\pi(s_t) \leftarrow \text{argmax}_a Q(s_t, a)$$

# Temporal Difference Learning

*Repeat:*

Q-Learning

$$t \leftarrow 0$$

*Initialize*  $s_0$

*Repeat until episode terminates:*

*Generate next step  $a_t, r_{t+1}, s_{t+1}$  by following policy  $\pi$*

$$\Delta \leftarrow r_{t+1} + \max_a Q(s_{t+1}, a) - Q(s_t, a_t)$$


$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \Delta$$

$$\pi(s_t) \leftarrow \operatorname{argmax}_a Q(s_t, a)$$


$$t \leftarrow t + 1$$

# Grid World

## Environment



## Q-Values



# Function Approximation in RL

Function approximation is needed if the number of states or actions is too large or even continuous.

Approximating values with neural networks parametrized by  $\theta$ :

$$v(s) \rightarrow v_\theta(s)$$

Training through backpropagation of value errors:

$$\Delta\theta = \frac{\partial v}{\partial \theta} \Delta v$$

# DDPG: An actor-critic method

## Deep deterministic policy gradient

Data generation

Replay buffer

Value iteration

Policy iteration

Target networks

Randomly initialize critic network  $Q(s, a|\theta^Q)$  and actor  $\mu(s|\theta^\mu)$  with weights  $\theta^Q$  and  $\theta^\mu$ .

Initialize target network  $Q'$  and  $\mu'$  with weights  $\theta^{Q'} \leftarrow \theta^Q$ ,  $\theta^{\mu'} \leftarrow \theta^\mu$

Initialize replay buffer  $R$

**for** episode = 1, M **do**

    Initialize a random process  $\mathcal{N}$  for action exploration

    Receive initial observation state  $s_1$

**for** t = 1, T **do**

        Select action  $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$  according to the current policy and exploration noise

        Execute action  $a_t$  and observe reward  $r_t$  and observe new state  $s_{t+1}$

        Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $R$

        Sample a random minibatch of  $N$  transitions  $(s_i, a_i, r_i, s_{i+1})$  from  $R$

        Set  $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'})|\theta^{Q'})$

        Update critic by minimizing the loss:  $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$

        Update the actor policy using the sampled policy gradient:

$$\nabla_{\theta^\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

Update the target networks:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$

# Part 2: Why RL Needs Specialized Techniques



Components of a learning system: one-step -> multi-step methods

- **Model:** autoregressive models -> value functions
- **Loss:** supervised -> residual-based
- **Optimization:** gradient methods-> non-gradient methods

**Stochasticity** is a key property to motivate these changes.

# Autoregressive Models vs Value Functions


# Prediction: One-Step and Multi-Step Methods

Autoregressive model  $s_t = F(s_{t-1})$

- Perfect a model on one-step predictions
- Apply it recursively to receive a multi-step prediction

Value function  $\sum_{t>t} r_i = F(s_t)$

- A multi-step method



# Multi-Step Predictors on Differential Equations



Differential equation and initial condition:

$$\partial_t u(t) = P(t, u, \dots) \quad \text{with} \quad u(0) = u_0$$

How to learn the solution of this differential equation?

Discretize on a grid  $u_i = u(t_i)$

- (One-step approach) Autoregressive model:  $u_{i+1} = F(u_i)$
- (Multi-step approach) Directly parametrize the solution  $u_i = F(t_i)$ .

## Autoregression

- Requires N model calls for an N-step prediction
- Learns the entire state trajectory  
(in stochastic settings, all occurring state sequences)
- Exponential error growth over time

# Example: Ising Model




Goal: learn to predict who wins

Autoregressive models learn the complete environment dynamics.

Value functions learn possible shortcuts (e.g. number of white cells minus number of black cells)

# Example: Cart Pole Prediction



Method	AR1	Unrolled AR1	AR10	Value-based
Parameters	9000	9000	9000	9000
Training Time [sec]	300	800	300	300
Prediction Time [sec]	270	270	295	6

# Supervised vs Residual-Based Losses

# Update Equations

## Supervised Learning

Monte Carlo (MC):

$$\Delta v(s_t) = \alpha \cdot \left( \sum_{i>t} r_i - v(s_t) \right)$$

## Trajectory Learning

$$\Delta s_t = \alpha \cdot \left( P^t(s_0) - s_t \right)$$

$v$  value     $s_t$  state at time  $t$      $r_t$  reward at time  $t$      $\alpha$  learning rate     $P$  physics operator

## Residual-Based Learning

Temporal Difference (TD):


$$\Delta v(s_t) = \alpha \cdot \left( r_t + v(s_{t+1}) - v(s_t) \right)$$

## Physics-Informed Training

$$\Delta s_{t+1} = \alpha \cdot \left( P(s_t) - s_{t+1} \right)$$

# Example: 3-State Value Estimation

How would you estimate the values in this Markov reward process based on these observed episodes?




Episode 1: state A, reward 0, state C, reward -1

Episode 2: state B, reward 0, state C, reward +1

# Example: 3-State Value Estimation


Episode 1: state A, reward 0, state C, reward -1

Episode 2: state B, reward 0, state C, reward +1



Monte Carlo

MC minimizes the error on [past data](#).




Temporal Difference

TD minimizes the error on [future data](#) (through the Markov property).

# Example: Website Design


Task: Select website design that leads to the largest sale rate.



Which one would you choose based on this data?

	Website 1	Website 2	Website 3
Visitors	10	8	7
Proceeded to Checkout	6	5	4
Purchase	4	2	2

# Variance of MC and TD Estimators



# Example: Going Home

	Current prediction	New observation	MC	TD
University	40 min left		+	0
University subway station	35 min left	took 5 min	+	-
Home subway station	5 min left	took 25 min	+	+
Home	0 min left	took 15 min		

# Gradient vs Non-Gradient Methods


## Classical Methods

- Newton's Method
- Gradient Descent
- Momentum

## Our Interest

- First-order
- Not Gradient Descent
- Not gradient-based

# TD Updates are Non-Gradient Updates



TD loss for step  $t$

$$L_t = (r_t + v(s_{t+1}) - v(s_t))^2$$

Non-gradient update

$$\Delta v(s_t) = \alpha \cdot (r_t + v(s_{t+1}) - v(s_t))$$

Gradient update


$$\Delta v(s_t) = \alpha \cdot (r_t + v(s_{t+1}) - v(s_t))$$

$$\Delta v(s_{t+1}) = -\alpha \cdot (r_t + v(s_{t+1}) - v(s_t))$$


$v$  value     $s_t$  state at time  $t$      $r_t$  reward at time  $t$      $\alpha$  learning rate

# Information Flow


Rewards (edges) and corresponding values (nodes)



New observation



Earlier values have to be changed, not later values.




# Information Flow II


One particle in a uniform force field:  $\ddot{x} = F$  with  $x(0) = \dot{x}(0) = 0$

$$(x_{i+1}, p_{i+1}) = (x_i + p_i \cdot dt, p_i + F \cdot dt) \approx (x_i + p_i, p_i + 1)$$

Current Approximation




Later predictions have to be changed, not earlier predictions.




# Convergence Rate

Value prediction on this  
Markov reward process



Approximate Convergence Rate



# Function Approximation

TD loss for step  $t$

$$L_t = (r_t + v_\theta(s_{t+1}) - v_\theta(s_t))^2$$

Gradient update

$$\Delta\theta = \alpha \cdot (r_t + v_\theta(s_{t+1}) - v_\theta(s_t)) \cdot (\partial_\theta v_\theta(s_t) - \partial_\theta v_\theta(s_{t+1}))$$

Non-gradient update

$$\Delta\theta = \alpha \cdot (r_t + v_\theta(s_{t+1}) - v_\theta(s_t)) \cdot \partial_\theta v_\theta(s_t)$$

$v$  value     $s_t$  state at time  $t$      $r_t$  reward at time  $t$      $\alpha$  learning rate     $\theta$  parameters


# Convergence Properties

Divergence of non-gradient TD

Solution:  $w = 0$

Non-gradient update step:  $\Delta w = \alpha w$

$$v(A, w) = w$$



Convergence theorem for non-gradient TD

Linear function approximation

On-Policy Sampling

Linear function approximation with quadratic loss

$$L(\theta) = (F\theta - c)^2$$

Gradient update

$$\partial_{\theta} L = F^{\dagger}(F\theta - c)$$

Gradient fixed-point

$$\theta_{Gradient} = (F^{\dagger}F)^{-1}F^{\dagger}c$$

Non-gradient update

$$B(F\theta - c)$$

Non-gradient fixed-point

$$\theta_{Non-Gradient} = (BF)^{-1}Bc$$

# The Problem with Gradients on Residual Objectives



- Even though we don't use it for training, our goal is to minimize a supervised loss.
- We are interested in a residual loss for training only because of its statistical benefits.
- Classical optimization methods attempt to find a minimum of the loss they are applied to.

Supervised objectives

+





Gradient methods

Residual objectives

+

Non-gradient methods

# Harmonic Oscillator, Physics-Informed Training



## Reinforcement Learning

- explores alternative approaches for multi-step tasks (value functions, temporal difference learning, non-gradient methods).
- shares deep analogies with physical learning techniques.
- addresses key computational, statistical and optimization bottlenecks.

# References

- R. S. Sutton, Andrew G. Barto: ‘Reinforcement Learning, An Introduction’
- R. S. Sutton’s personal homepage: <http://www.incompleteideas.net/>
- L. C. Baird: ‘Reinforcement Learning through Gradient Descent’
- T. P. Lillicrap et al.: ‘Continuous Control with Deep Reinforcement Learning’
- D. Cheikhi & D. Russo: ‘On the Statistical Benefits of Temporal Difference Learning’
- P. Schnell et al.: ‘Temporal Difference Learning: Why It Can Be Fast and How It Will Be Faster’